

FLAG: Adversarial Data Augmentation for Graph Neural Networks

Kezhi Kong¹, Guohao Li², Mucong Ding¹, Zuxuan Wu¹, Chen Zhu¹,
Bernard Ghanem², Gavin Taylor³, Tom Goldstein¹

¹University of Maryland, ²KAUST, ³US Naval Academy

The Learning on Graphs and Geometry Reading Group (LoGaG) Presentation
2021.09.28

Background

- Data Augmentation is important to generalization
 - For graphs, we have **structural-based** augmentations:
 - Neighbor Sampling, DropEdge, Virtual Node, Edge Permutation...
- Question: **feature-based** augmentations for graphs?

Background




Figure 1: Generalization performance of FLAG on all three tasks. Left: node classification with GAT as baseline on ogbn-products; Middle: link prediction with hits@20 as metric (the higher the better) and GraphSAGE as baseline on ogbl-ddi; Right: graph classification with GIN as baseline on ogbg-molhiv. Plotted lines are attained by smoothing the original lines (the shallow ones), where smooth weights are 0.75, 0.75, and 0.5 respectively.

Adversarial Training Is the Cure!

Contributions

- Our work is the **first** general-purpose feature-based data augmentation method on graph data
- The method works on all the three major graph **tasks** (node, link, and graph)
- The method has good **scalability** and works on large-scale datasets
- The method is **easy** to implement and use
- The method is **complementary** to existing regularizers (Dropout) and graph structure augmentations (Neighbor sampling & Virtual node)

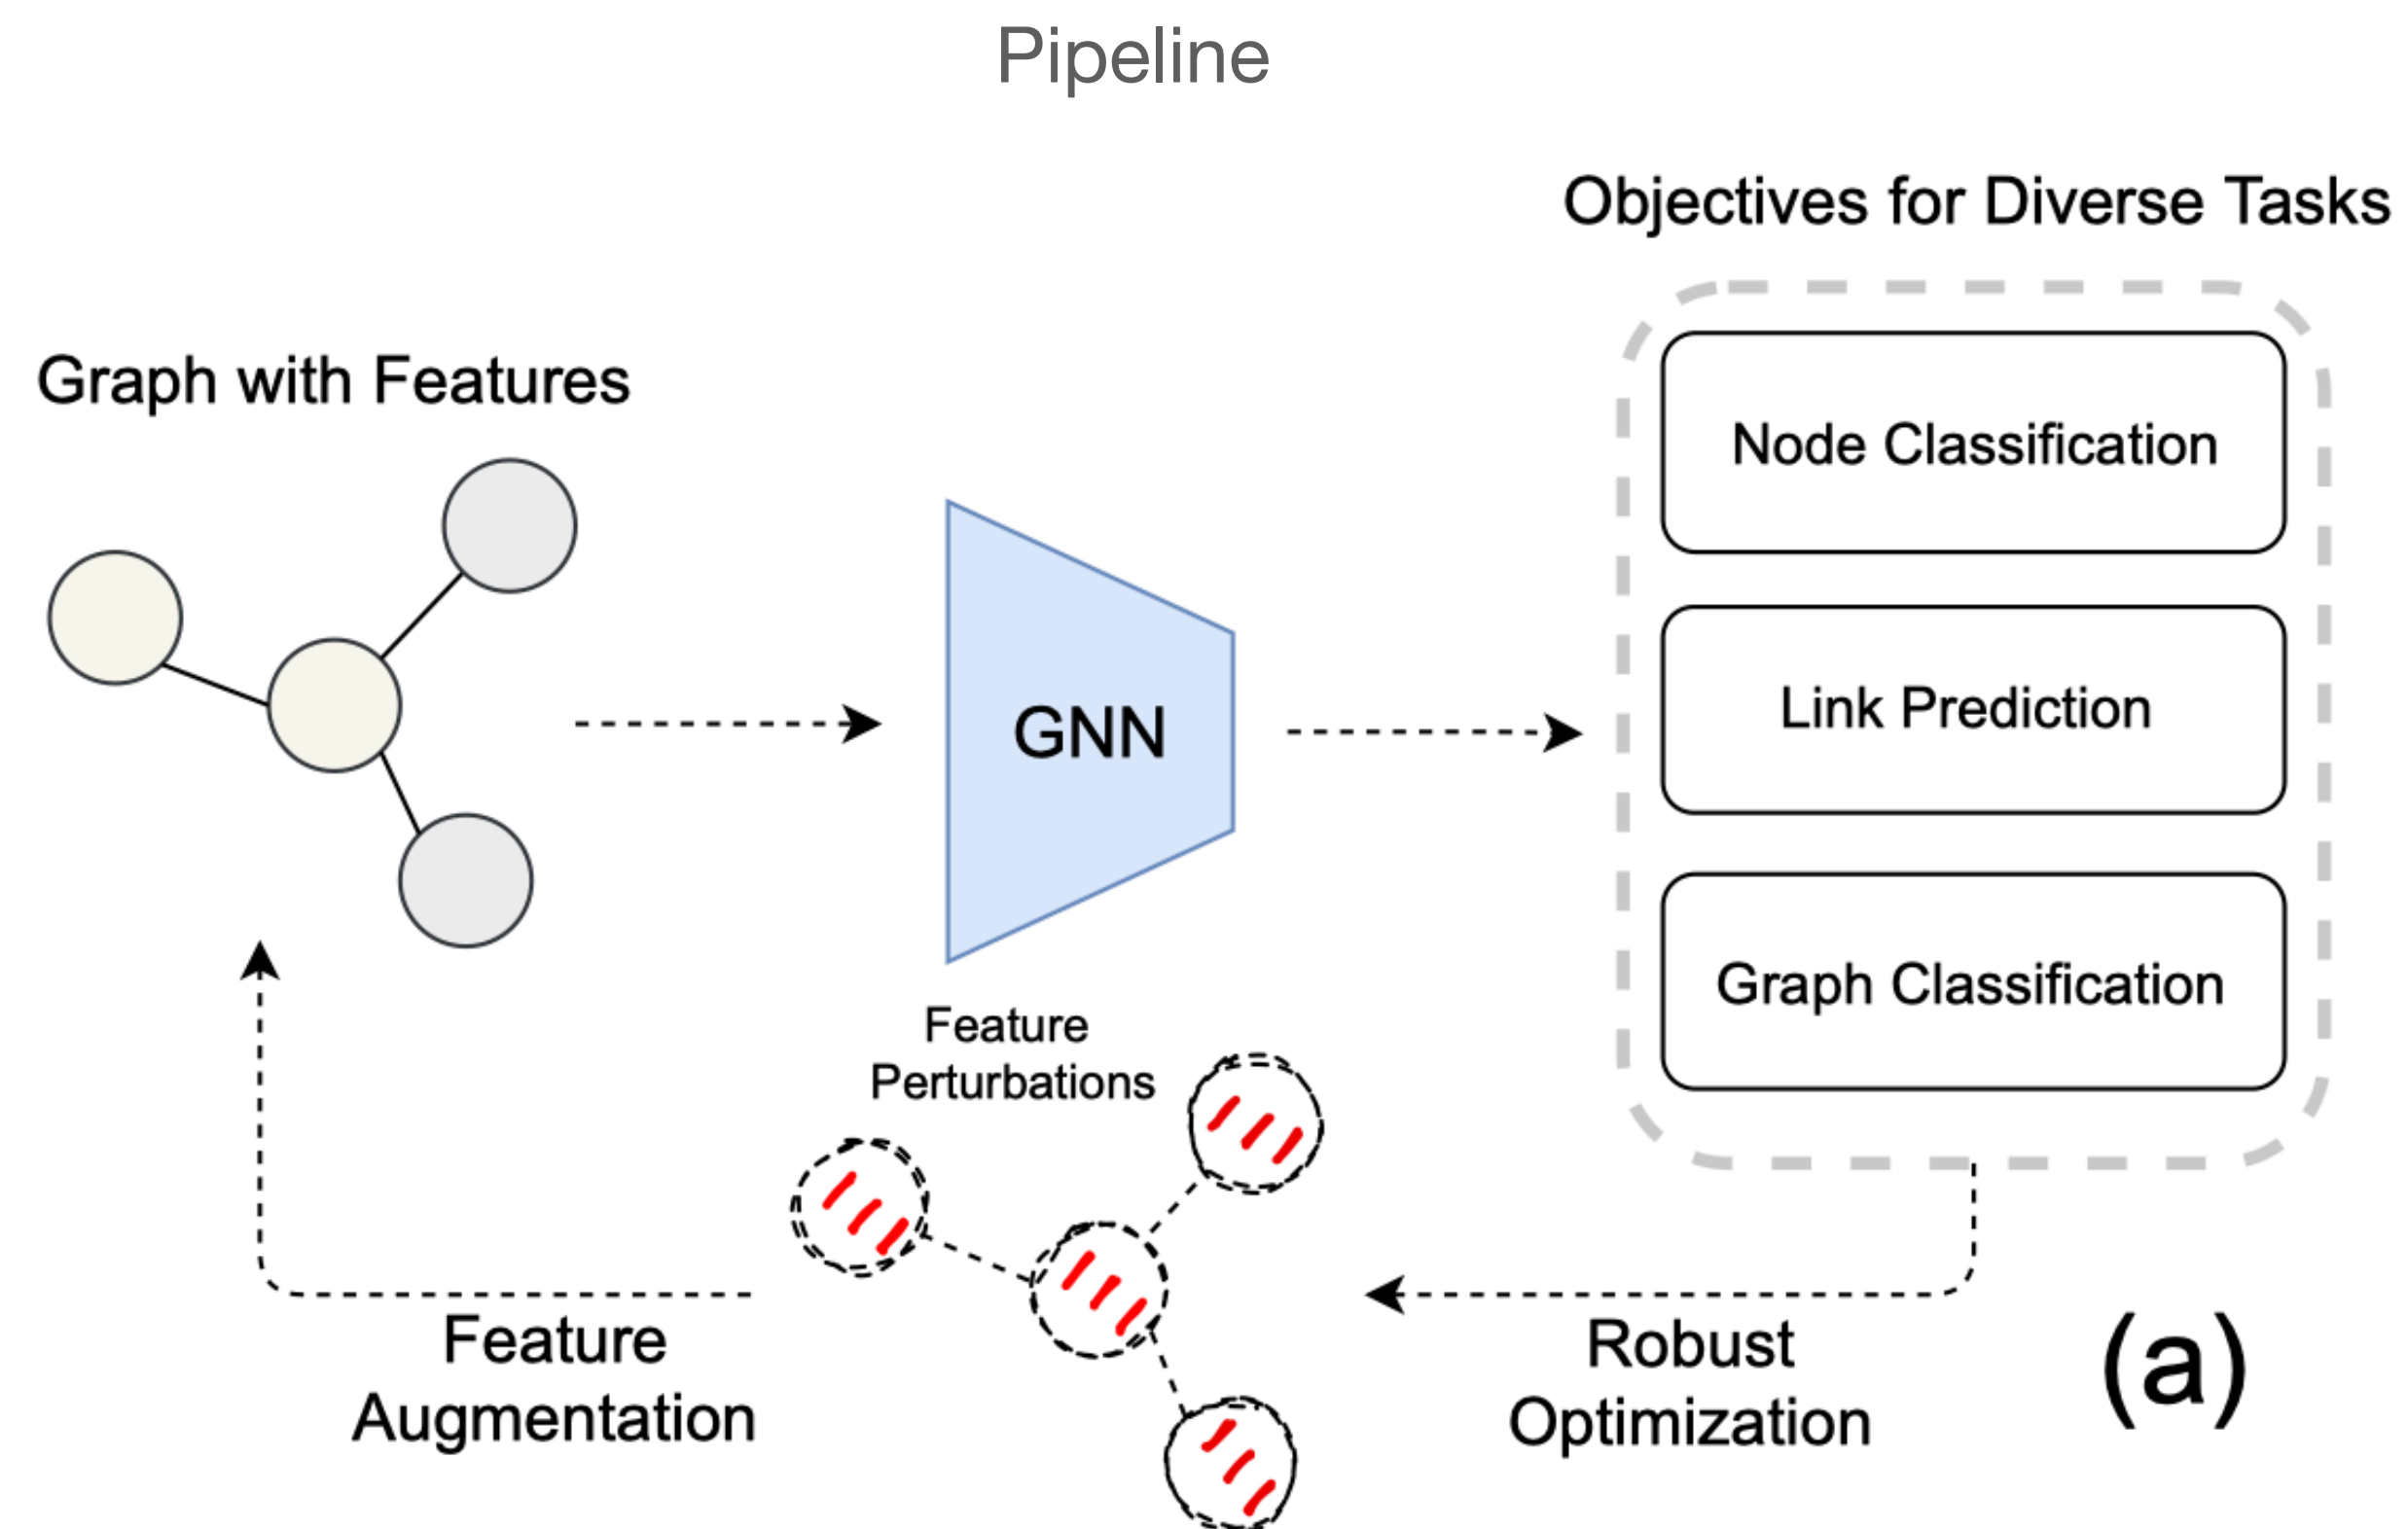
Method: FLAG

Min-Max Optimization

$$\min_{\theta} E_{(x,y) \sim \mathcal{D}} \left[\max_{\|\delta\|_p \leq \epsilon} L(f_{\theta}(x + \delta), y) \right],$$

Inner Gradient Ascent

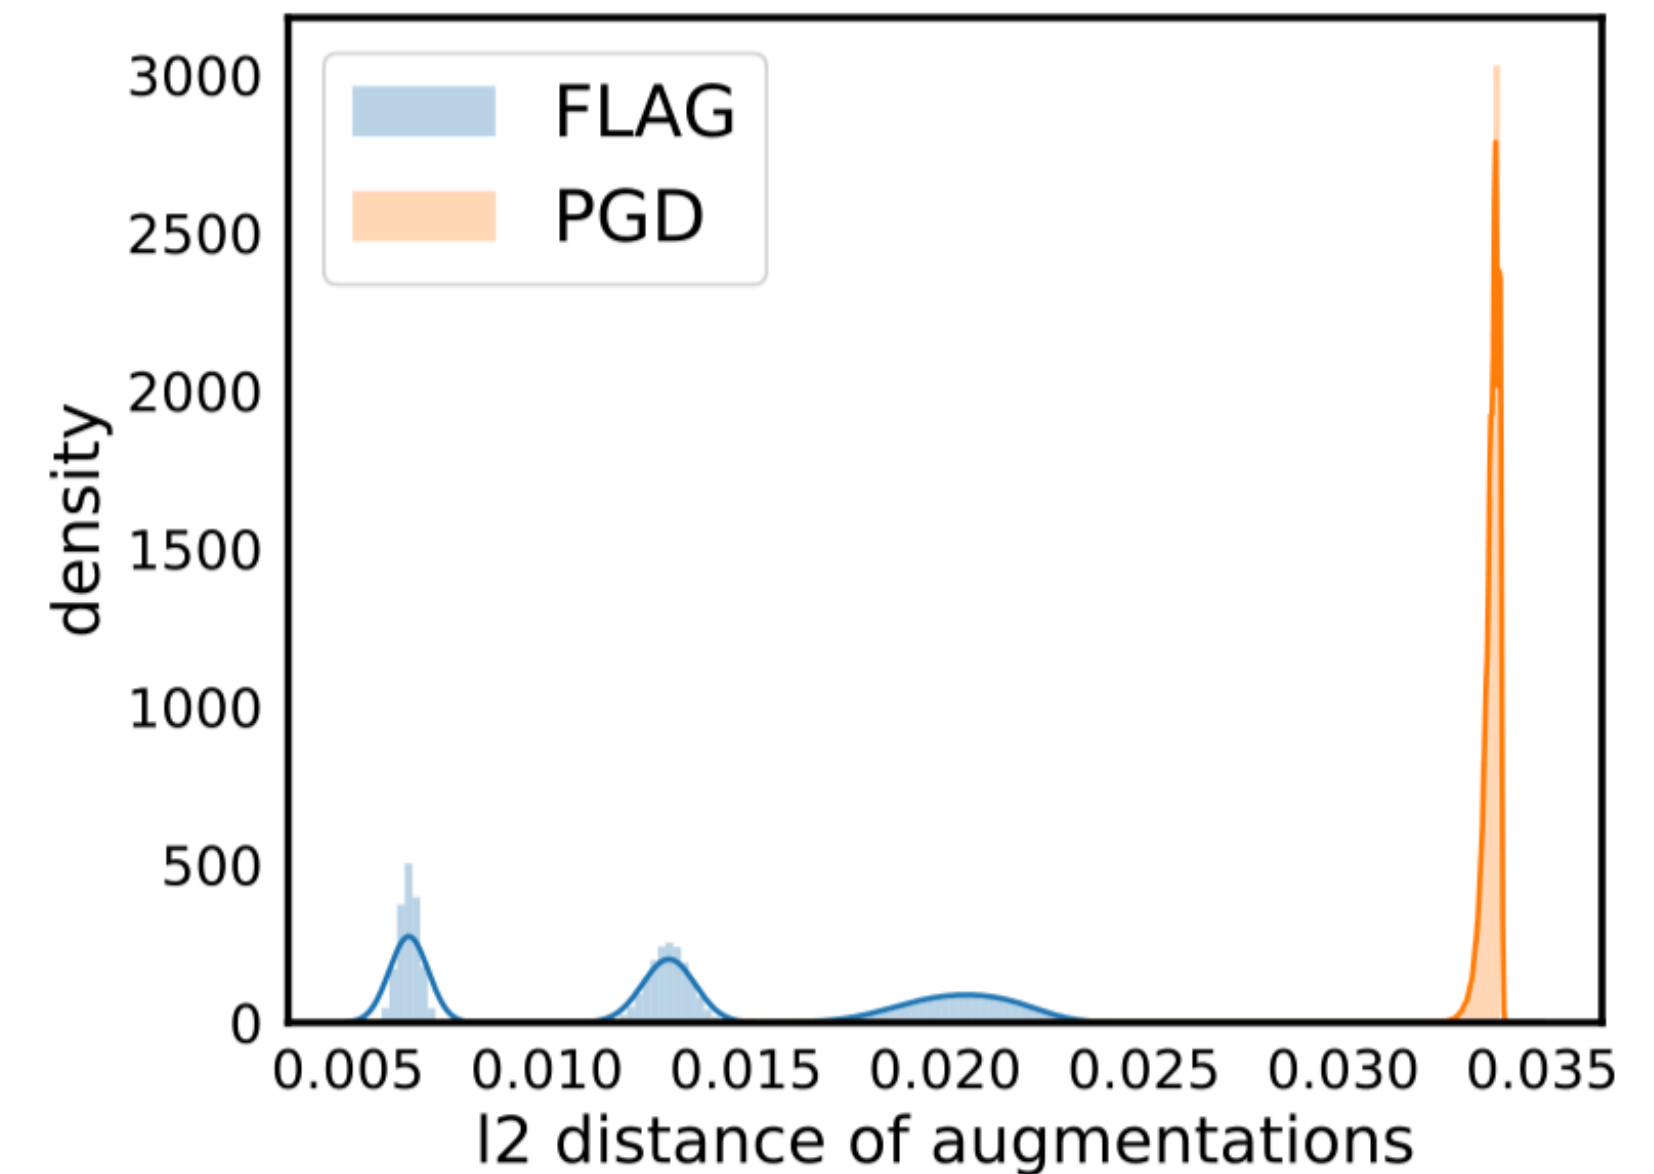
$$\delta_{t+1} = \Pi_{\|\delta\|_{\infty} \leq \epsilon} (\delta_t + \alpha \cdot \text{sign}(\nabla_{\delta} L(f_{\theta}(x + \delta_t), y)))$$



Method: FLAG


Principle: Multi-scale Augmentation

- “Free” Training
 - Leverage the batch replay technique but improve with gradient accumulation
- Weighted Perturbation
 - For node classification, augment labeled vs. unlabeled with diverse magnitudes (typically larger for unlabeled ones)



Method: FLAG

Principle: Multi-scale Augmentation



Method: FLAG

```
1  #M as ascent steps, alpha as ascent step size
2  #X denotes input node features, y denotes labels
3  def flag(gnn, X, y, optimizer, criterion, M, alpha) :
4      gnn.train()
5      optimizer.zero_grad()
6
7      pert = torch.FloatTensor(*X.shape).uniform_(-alpha, alpha)
8      pert.requires_grad_()
9      out = gnn(X+pert)
10     loss = criterion(out, y)/M
11
12     for _ in range(M-1):
13         loss.backward()
14         pert_data = pert.detach() + alpha*torch.sign(pert.grad.detach())
15         pert.data = pert_data.data
16         pert.grad[:] = 0
17         out = gnn(X+pert)
18         loss = criterion(out, y)/M
19
20     loss.backward()
21     optimizer.step()
```


Experiments

Node Classification

Table 1: Node property prediction test performance on ogbn-products, ogbn-proteins, and ogbn-arxiv datasets. Blank denotes no statistics on the leaderboard.

	ogbn-products	ogbn-proteins	ogbn-arxiv
Backbone	Test Acc	Test ROC-AUC	Test Acc
GCN	-	72.51±0.35	71.74±0.29
+FLAG	-	71.71±0.50	72.04±0.20
GraphSAGE	78.70±0.36	77.68 ±0.20	71.49±0.27
+FLAG	79.36±0.57	76.57±0.75	72.19±0.21
GAT	79.45±0.59	-	73.65±0.11
+FLAG	81.76±0.45	-	73.71±0.13
DeeperGCN	80.98±0.20	85.80±0.17	71.92±0.16
+FLAG	81.93±0.31	85.96±0.27	72.14±0.19

Table 2: Test performance on the heterogeneous OGB node property prediction dataset ogbn-mag.

	ogbn-mag
Backbone	Test Acc
R-GCN	46.78±0.67
+FLAG	47.37±0.48

Experiments

Link Prediction

Table 3: Link property prediction test performance on ogbl-ddi and ogbl-collab datasets.

Backbone	ogbl-ddi	ogbl-collab
	Hits@20	Hits@50
GCN	37.07 \pm 5.07	44.75 \pm 1.07
+FLAG	51.41 \pm 3.76	46.22 \pm 0.81
GraphSAGE	53.90 \pm 4.74	48.10 \pm 0.81
+FLAG	63.31 \pm 6.06	48.44 \pm 0.40

Experiments

Graph Classification

Table 4: Graph property test performance on ogbg-molhiv, ogbg-molpcba, ogbg-ppa, and ogbg-code datasets. \dagger denotes the existence of virtual nodes; blank denotes no statistics on the leaderboard.

Backbone	ogbg-molhiv Test ROC-AUC	ogbg-molpcba Test AP	ogbg-ppa Test Acc	ogbg-code Test F1
GCN	76.06 ± 0.97	20.20 ± 0.24	68.39 ± 0.34	31.63 ± 0.18
+FLAG	76.83 ± 1.02	21.16 ± 0.17	68.38 ± 0.47	32.09 ± 0.19
GCN-Virtual	75.99 ± 1.19	24.24 ± 0.34	68.57 ± 0.61	32.63 ± 0.13
+FLAG	75.45 ± 1.58	24.83 ± 0.37	69.44 ± 0.52	33.16 ± 0.25
GIN	75.58 ± 1.40	22.66 ± 0.28	68.92 ± 1.00	31.63 ± 0.20
+FLAG	76.54 ± 1.14	23.95 ± 0.40	69.05 ± 0.92	32.41 ± 0.40
GIN-Virtual	77.07 ± 1.49	27.03 ± 0.23	70.37 ± 1.07	32.04 ± 0.18
+FLAG	77.48 ± 0.96	28.34 ± 0.38	72.45 ± 1.14	32.96 ± 0.36
DeeperGCN	78.58 ± 1.17	$27.81^{\dagger} \pm 0.38$	77.12 ± 0.71	-
+FLAG	79.42 ± 1.20	$28.42^{\dagger} \pm 0.43$	77.52 ± 0.69	-

Ablation Studies

Table 6: Test accuracy of GAT on ogbn-products trained with different adversarial augmentations. FLAG (fast) means the training epoch number is decreased to make our method trained as fast as the baseline.

Backbone	Test Acc
GAT	79.45 \pm 0.59
GAT+PGD	80.96 \pm 0.41
GAT+“Free”	79.42 \pm 0.84
GAT+FLAG	81.76\pm0.45
GAT+FLAG (fast)	80.64 \pm 0.74

Table 5: Test Accuracy on the ogbn-arxiv dataset with different BN methods.

Method	GCN	GraphSAGE
w/o BN	71.09 \pm 0.22	69.58 \pm 0.76
w/ BN	71.74 \pm 0.29	71.49 \pm 0.27
w/ BN +FLAG	72.04 \pm 0.20	72.19 \pm 0.21
w/ Dual BN +FLAG	72.11\pm0.23	72.21\pm0.20

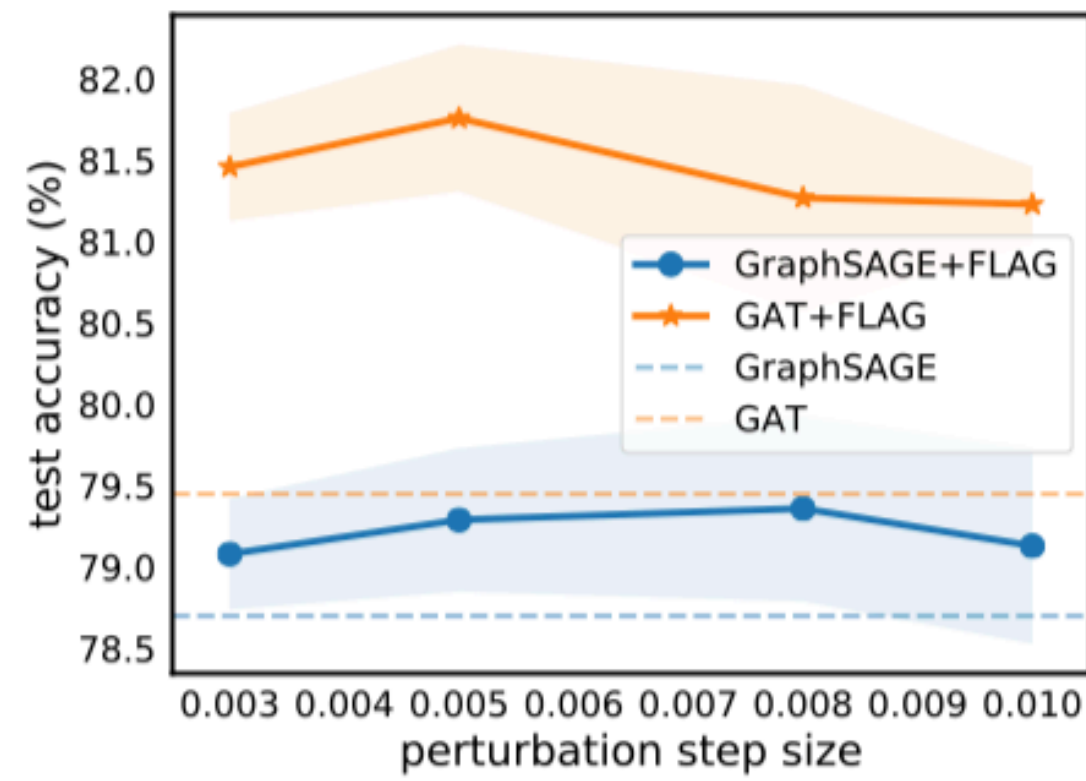
Table 7: Test Accuracy on the ogbn-products dataset.

Backbone	Test Acc
GAT w/o dropout	75.67 \pm 0.27
GAT w/ dropout	79.45 \pm 0.59
GAT w/ dropout +FLAG	81.76\pm0.45

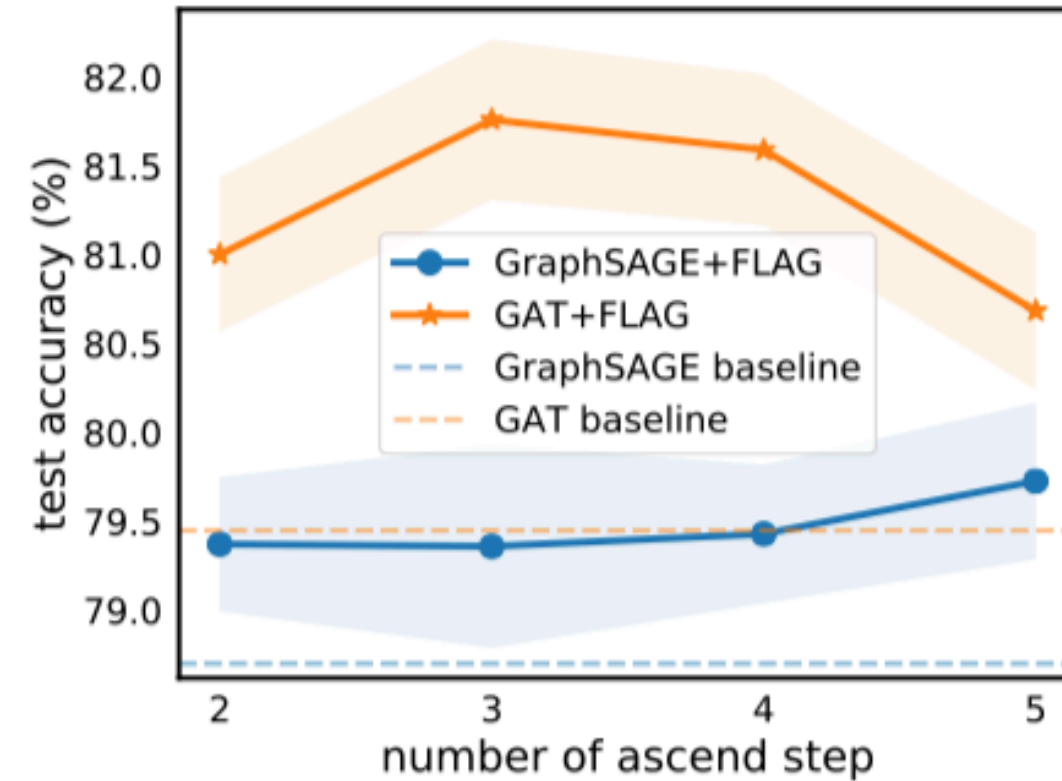
Table 8: Test accuracy on ogbn-products with GraphSAGE trained with diverse mini-batch algorithms.

Backbone	ogbn-products
	Test Acc
GraphSAGE w/ NS	78.70 \pm 0.36
+FLAG	79.36\pm0.57
GraphSAGE w/ Cluster	78.97\pm0.33
+FLAG	78.60 \pm 0.27
GraphSAGE w/ SAINT	79.08 \pm 0.24
+FLAG	79.60\pm0.19

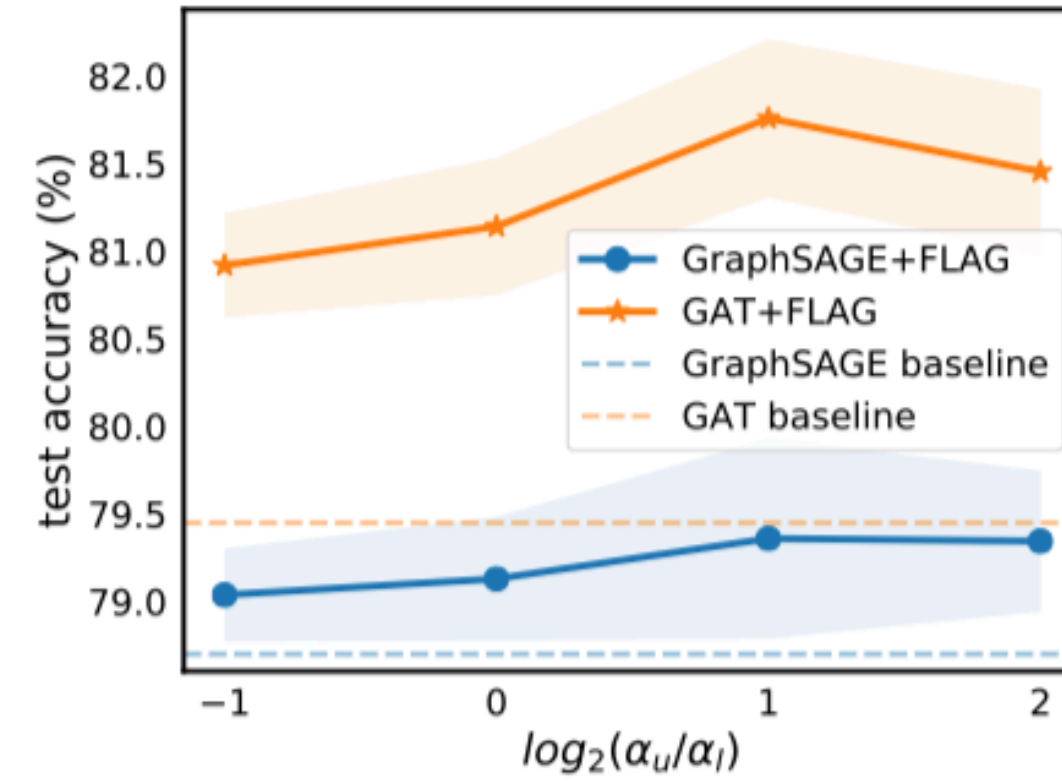
Ablation Studies




(a) step size



(b) ascent steps



(c) weighted perturbation



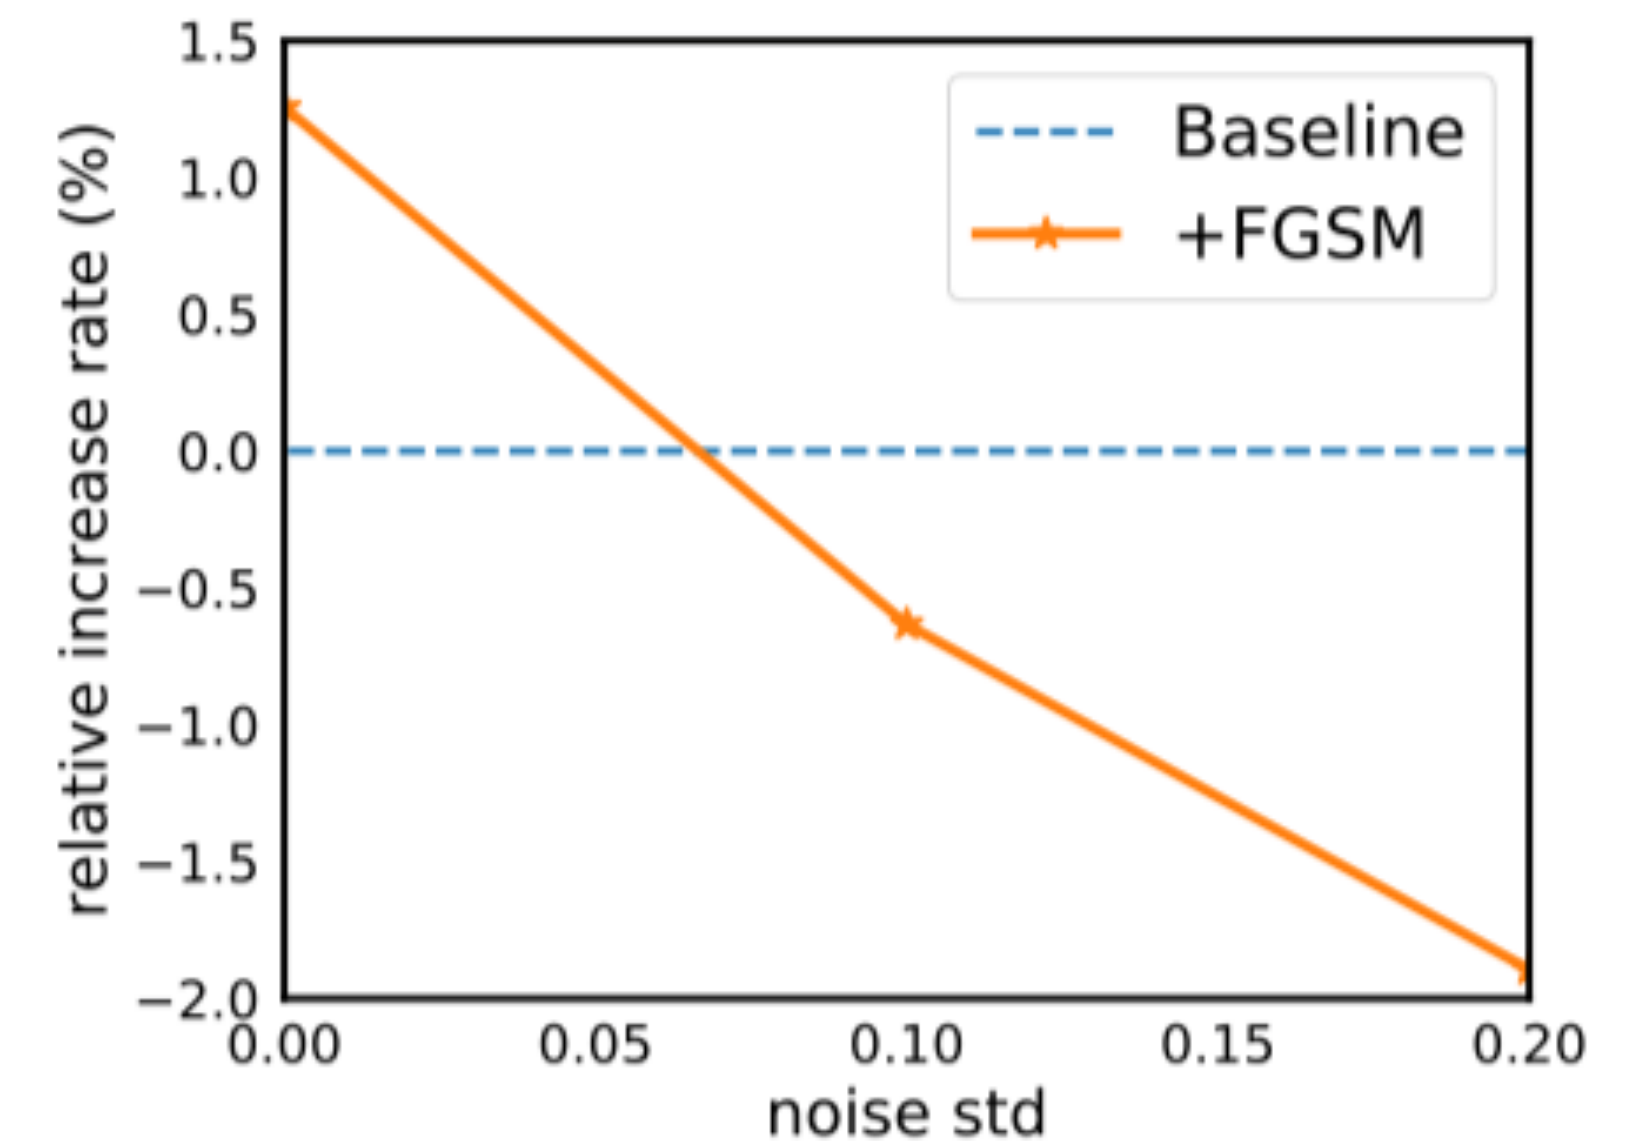
(a) model depth

Figure 3: Results of GraphSAGE and GAT on the ogbn-products dataset.

Analysis

- Guess: Discrete features are the ones that should be adversarially augmented
- Observations
 1. FLAG boosts MLPs on graph datasets
 - A. $61.06 \pm 0.08\%$ to $62.41 \pm 0.16\%$ on ogbn-product
 - B. $55.50 \pm 0.23\%$ to $56.02 \pm 0.19\%$ on ogbn-arxiv
 2. Adversarial Training boosts performance of CNNs on MNIST [Tsipras et al. 2018]
 3. Gaussian noises destroy the boost

Cora toy example



(b) noise std

Limitations

- FLAG is empirically oriented, which lacks theoretical motivation
- FLAG introduces time overhead
- FLAG may not work when there are no initial node features

Thanks for listening!

- Paper: <https://arxiv.org/abs/2010.09891>
- Code: <https://github.com/devnkong/FLAG>