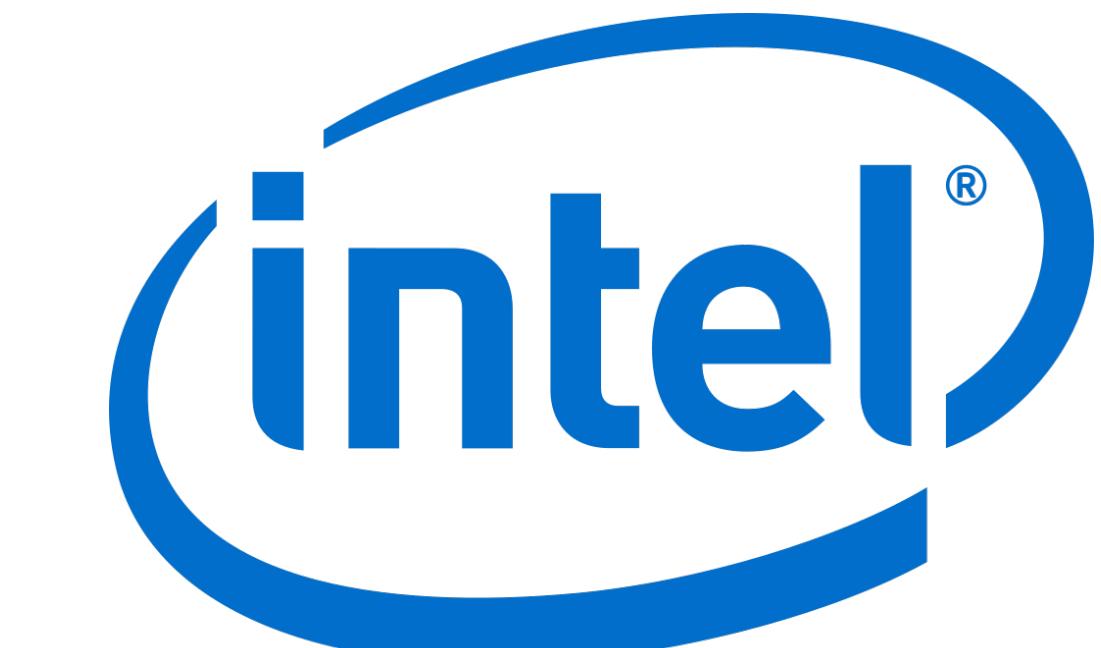
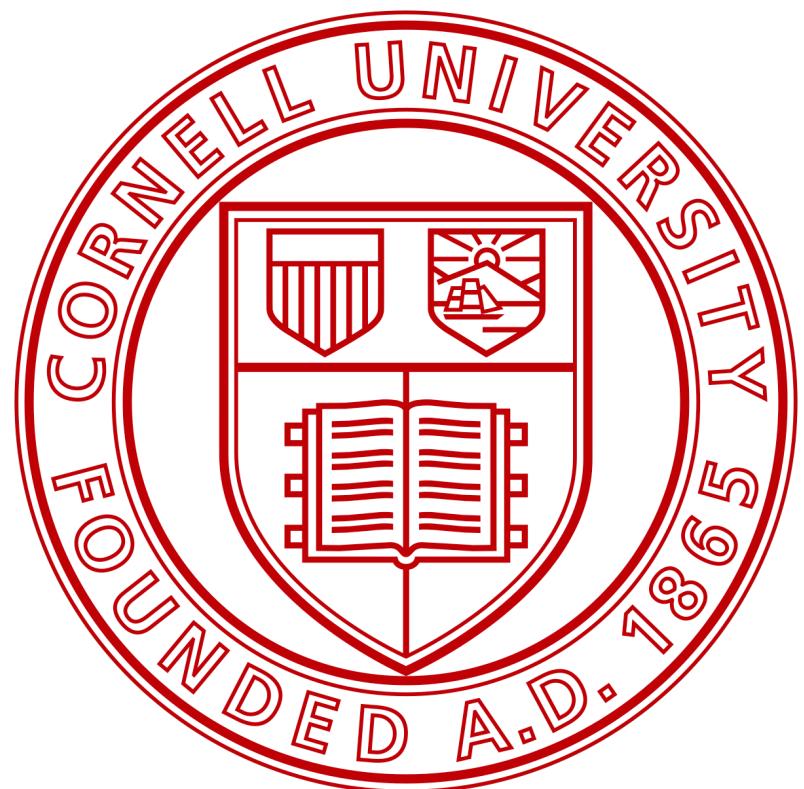


Geometry Processing with Neural Fields

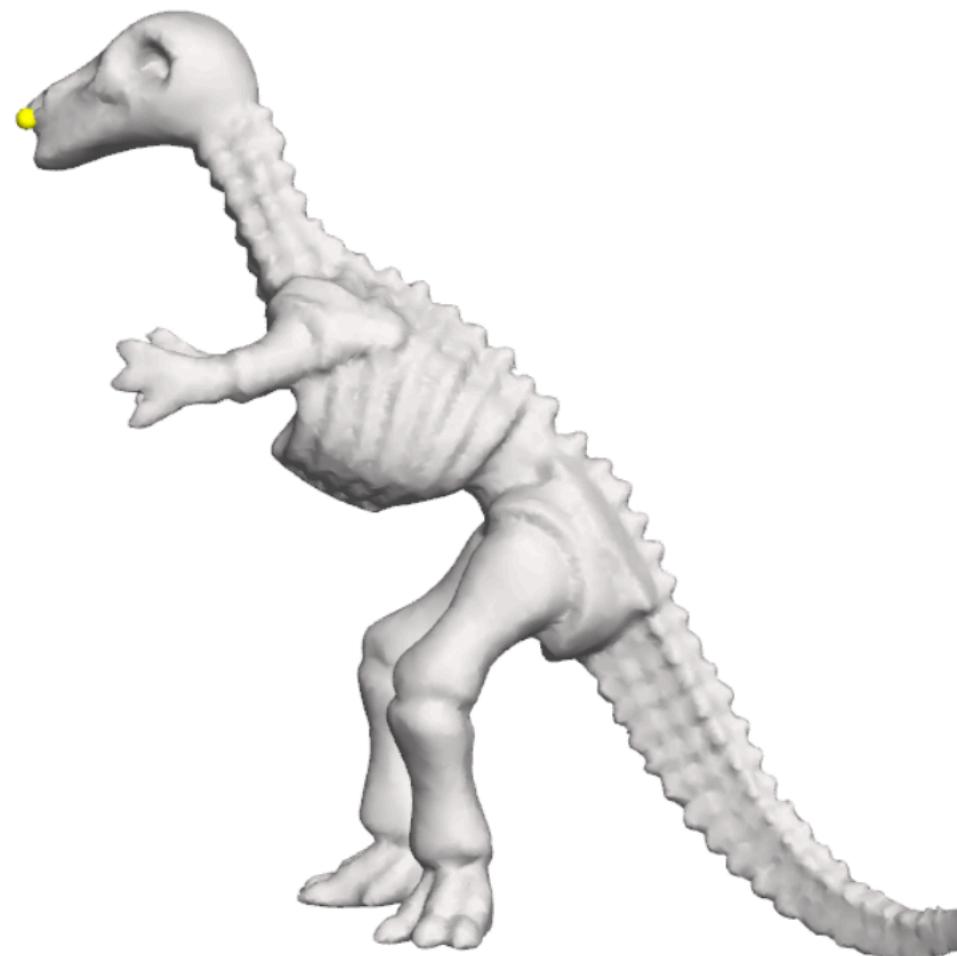
Guandao Yang, Serge Belongie, Bharath Hariharan, Vladlen Koltun



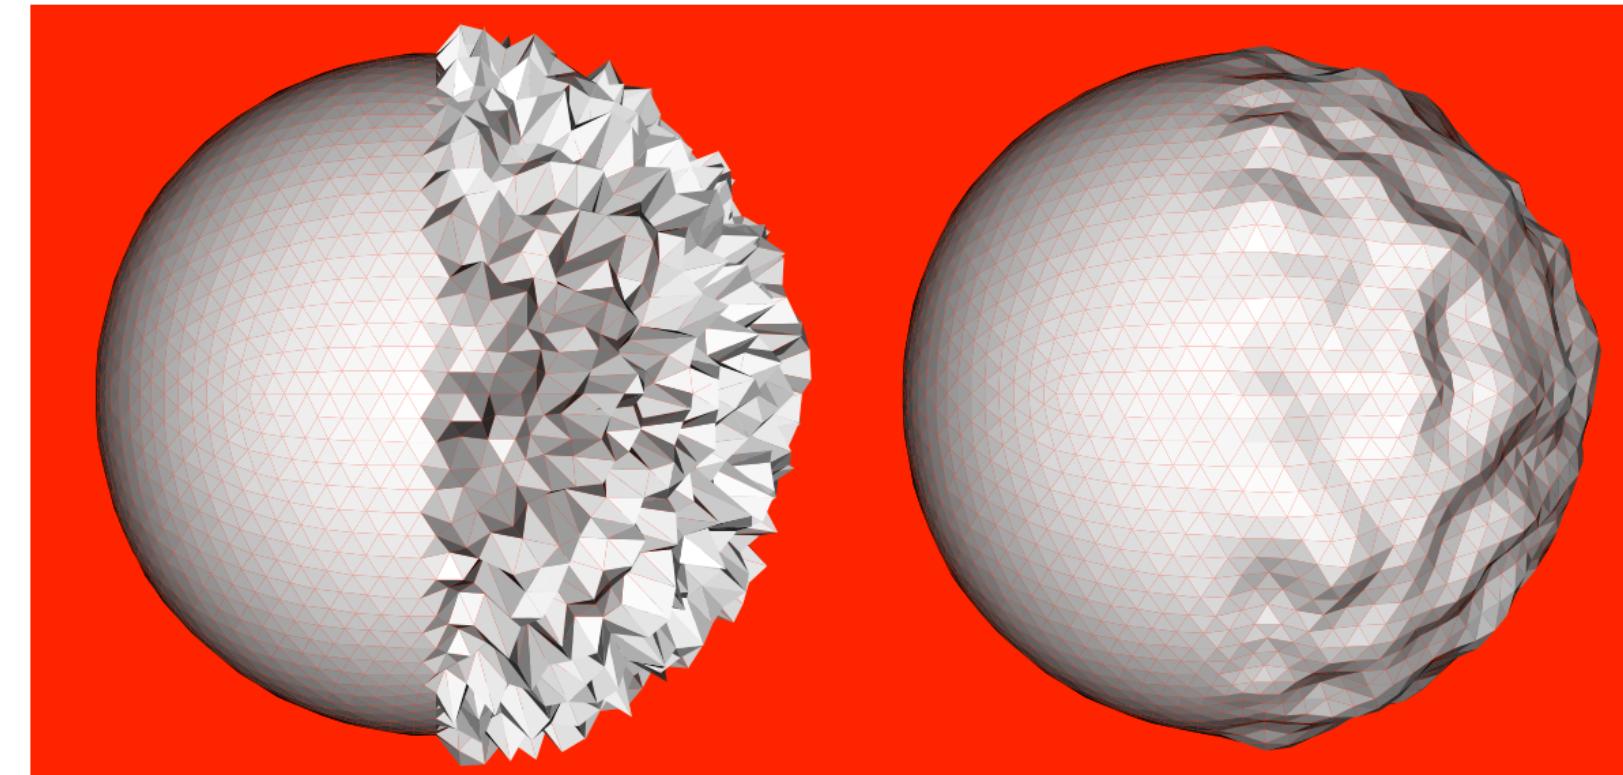
Geometry Processing

- manipulation of geometry data

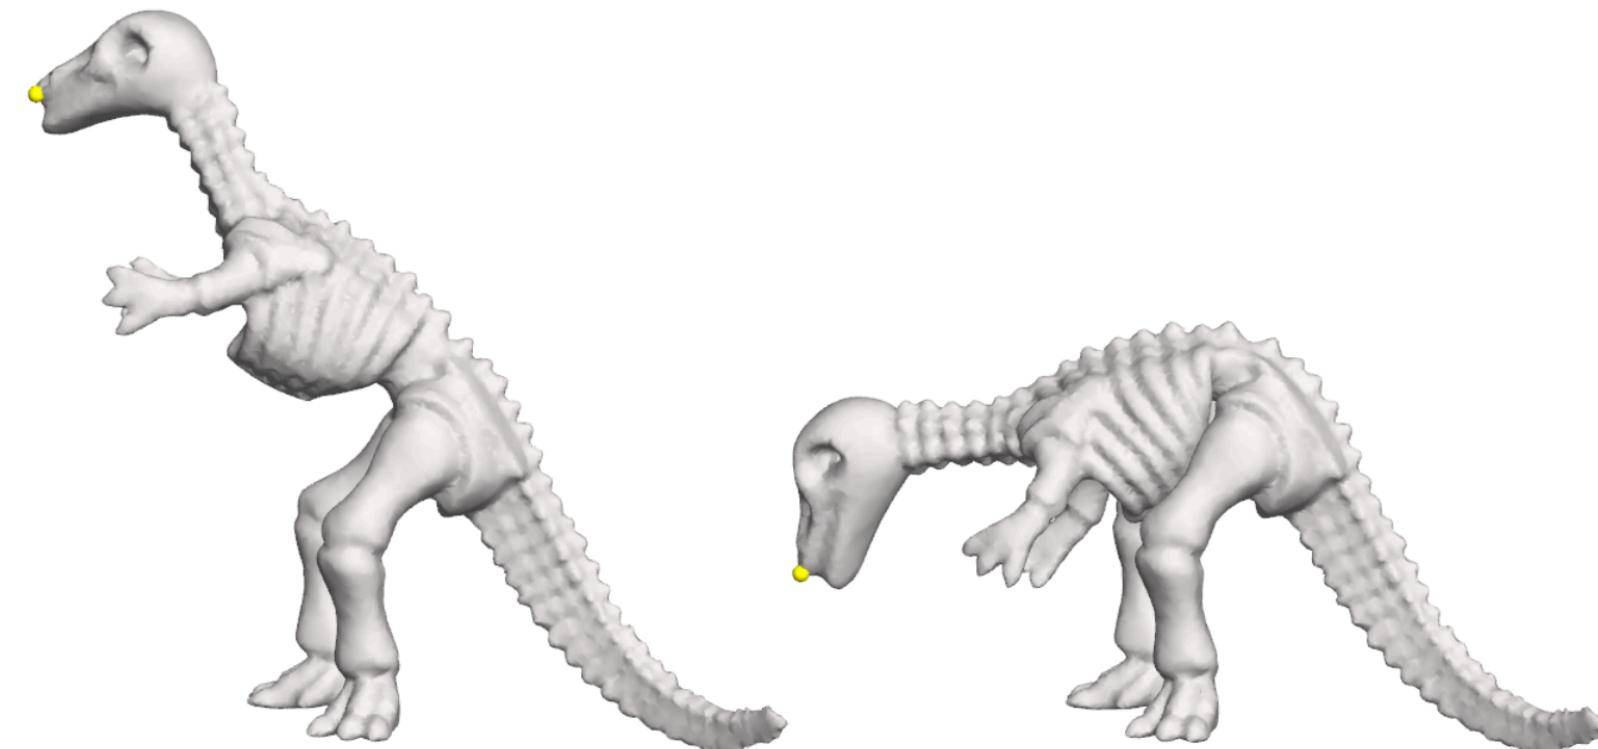
Geometry Data



Geometry Processing Tasks



(Gabriel Taubin, 1995)

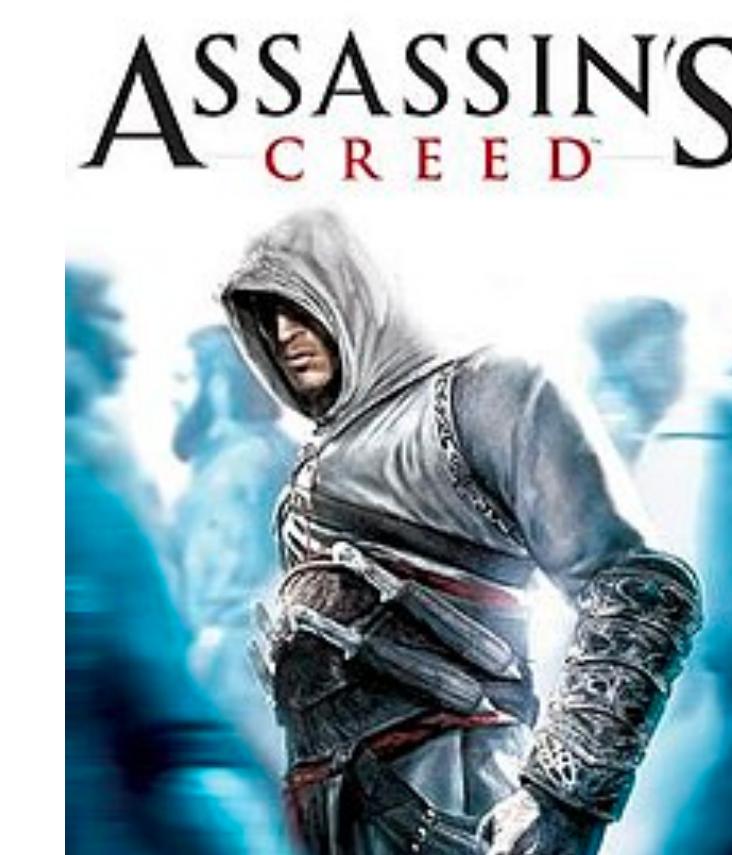


(Sorkine and Alexa, 2007)

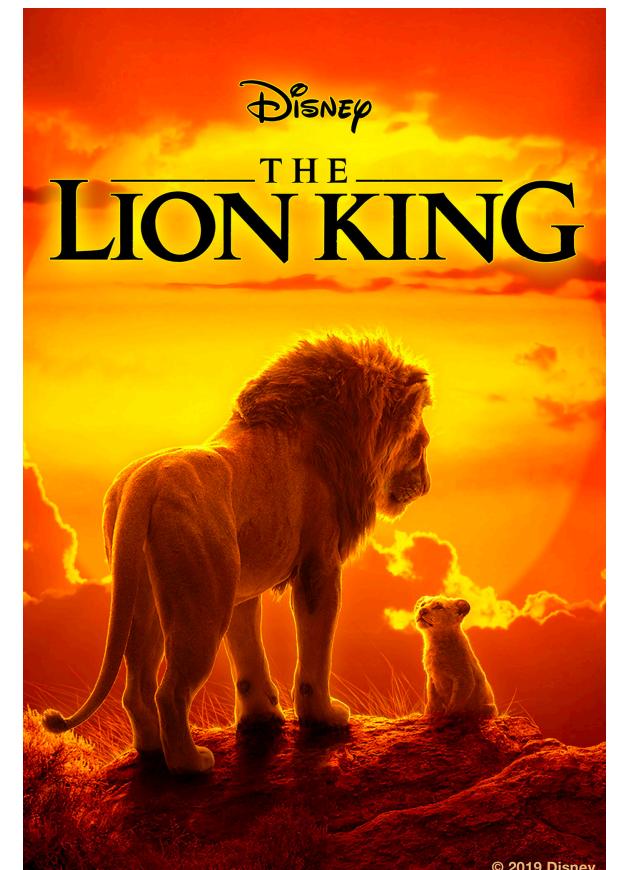
Applications



Computer Aided Design

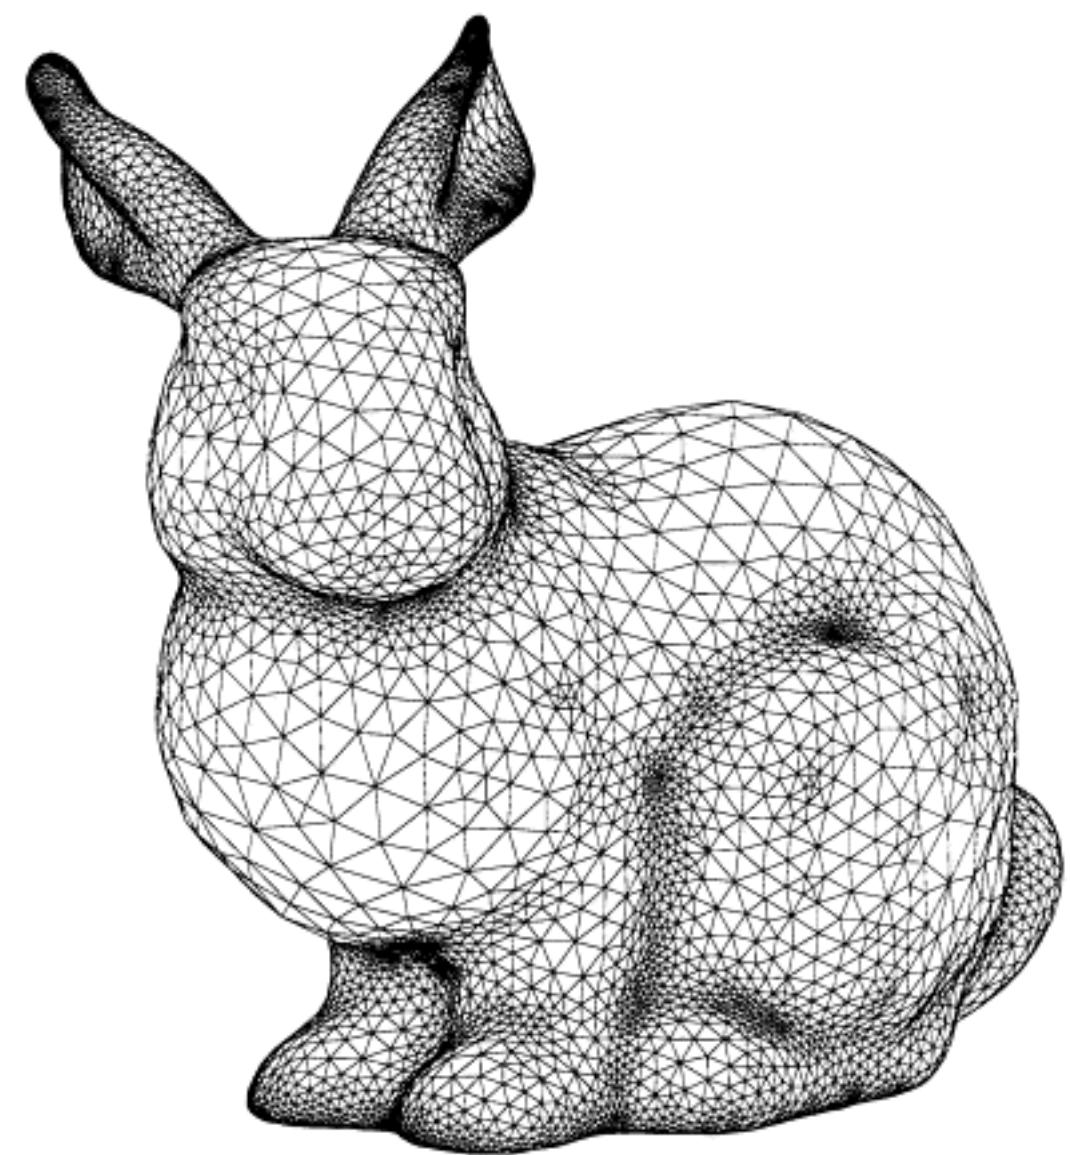


Gaming and Movies



© 2019 Disney

Representing Geometry Data



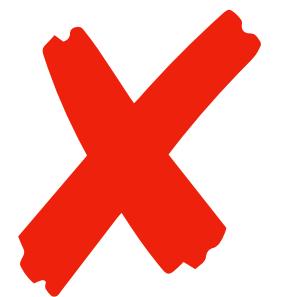
$$V = \{(x_i, y_i, z_i)\}_{i=1}^n$$
$$F = \{(u, v, w) \mid 1 \leq u, v, w, \leq n\}$$



Easy to understand



Fast access to local neighbor

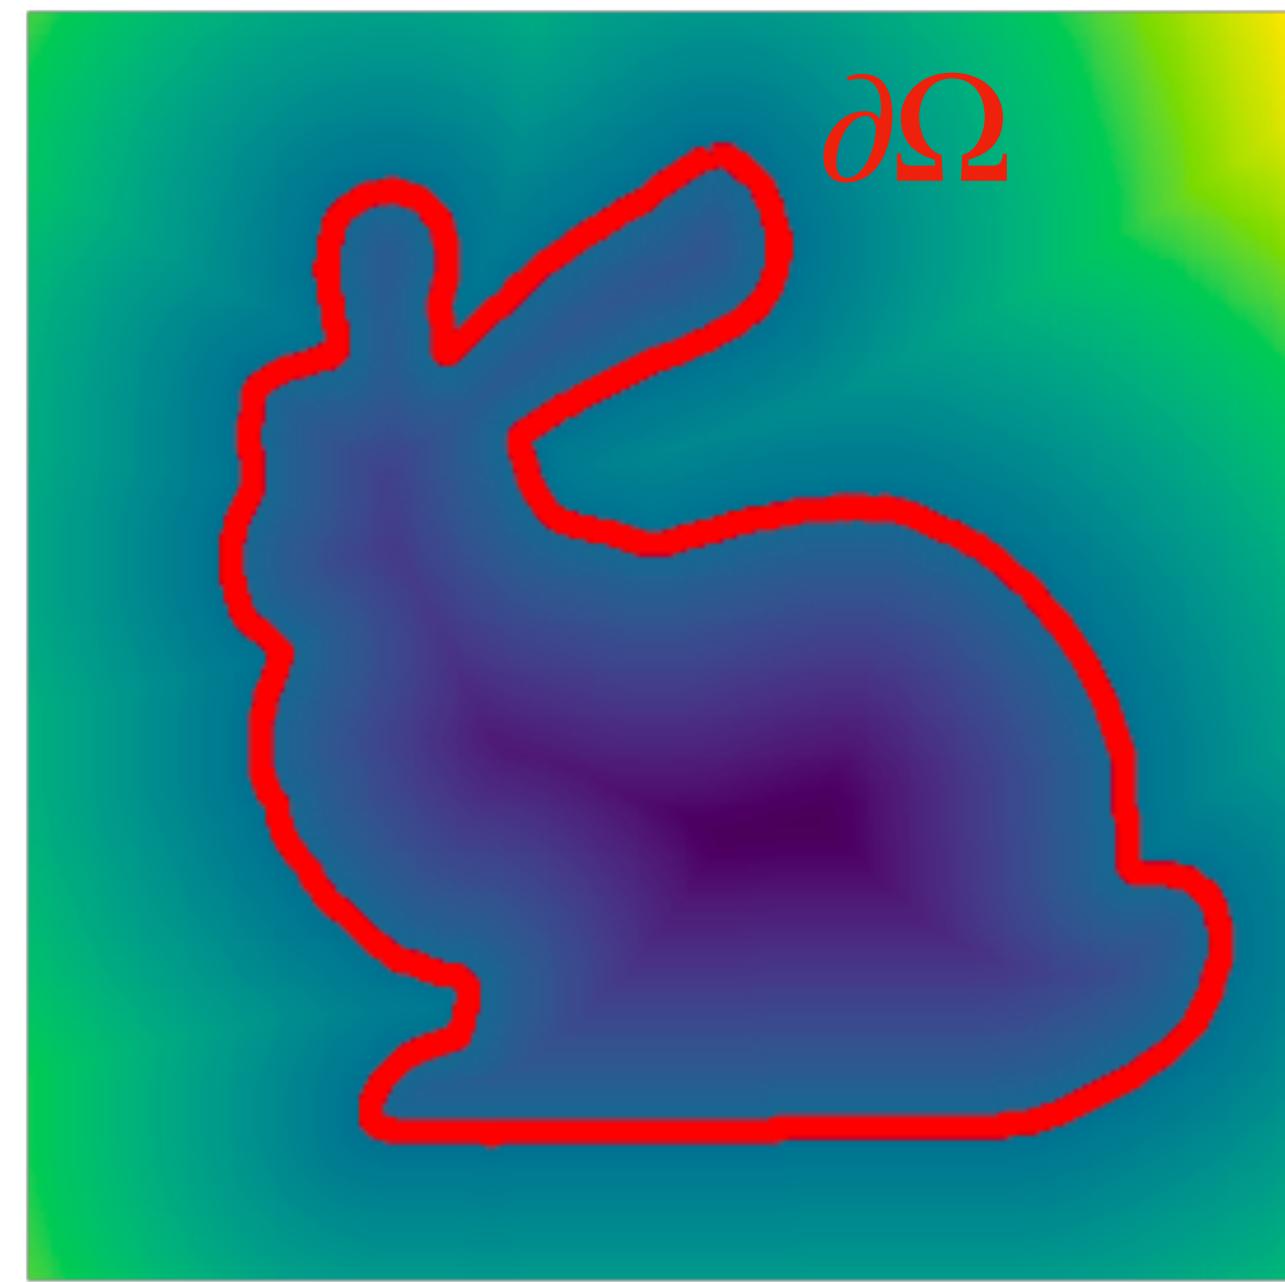


Topological changes



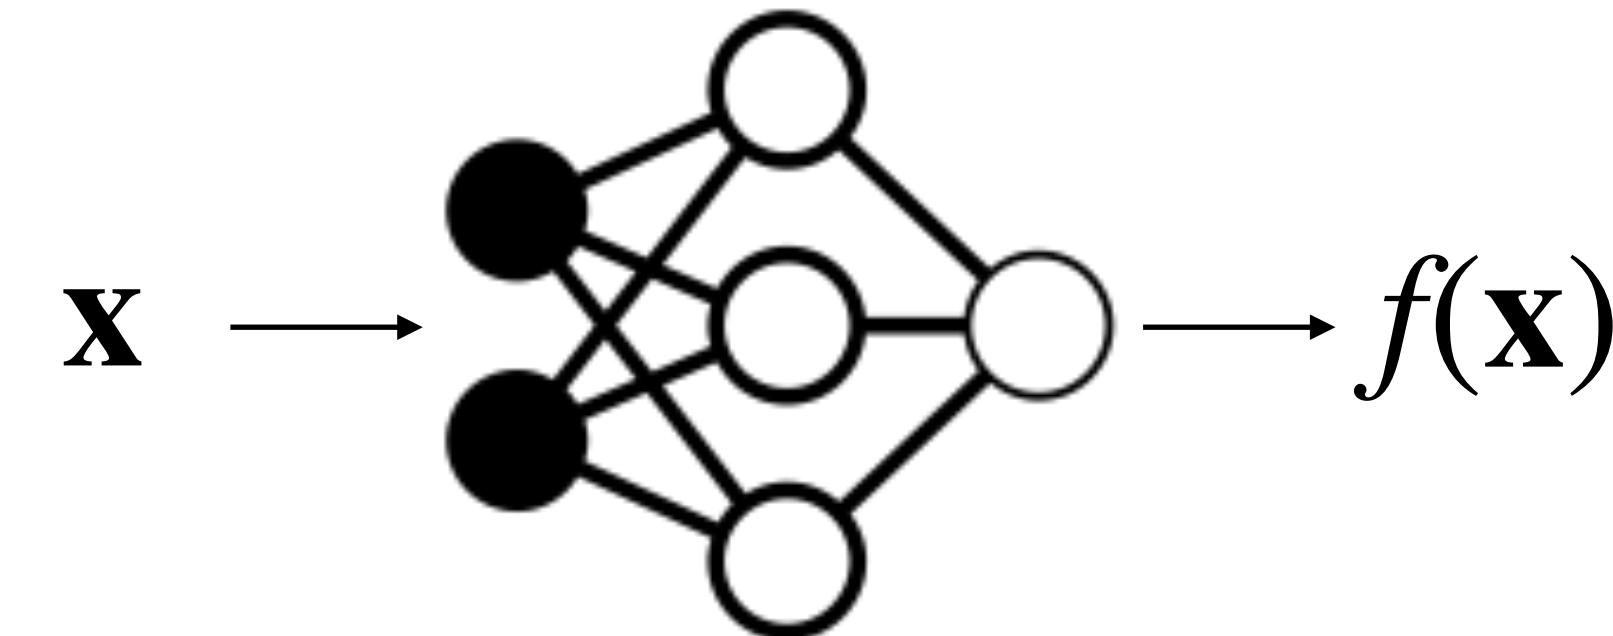
Need to maintain good surface discretization

Representing Geometry Data



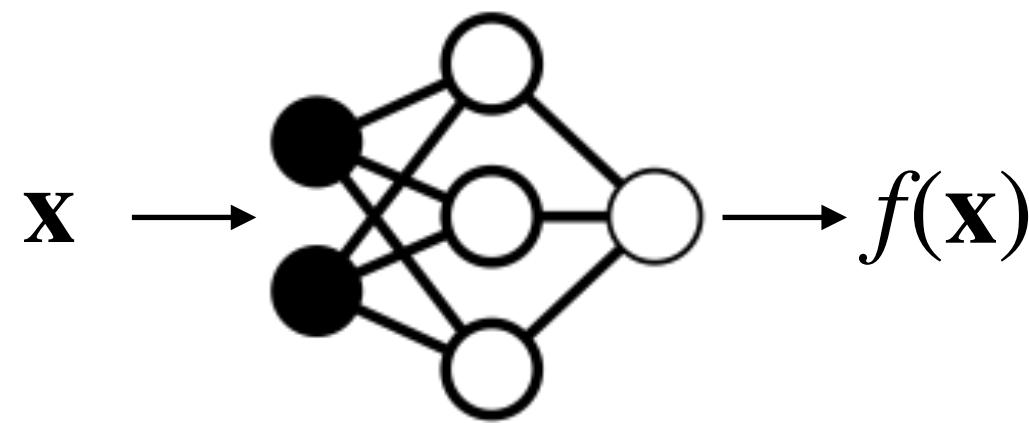
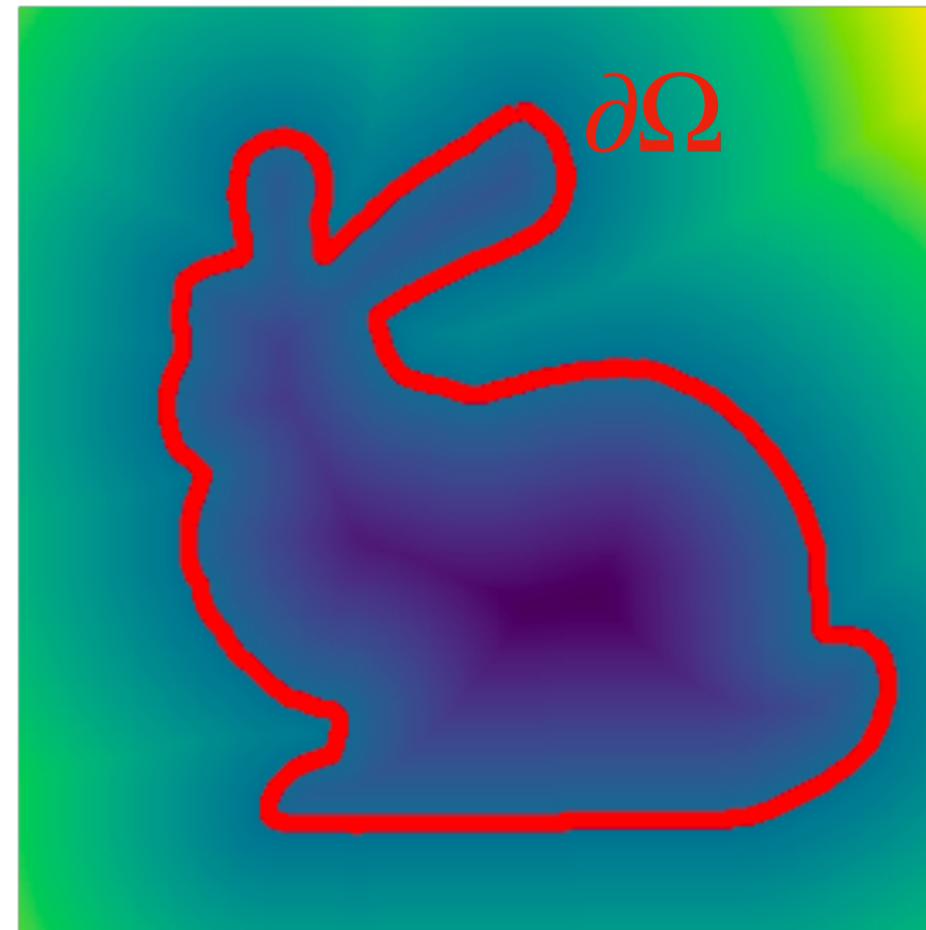
$$f: \mathbb{R}^3 \rightarrow \mathbb{R}$$

$$\partial\Omega = \{\mathbf{x} \mid f(\mathbf{x}) = c\}$$

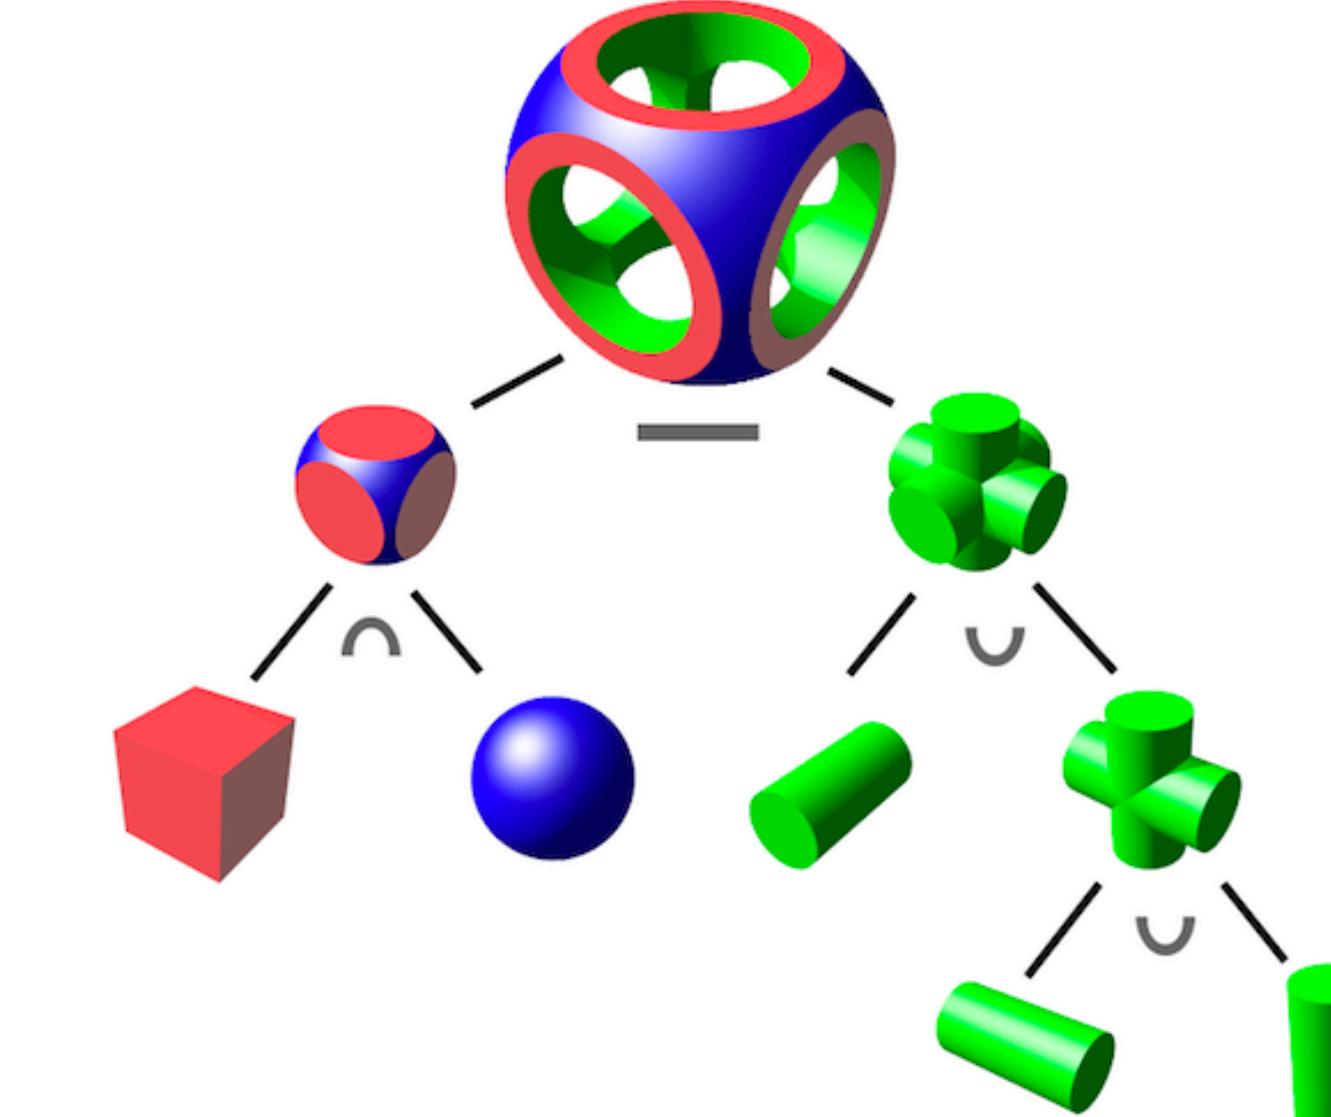


Neural Fields

Representing Geometry Data



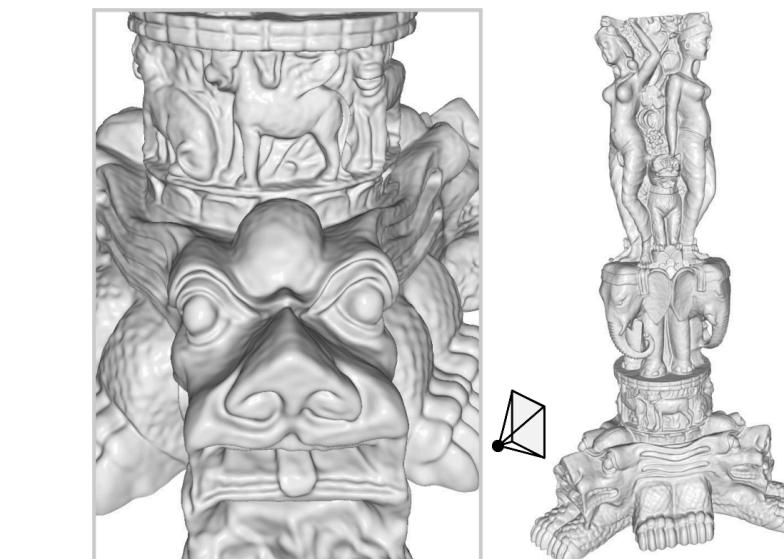
Neural Fields



✓ Topological changes



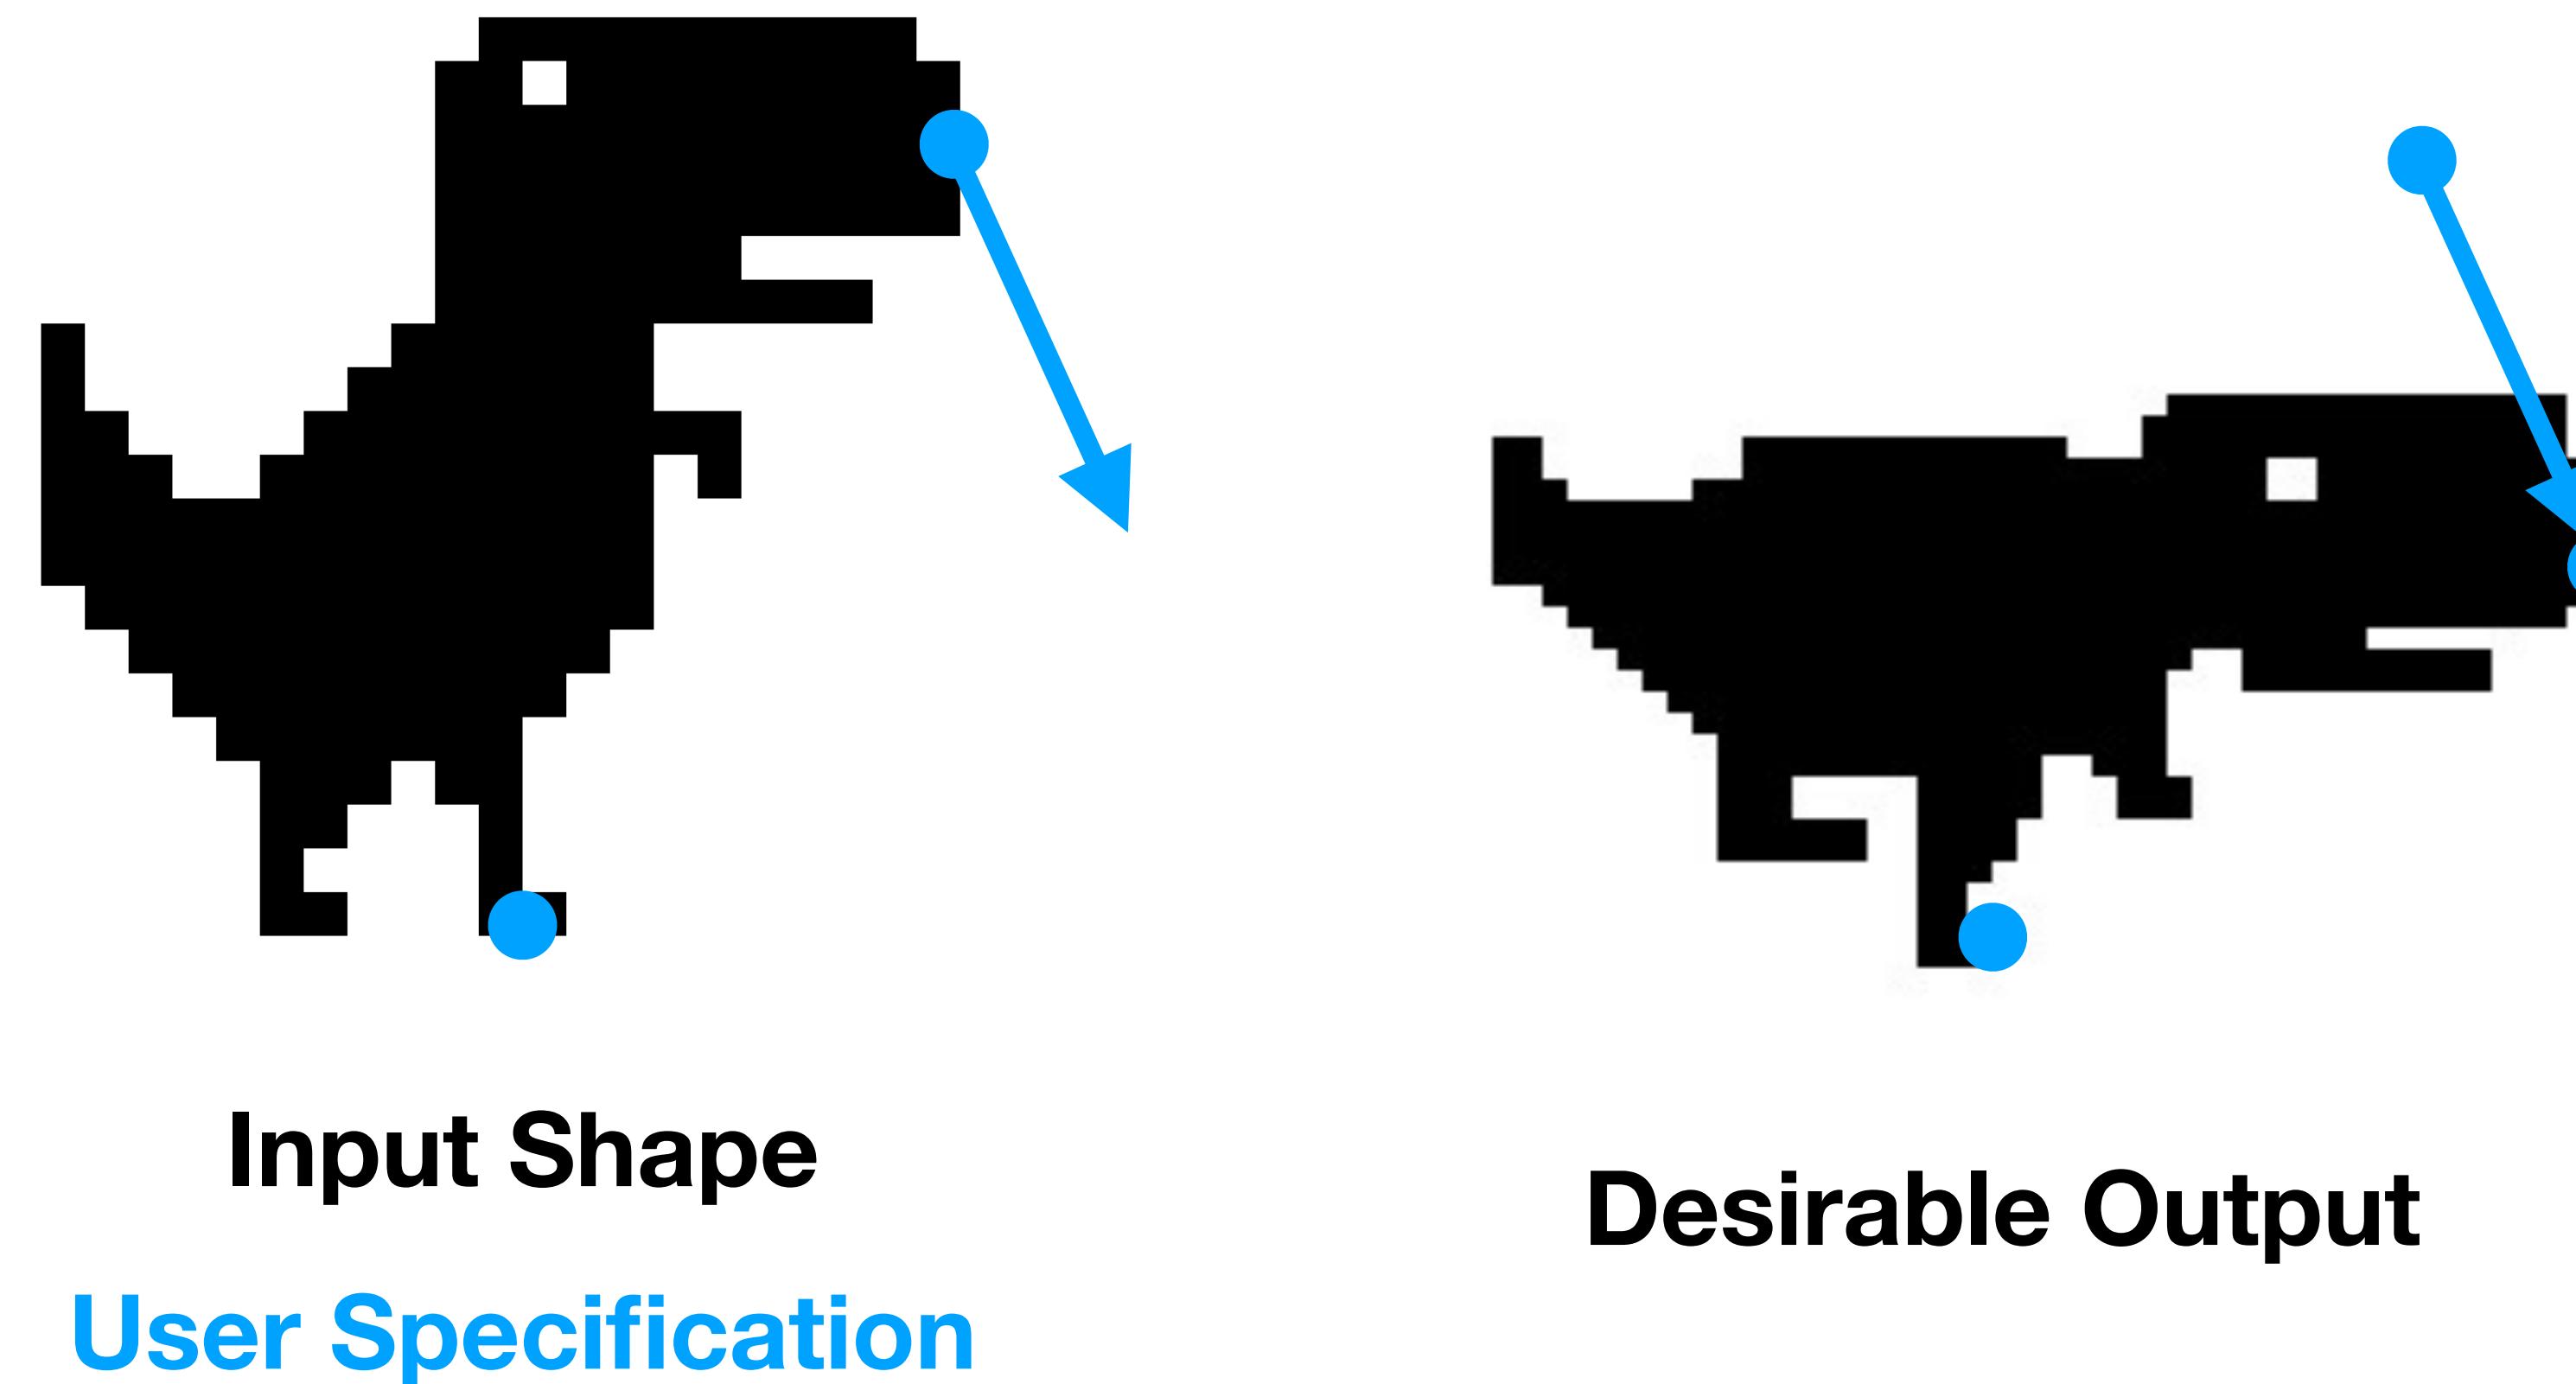
✓ Compact



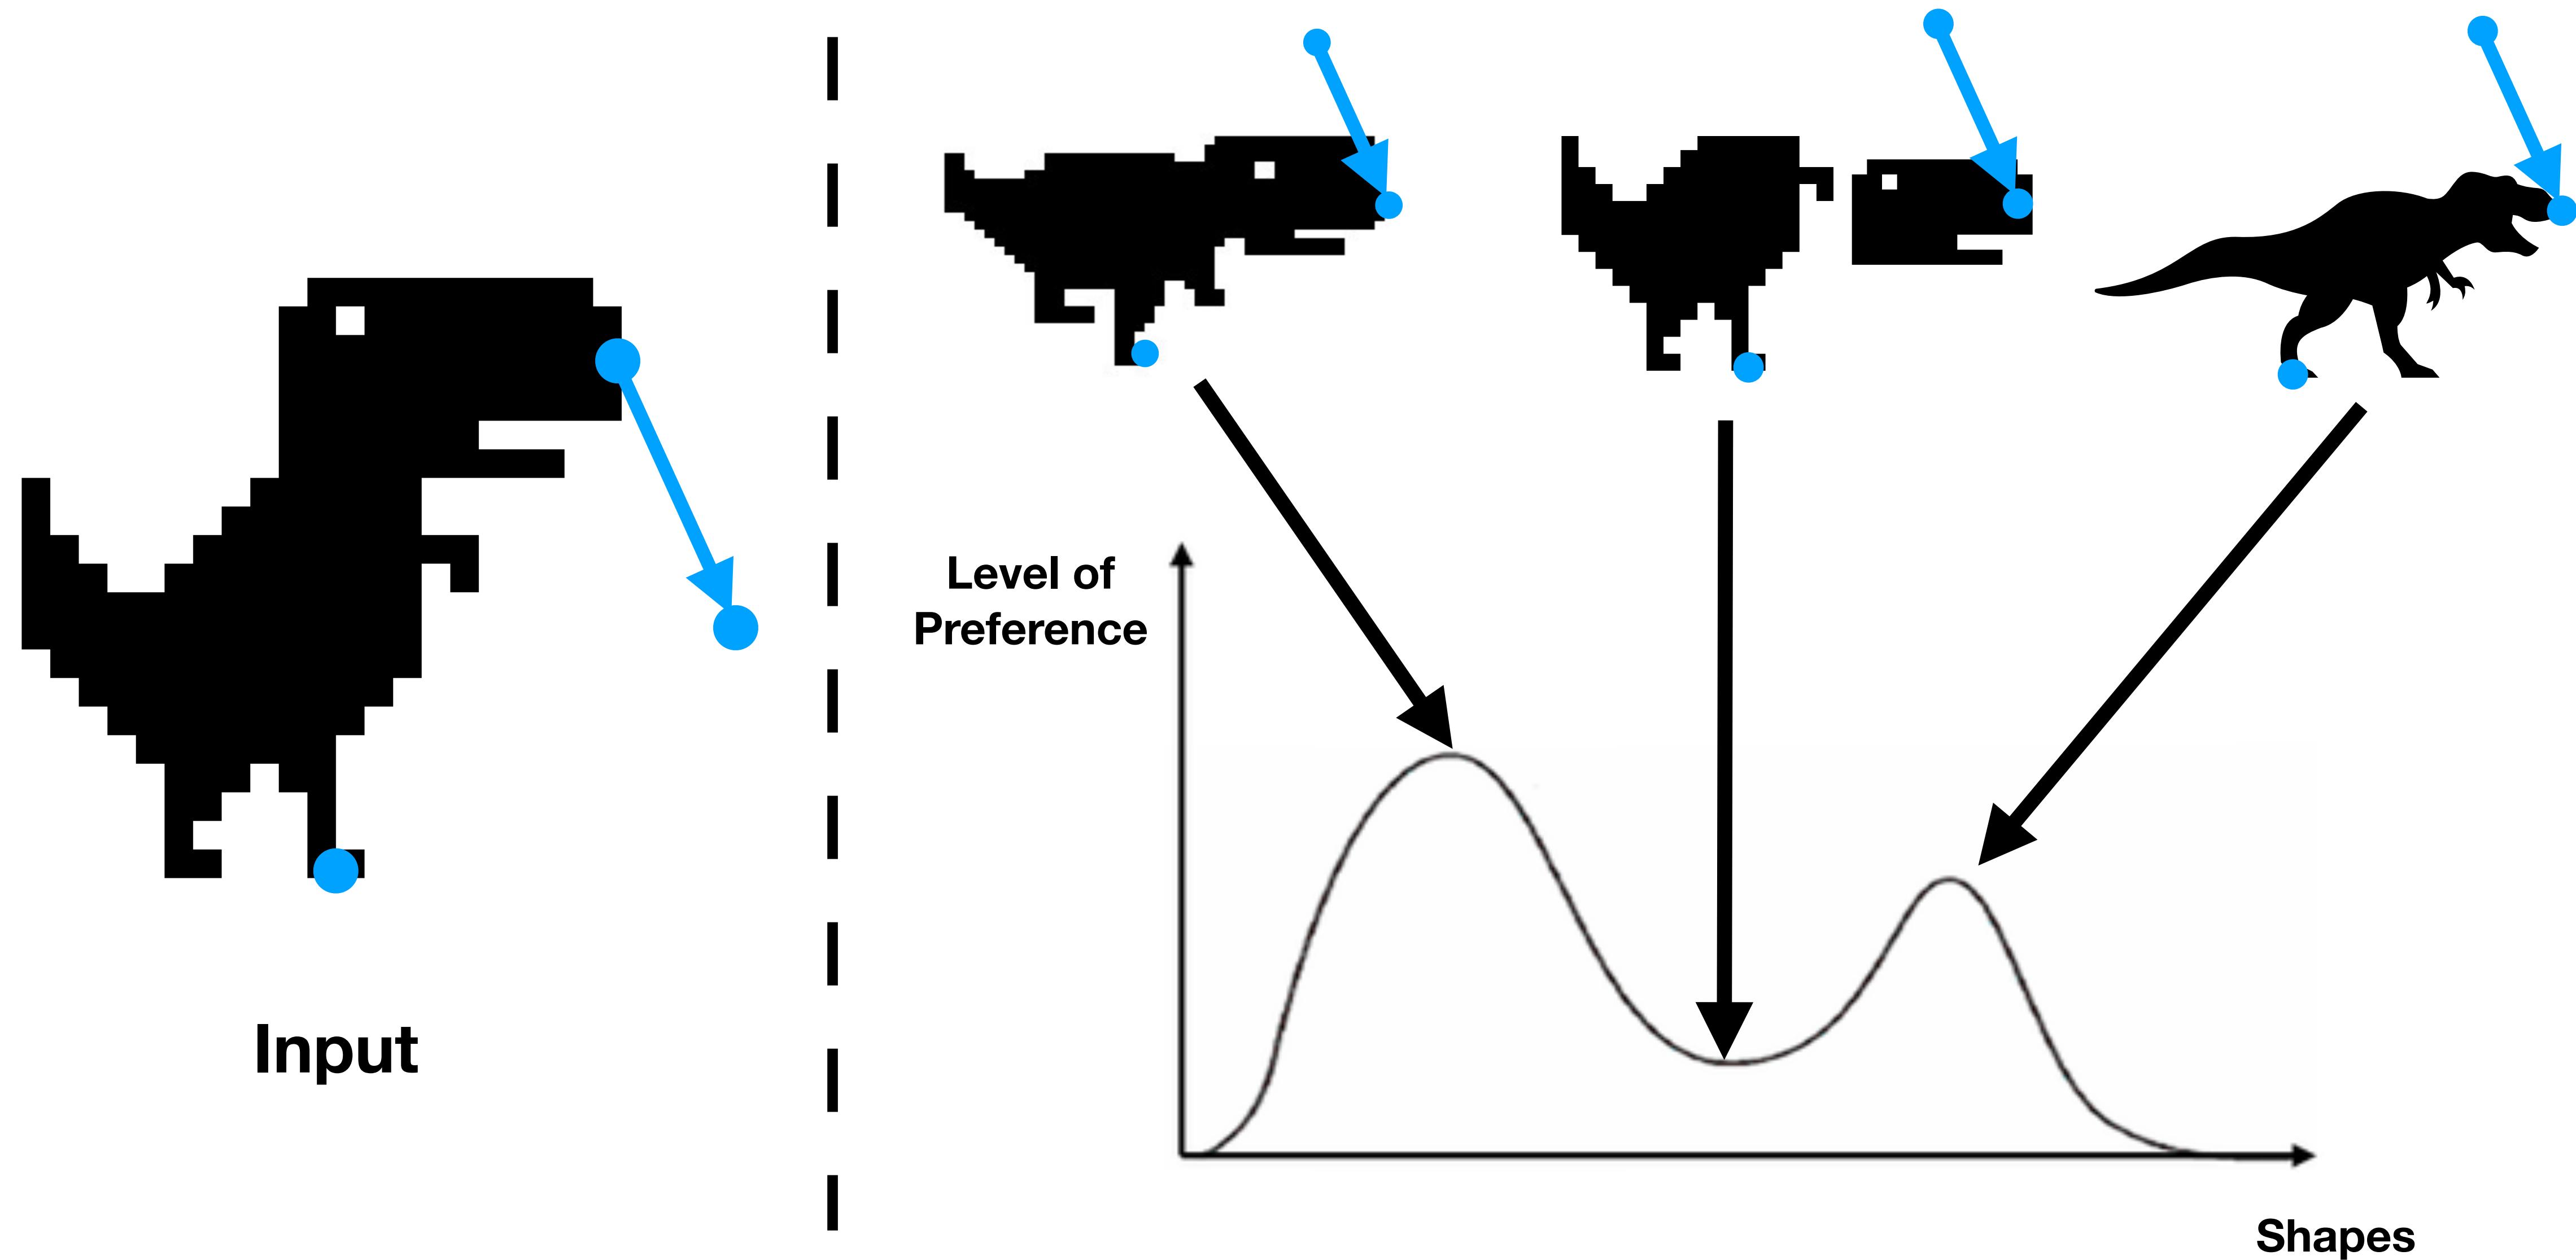
✓ High resolution

Can geometry processing be done entirely using Neural Fields?

Can geometry processing be done entirely using Neural Fields?



Geometry Processing is under-constrained



Physics-inspired Prior



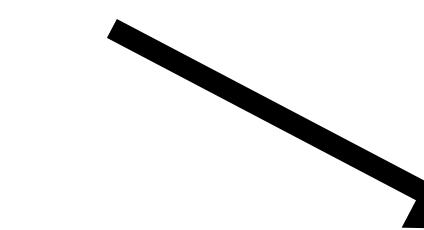
Smooth



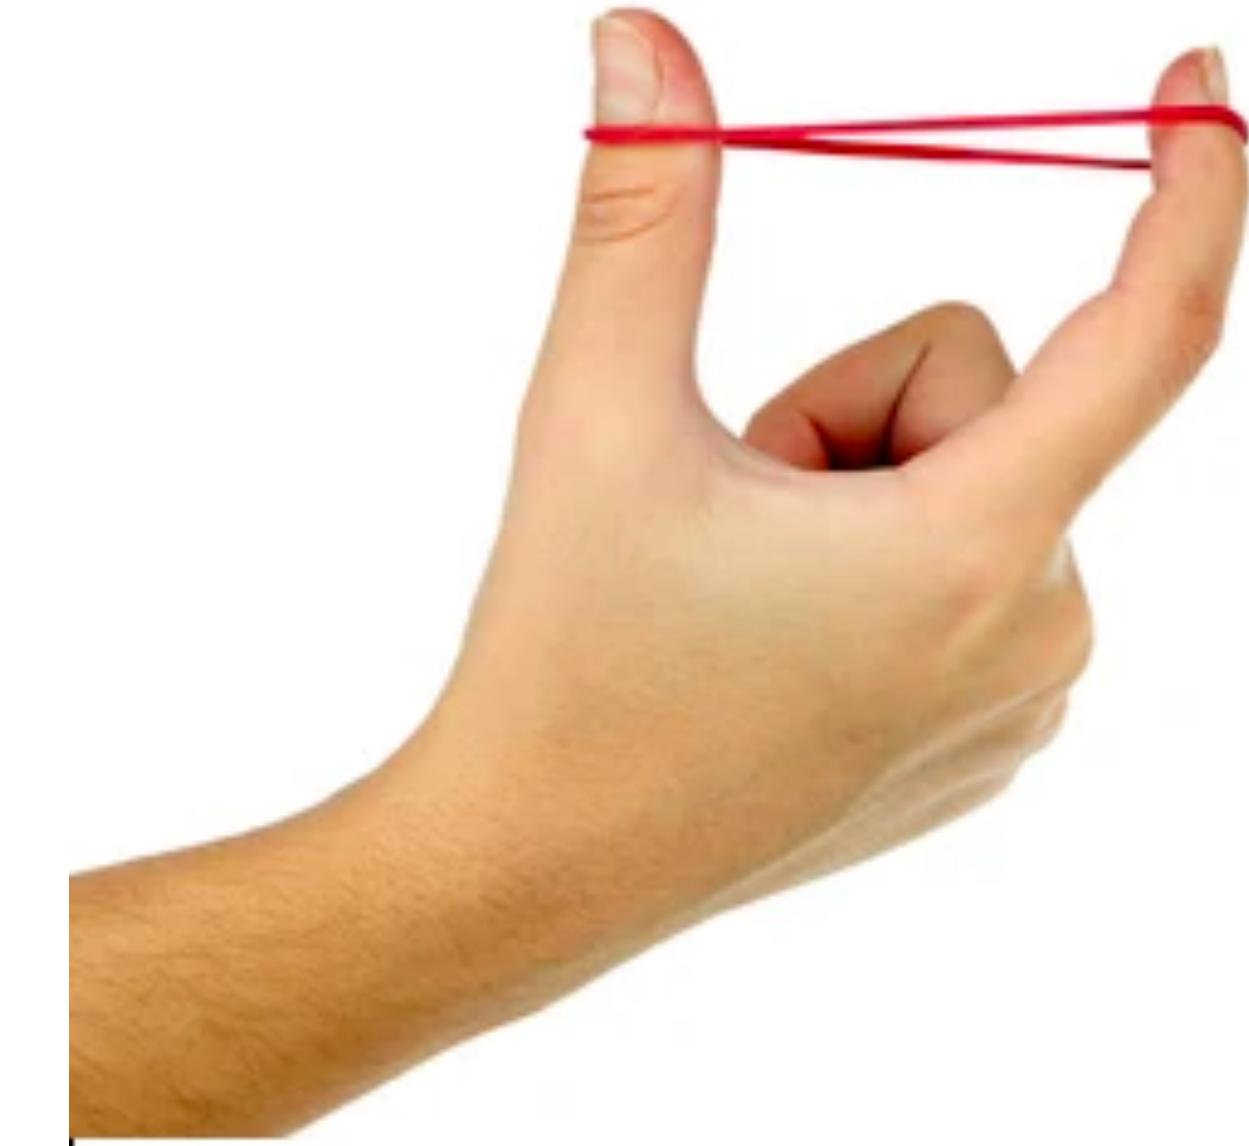
Kazhdan et. al., 2006;
Kazhdan and Hoppe, 2013



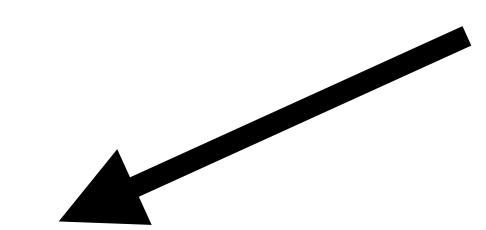
**Bending
Resistance**



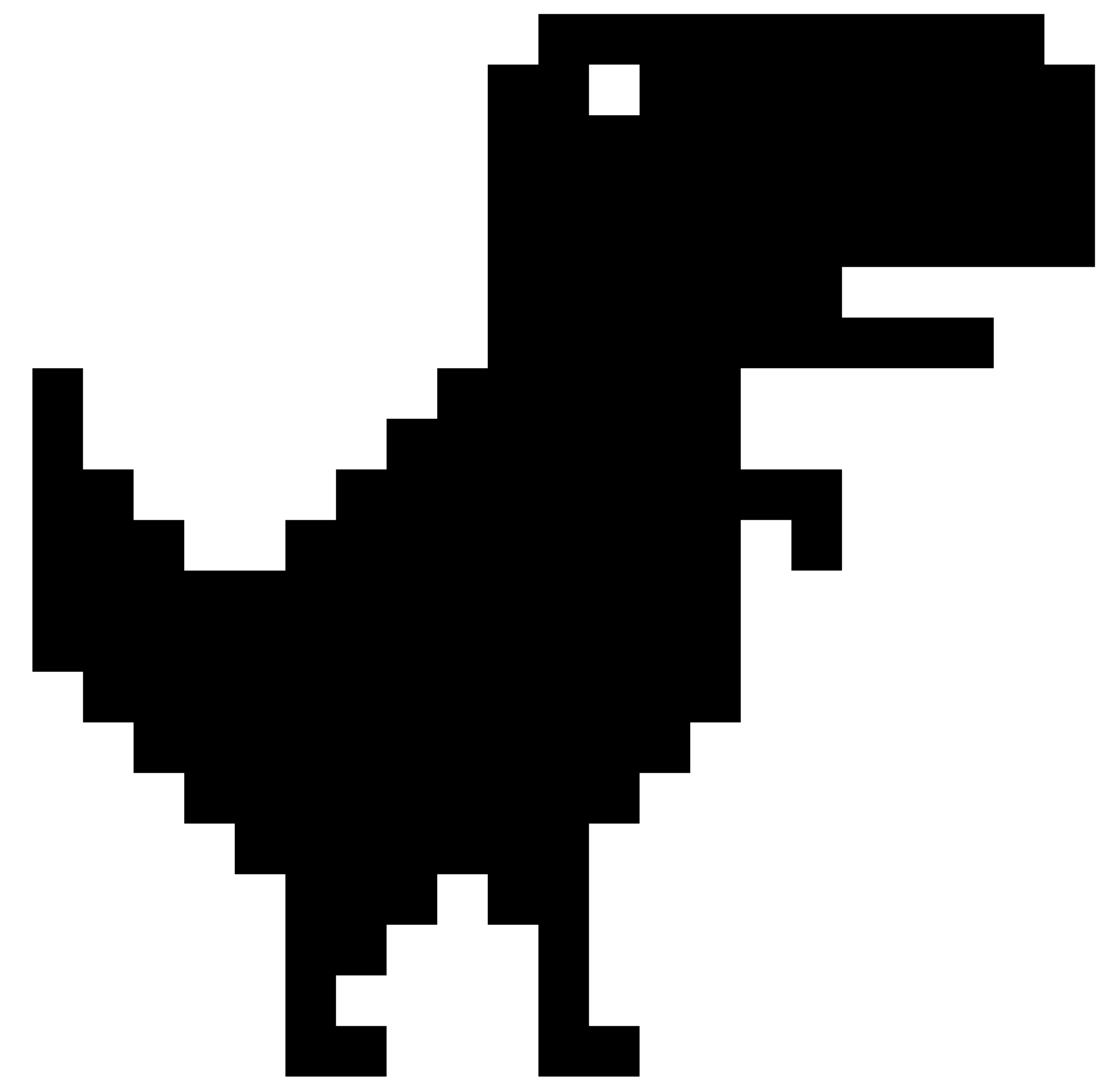
Terzopoulos et. al., 1987;
Sorkine and Alexa, 2007;
Levi and Gotsman, 2015



**Stretching
Resistance**



Quantify Priors



$$f: (u, v) \mapsto (x, y, z)$$

Normal

Area

Curvature

$$\mathbf{n} = \frac{\mathbf{f}_u \times \mathbf{f}_v}{\|\mathbf{f}_u \times \mathbf{f}_v\|}$$

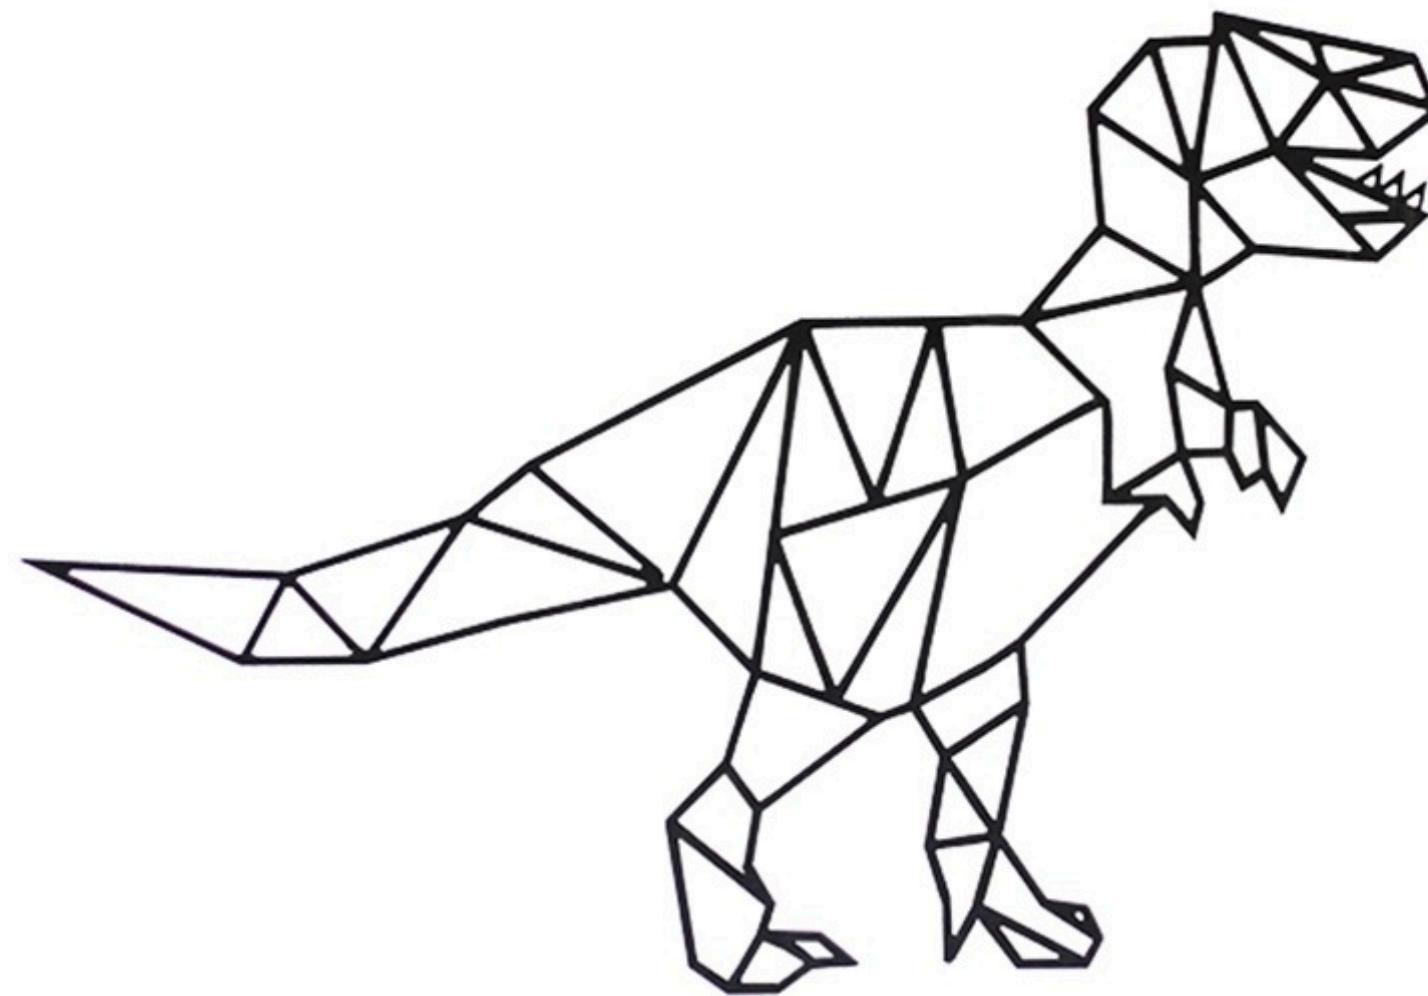
$$\mathbf{I} = \begin{bmatrix} \mathbf{f}_u^T \mathbf{f}_u & \mathbf{f}_u^T \mathbf{f}_v \\ \mathbf{f}_v^T \mathbf{f}_u & \mathbf{f}_v^T \mathbf{f}_v \end{bmatrix}$$

$$\mathbf{II} = \begin{bmatrix} \mathbf{f}_{uu}^T \mathbf{n} & \mathbf{f}_{uv}^T \mathbf{n} \\ \mathbf{f}_{vu}^T \mathbf{n} & \mathbf{f}_{vv}^T \mathbf{n} \end{bmatrix}$$

$$S = \iint |\mathbf{I}|^{\frac{1}{2}} du dv$$

$$\kappa = \frac{|\mathbf{II}|}{|\mathbf{I}|}$$

Quantify Prior - Mesh



$$V = \{(x_i, y_i, z_i)\}_{i=1}^n$$

$$F = \{(u, v, w) \mid 1 \leq u, v, w, \leq n\}$$

Normal

$$\mathbf{n}_{i,j,k} = \frac{(V_j - V_i) \times (V_k - V_i)}{|(V_j - V_i) \times (V_k - V_i)|}$$

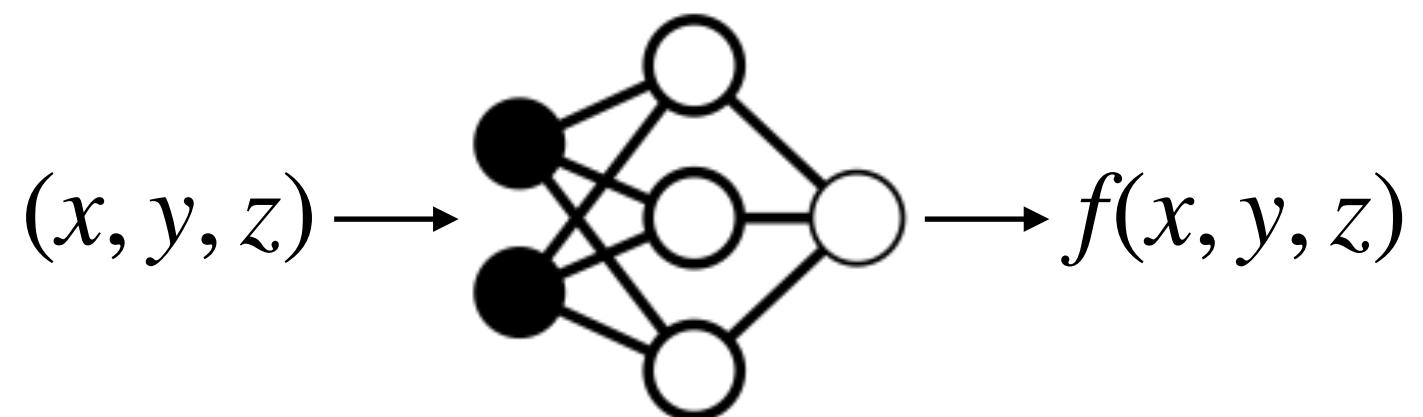
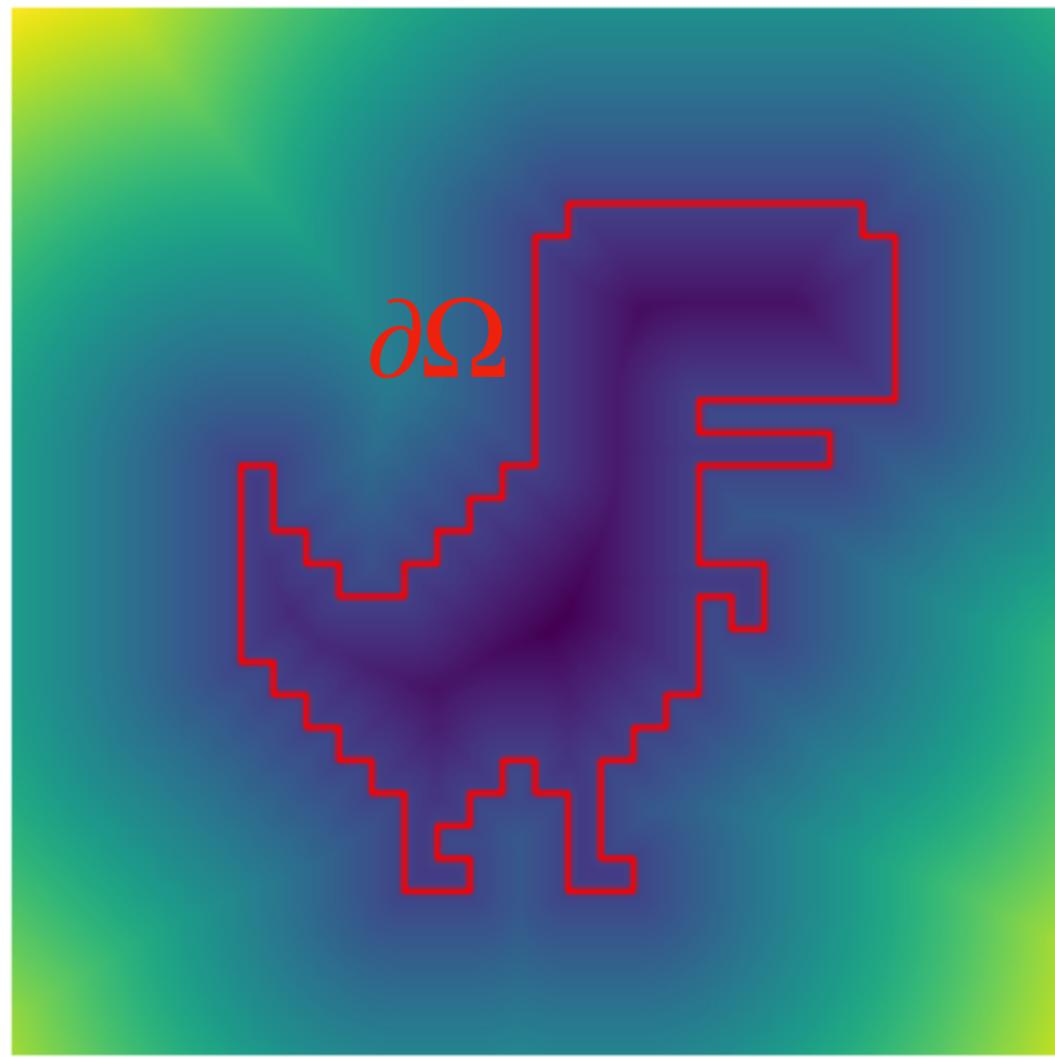
Area

$$S = \sum_{(i,j,k) \in F} \frac{1}{2} |(V_j - V_i) \times (V_k - V_i)|$$

Curvature

$$\kappa_i = \frac{1}{2} \sum_{j \in \mathcal{N}(i)} (\cot \alpha_{ij} + \cot \beta_{ij})(V_i - V_j)$$

Quantify Prior - Neural Fields



$$\partial\Omega = \{(x, y, z) \mid f(x, y, z) = 0\}$$

Normal



Area

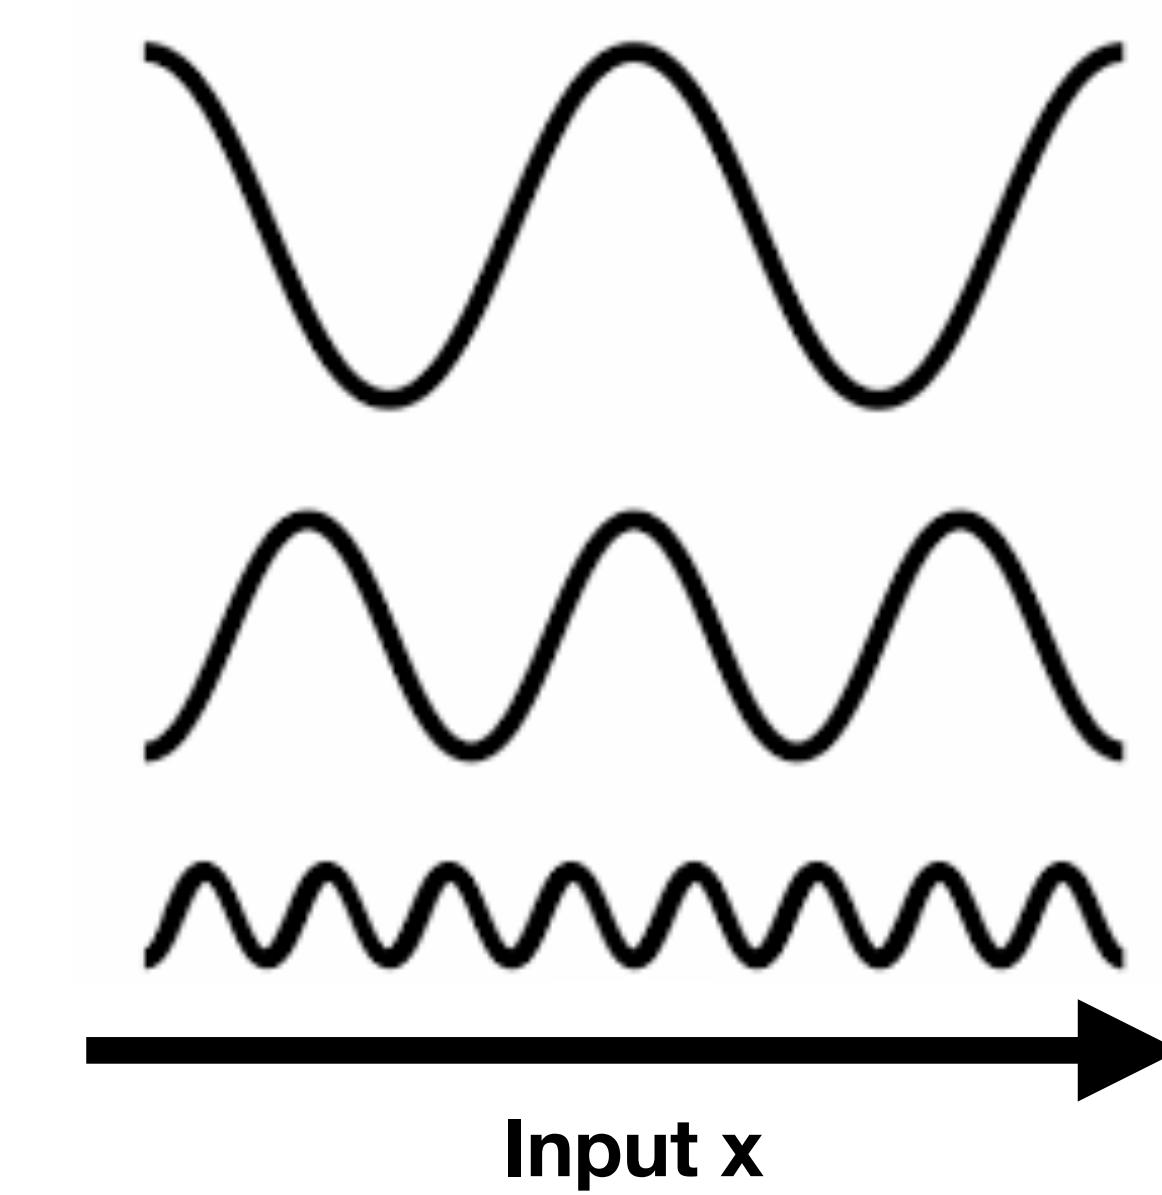
Curvature

Differentiability of Neural Fields

PyTorch

TensorFlow

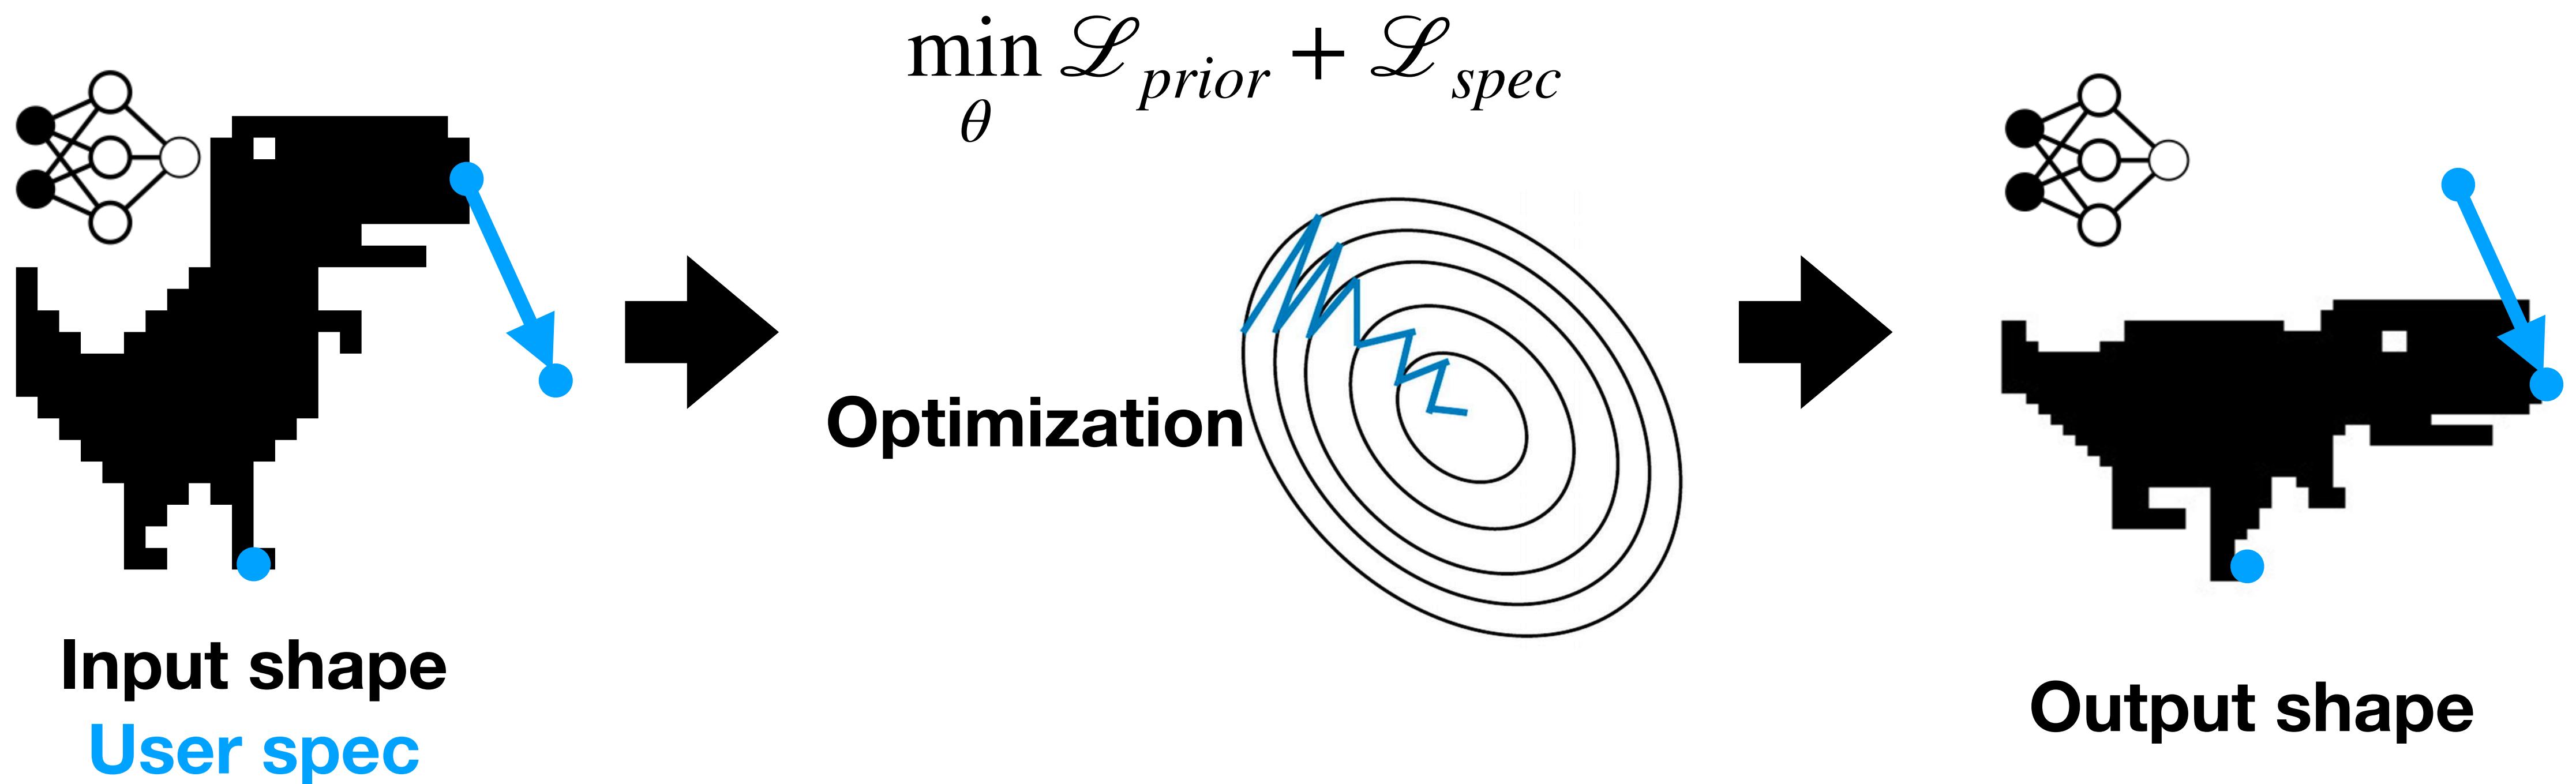
AutoGrad Frameworks



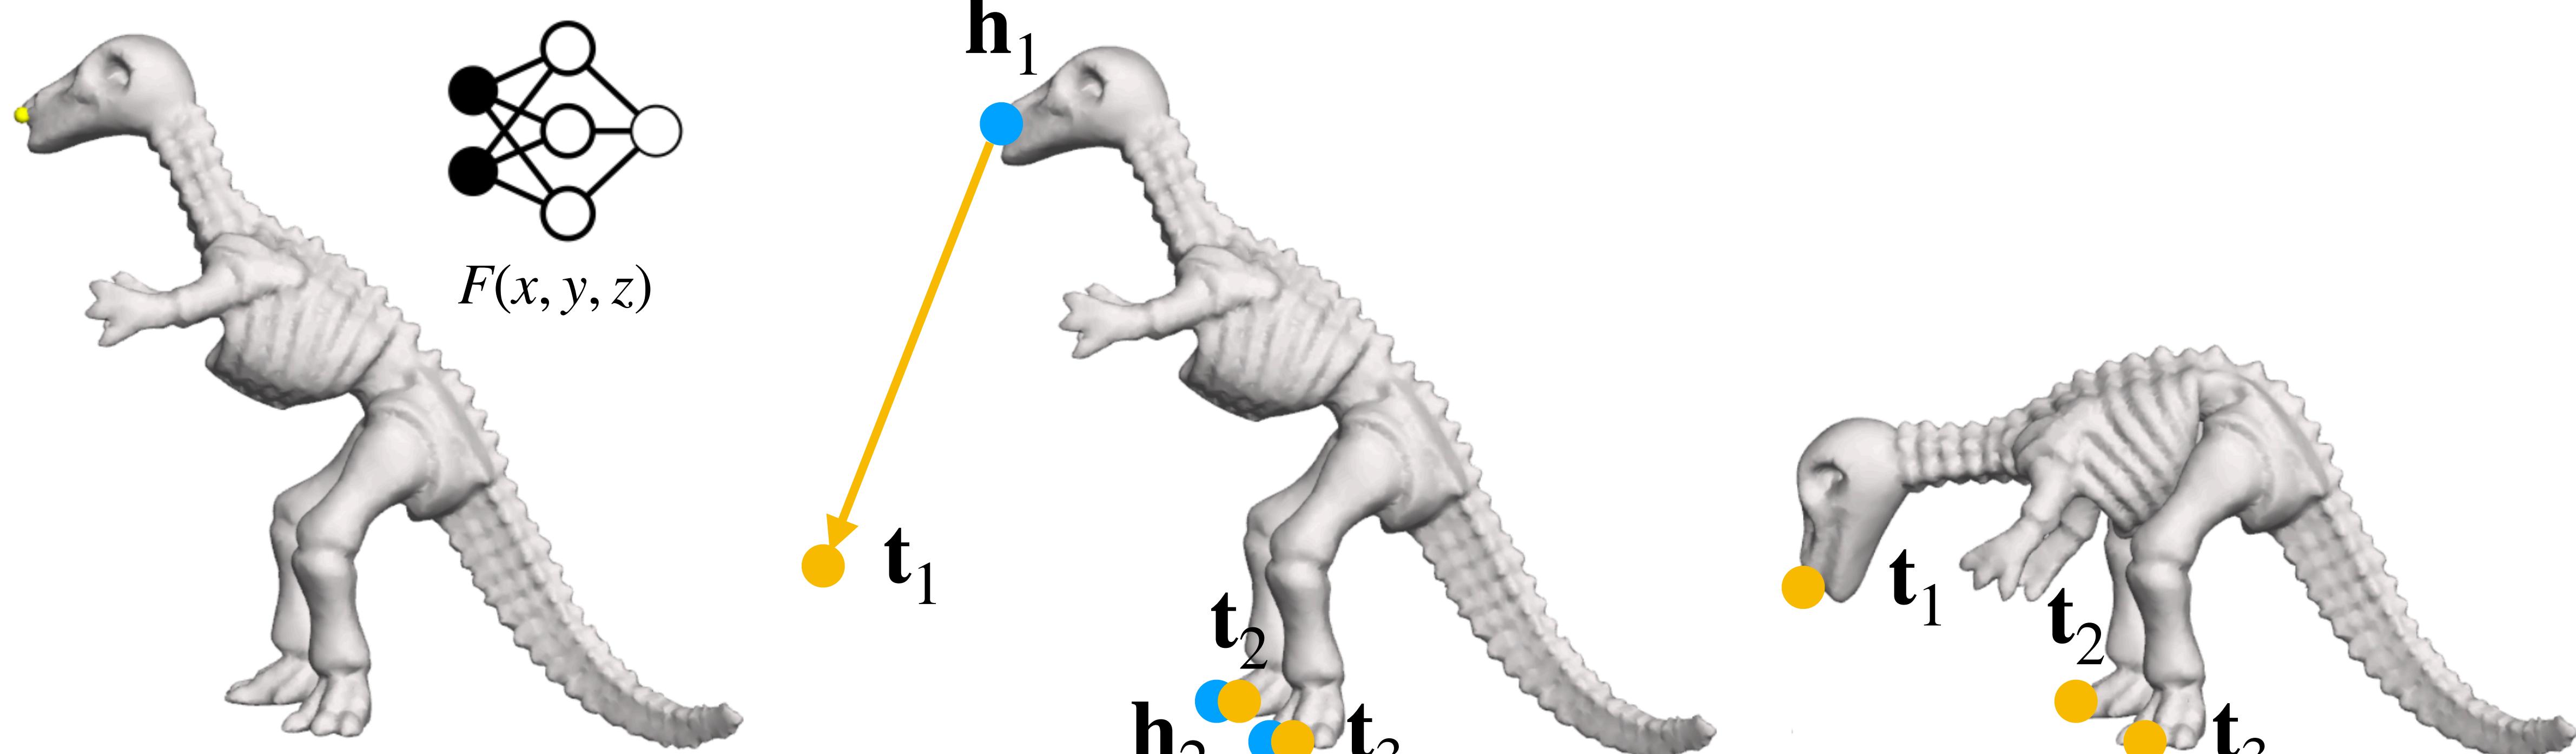
Positional Encoding

Smooth Activation

Geometry Processing with Neural Fields



Deformation with Neural Fields



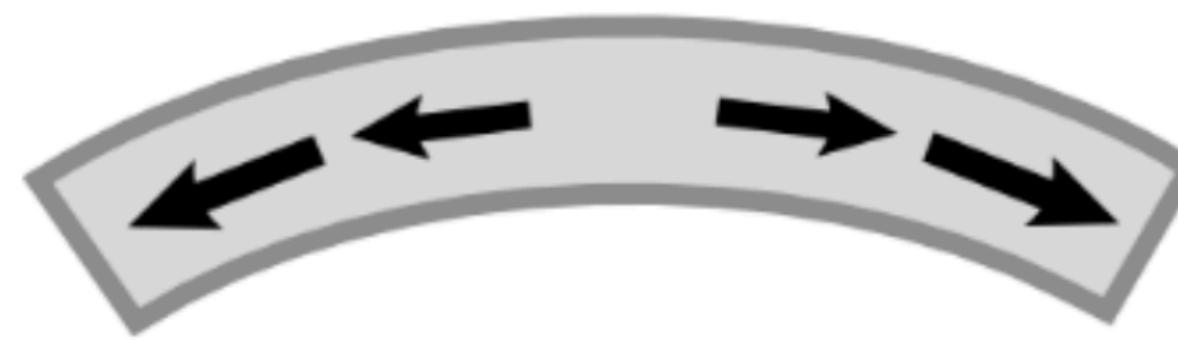
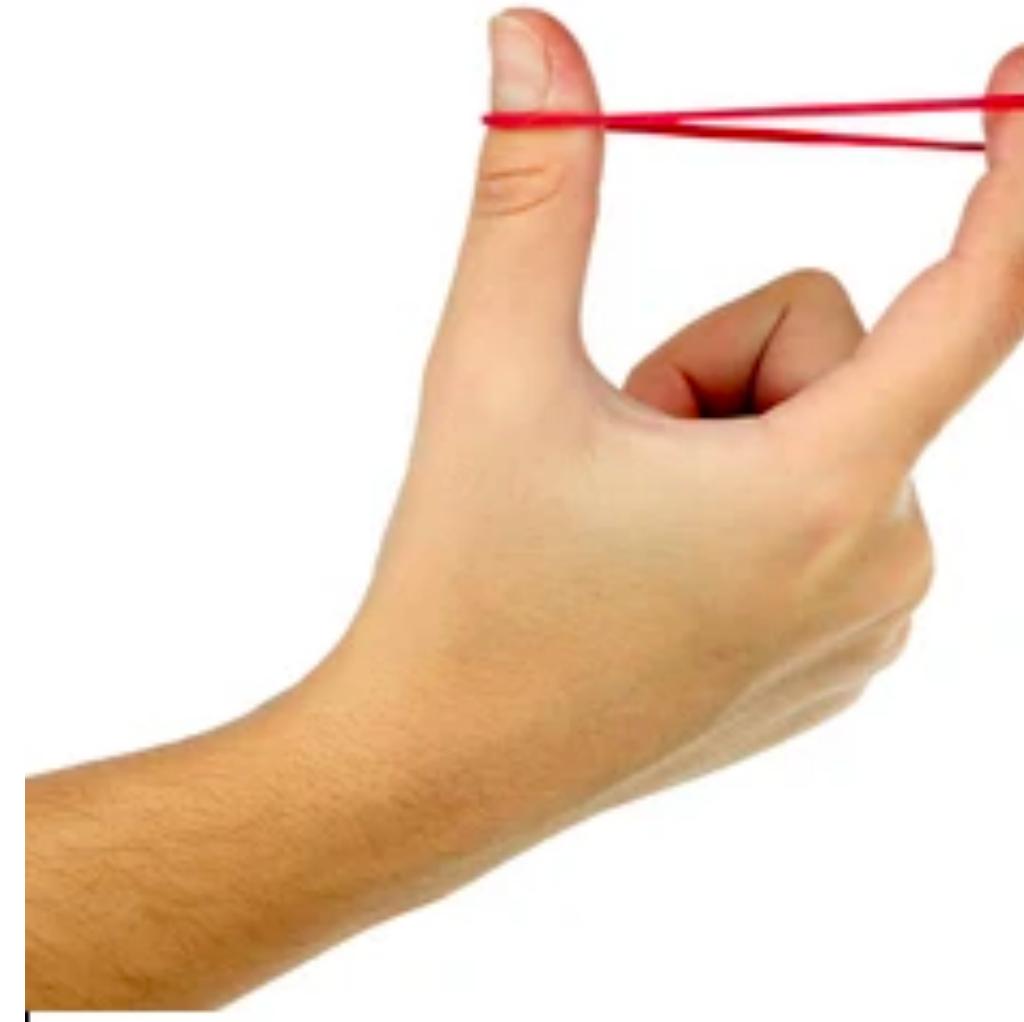
$$F(\mathbf{x}) \approx \sigma(\mathbf{x}) \min_{\mathbf{p} \in \partial\Omega} \|\mathbf{p} - \mathbf{x}\|$$

$$\sigma(\mathbf{x}) = \begin{cases} -1 & \text{if } \mathbf{x} \in \Omega \\ 1 & \text{otherwise} \end{cases}$$

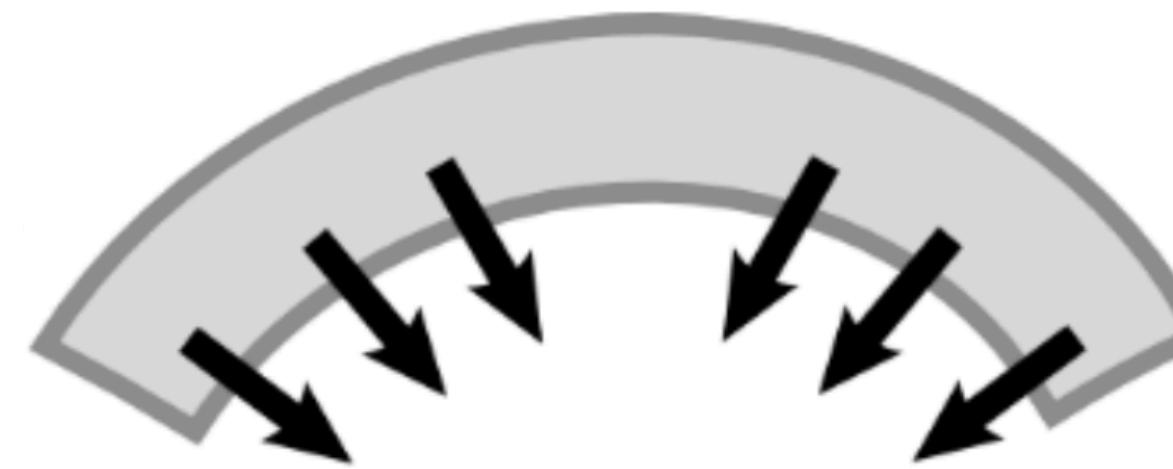
User specification

Desired Output
Natural Deformation

Elastic Deformation

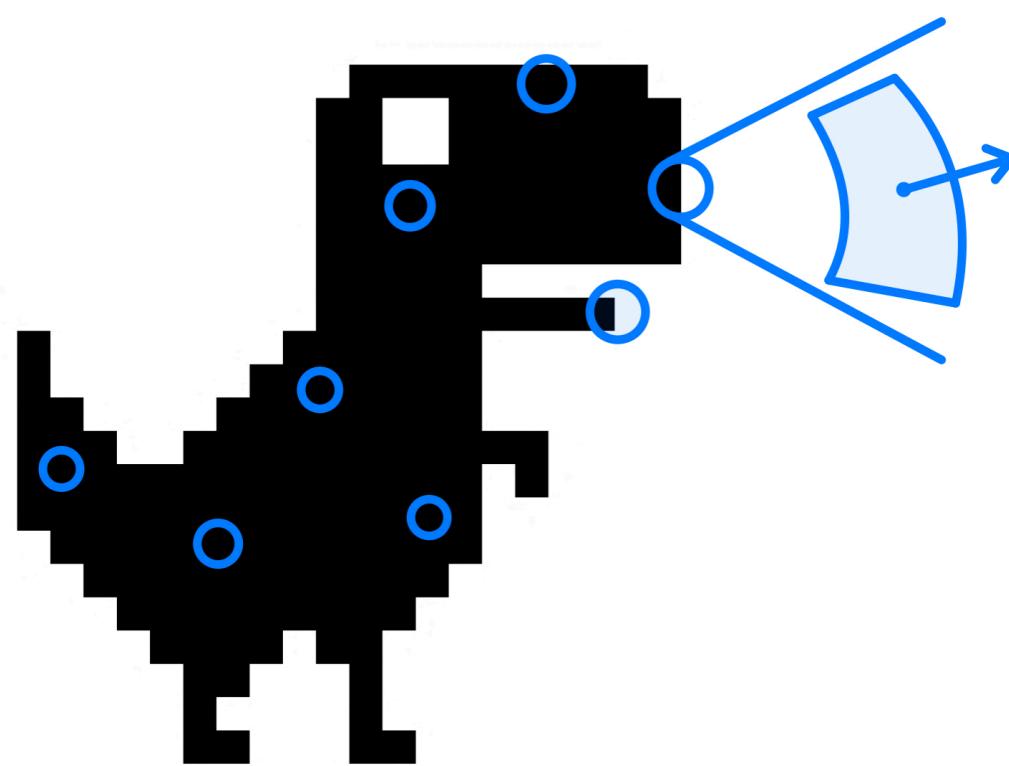
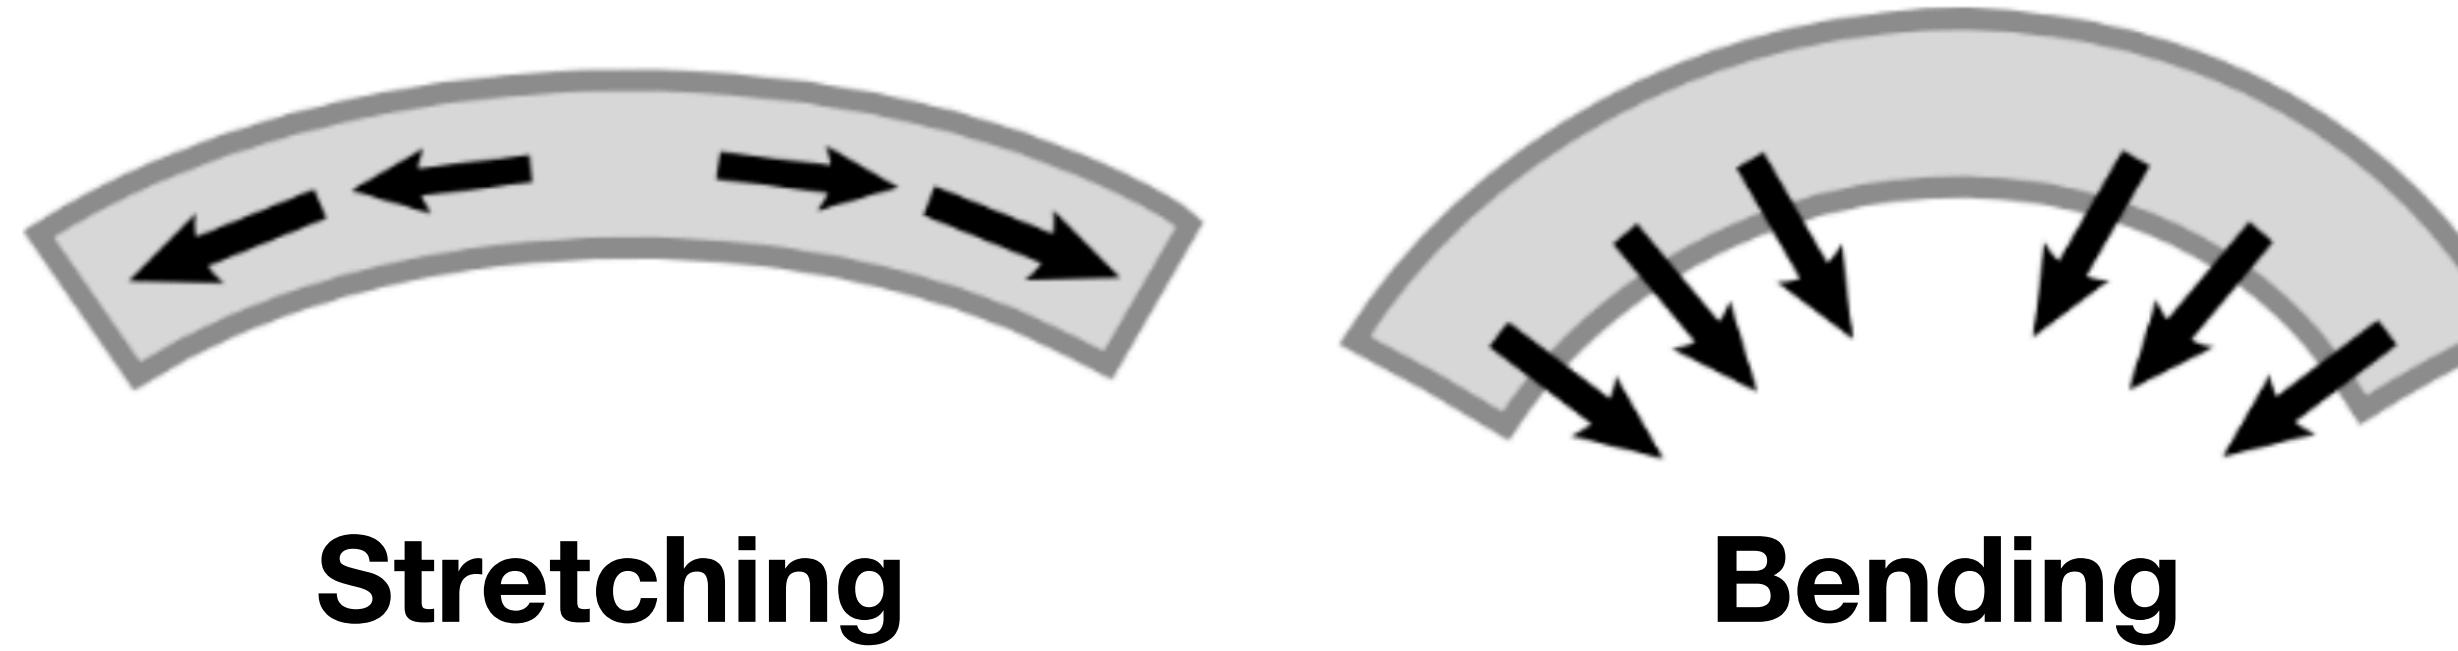


Stretching



Bending

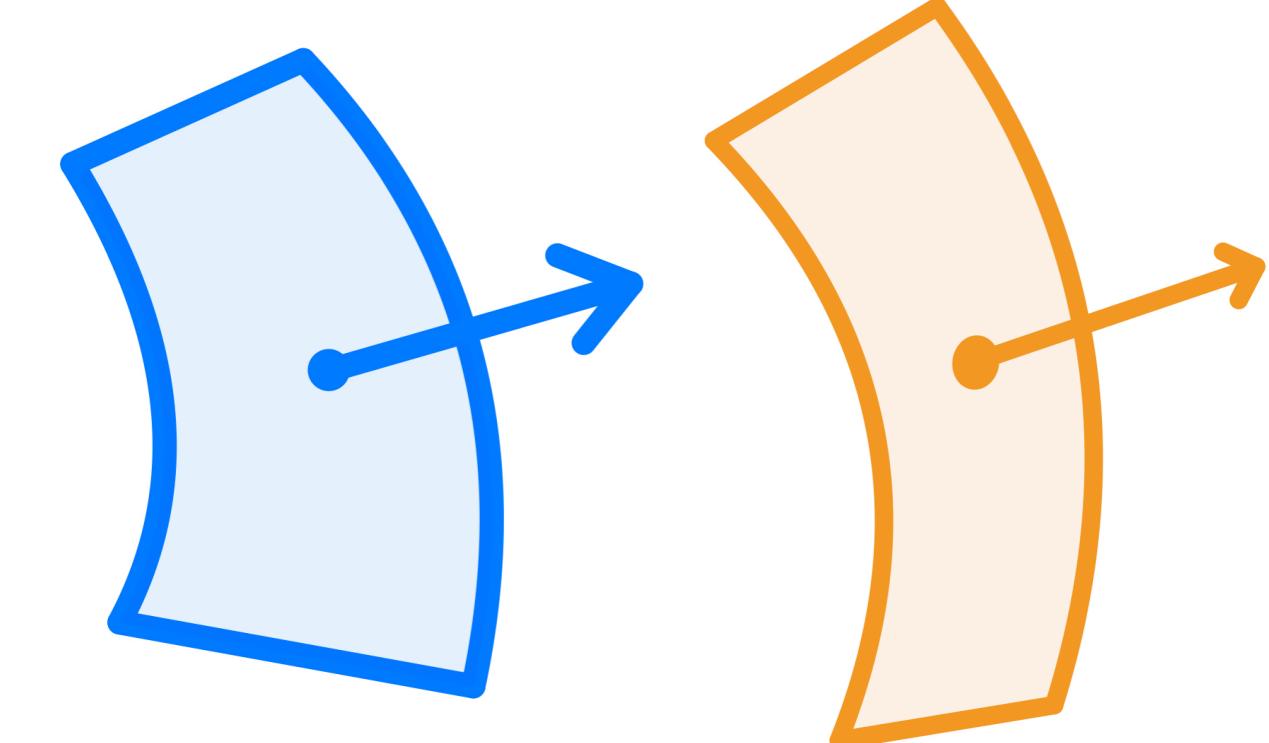
Elastic Deformation



1. Sampling



2. Correspondences

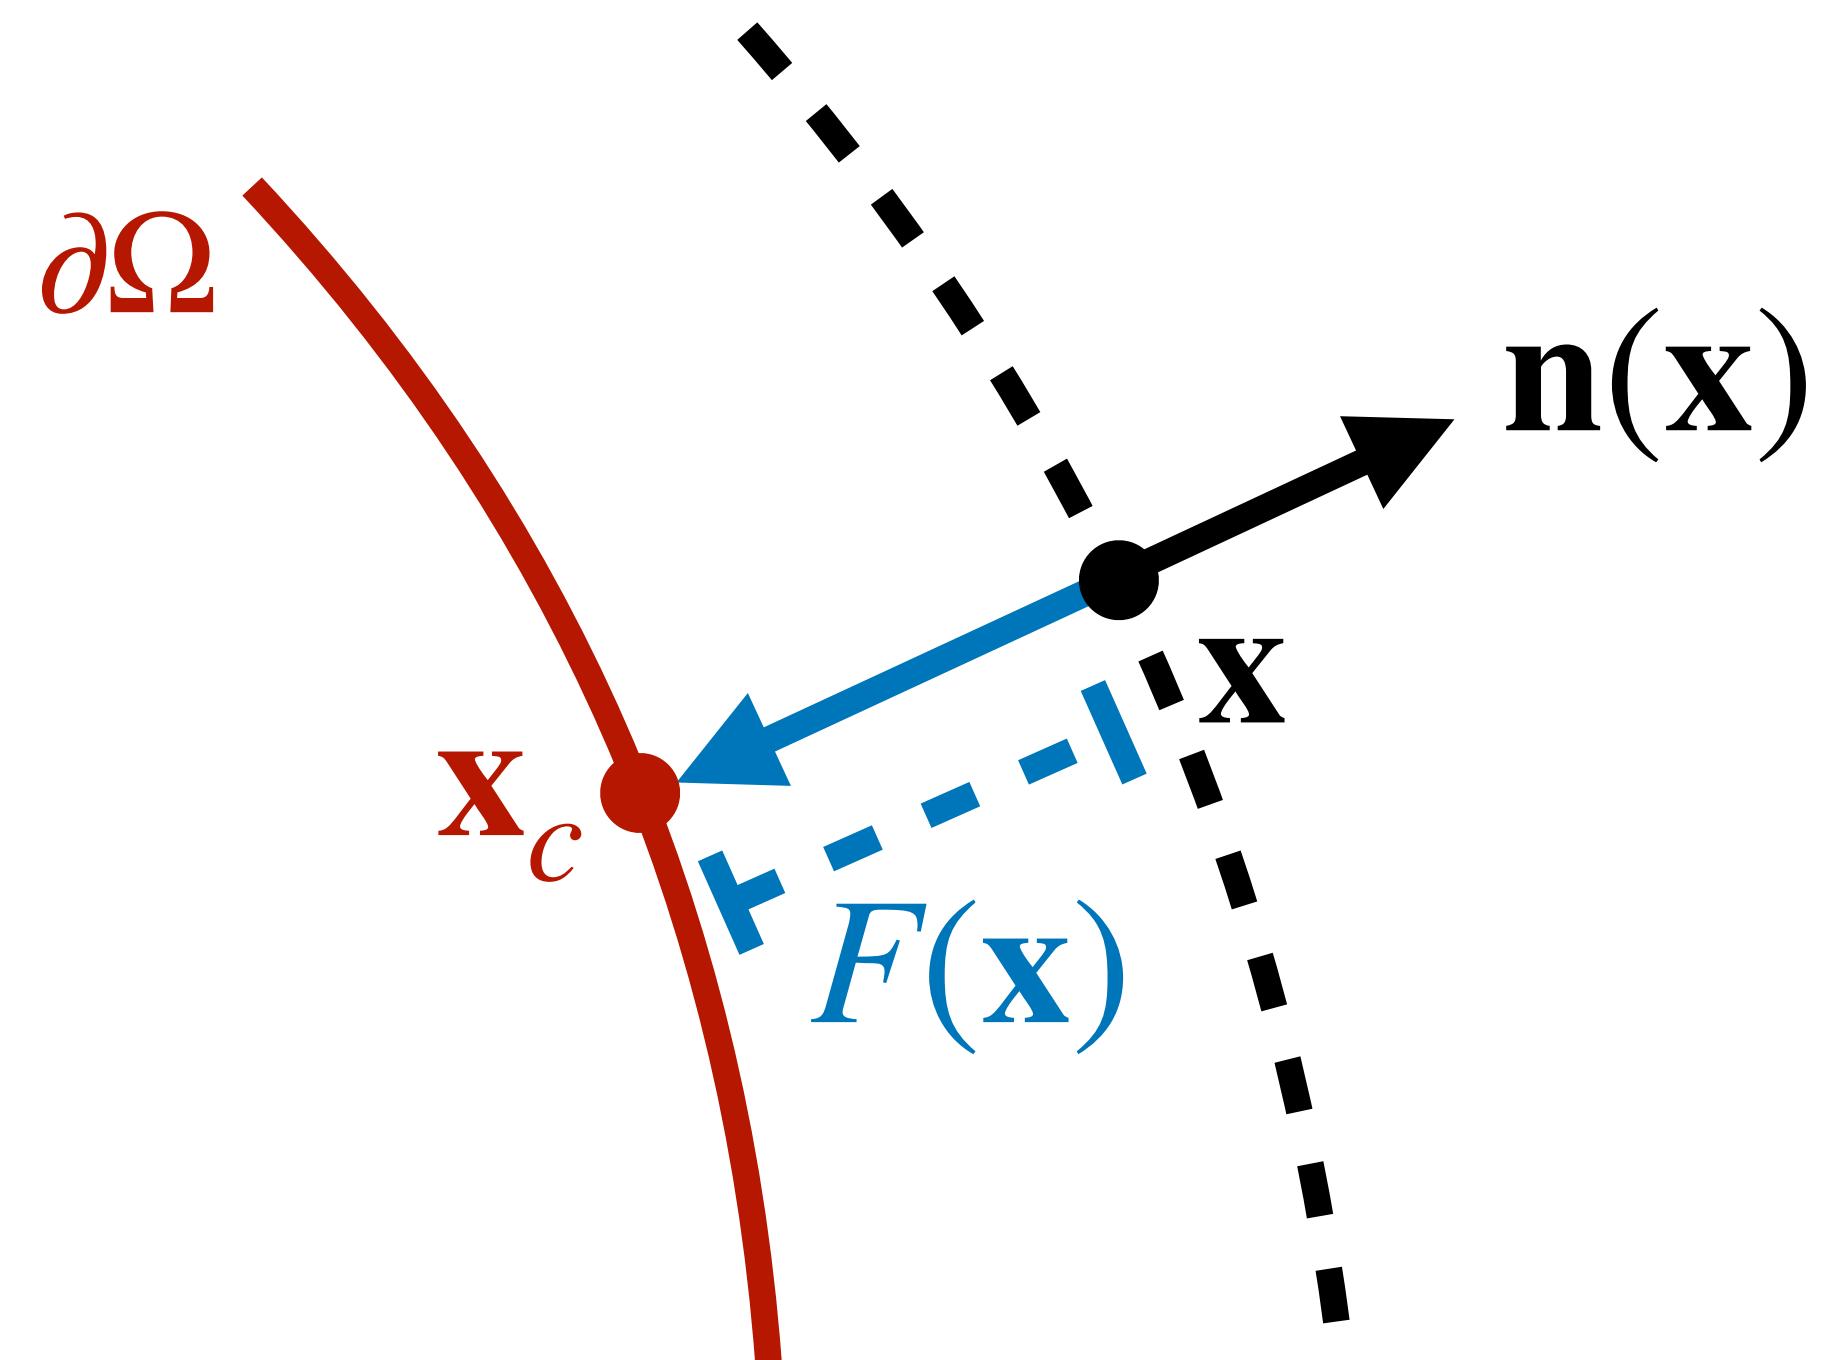


3. Comparing

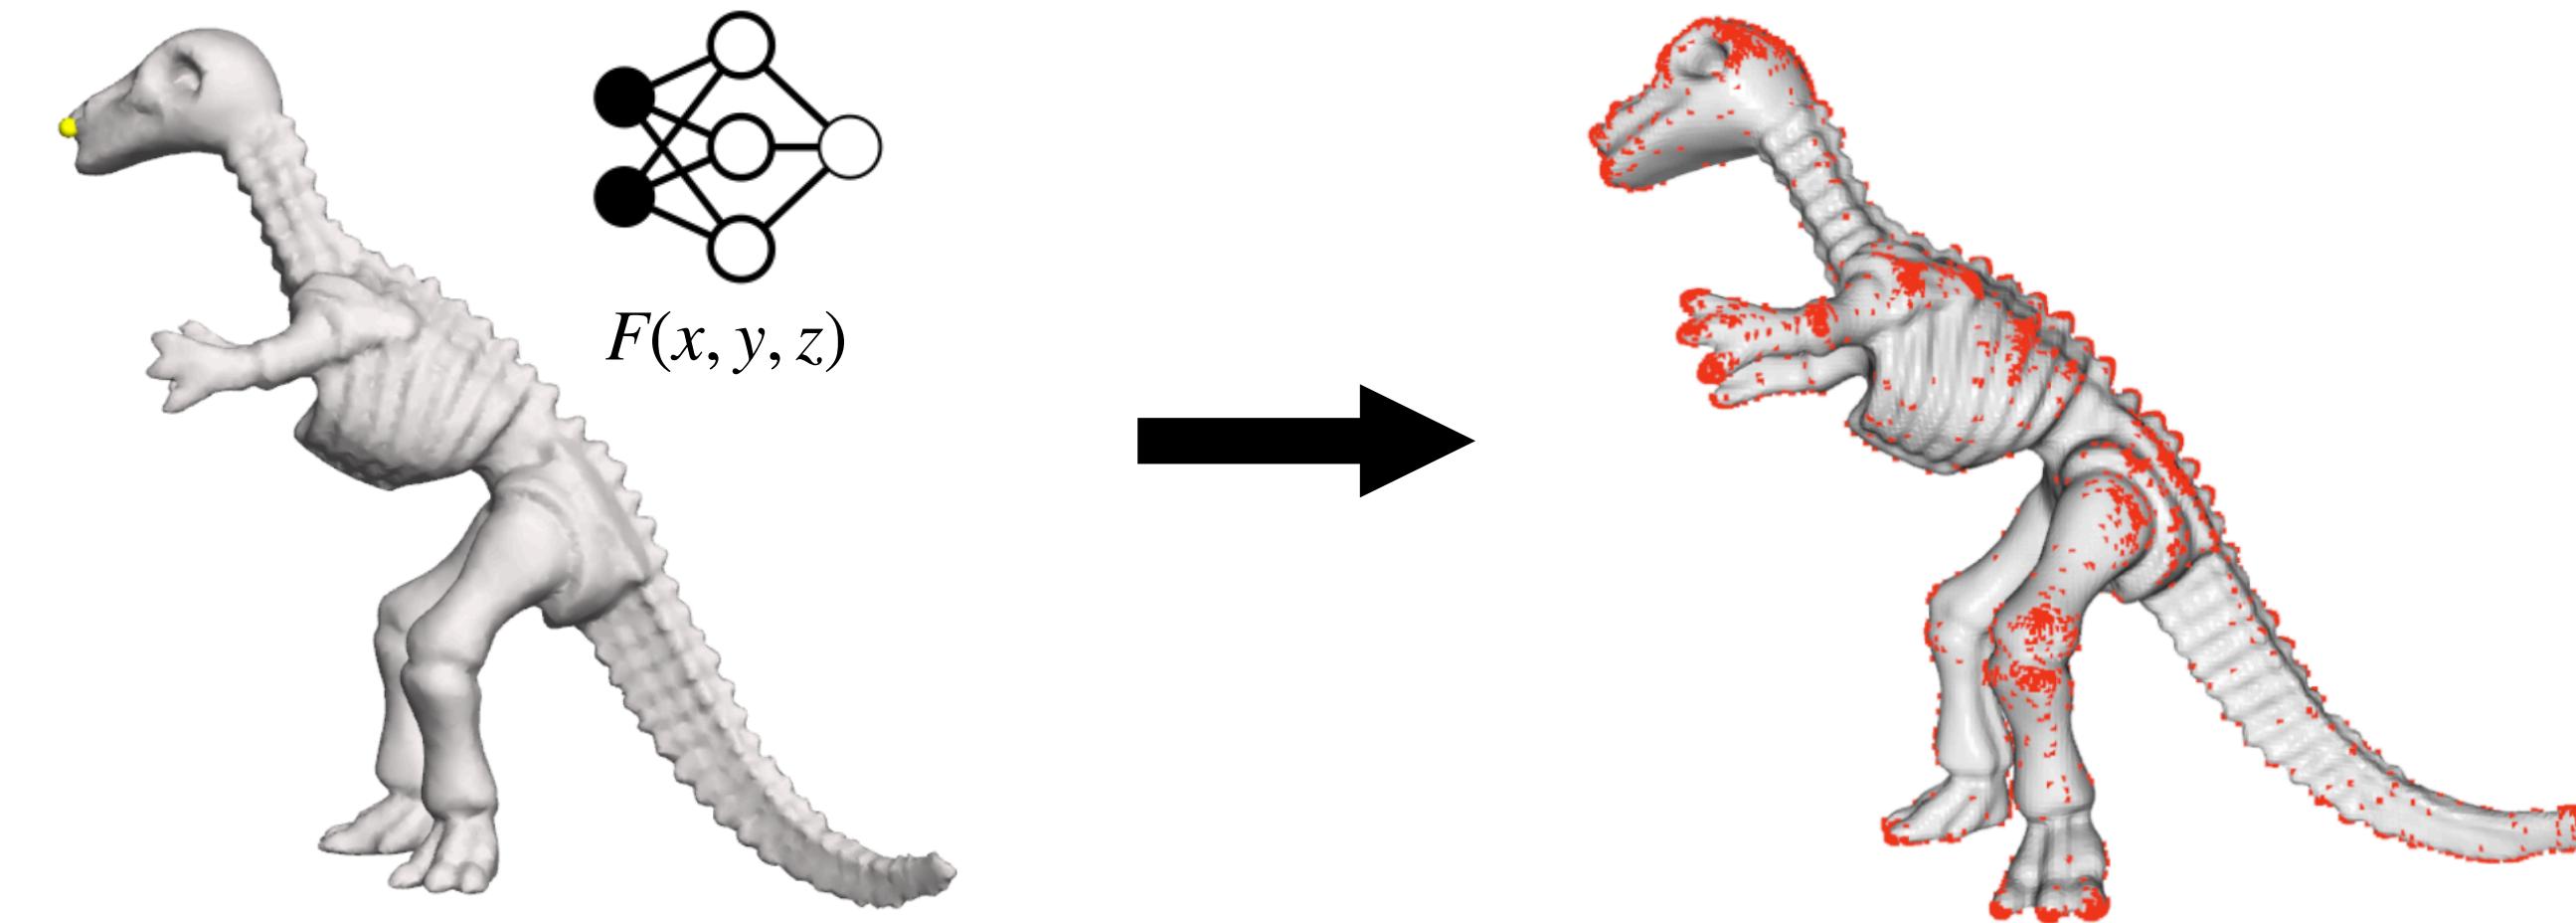
Step 1: Sampling

$$F(\mathbf{x}) \approx SDF(\mathbf{x})$$

$$\begin{aligned}\mathbf{x}_c &= \arg \min_{\mathbf{p} \in \partial\Omega} |\mathbf{p} - \mathbf{x}| \\ &= \mathbf{x} - F(\mathbf{x})\mathbf{n}(\mathbf{x})\end{aligned}$$



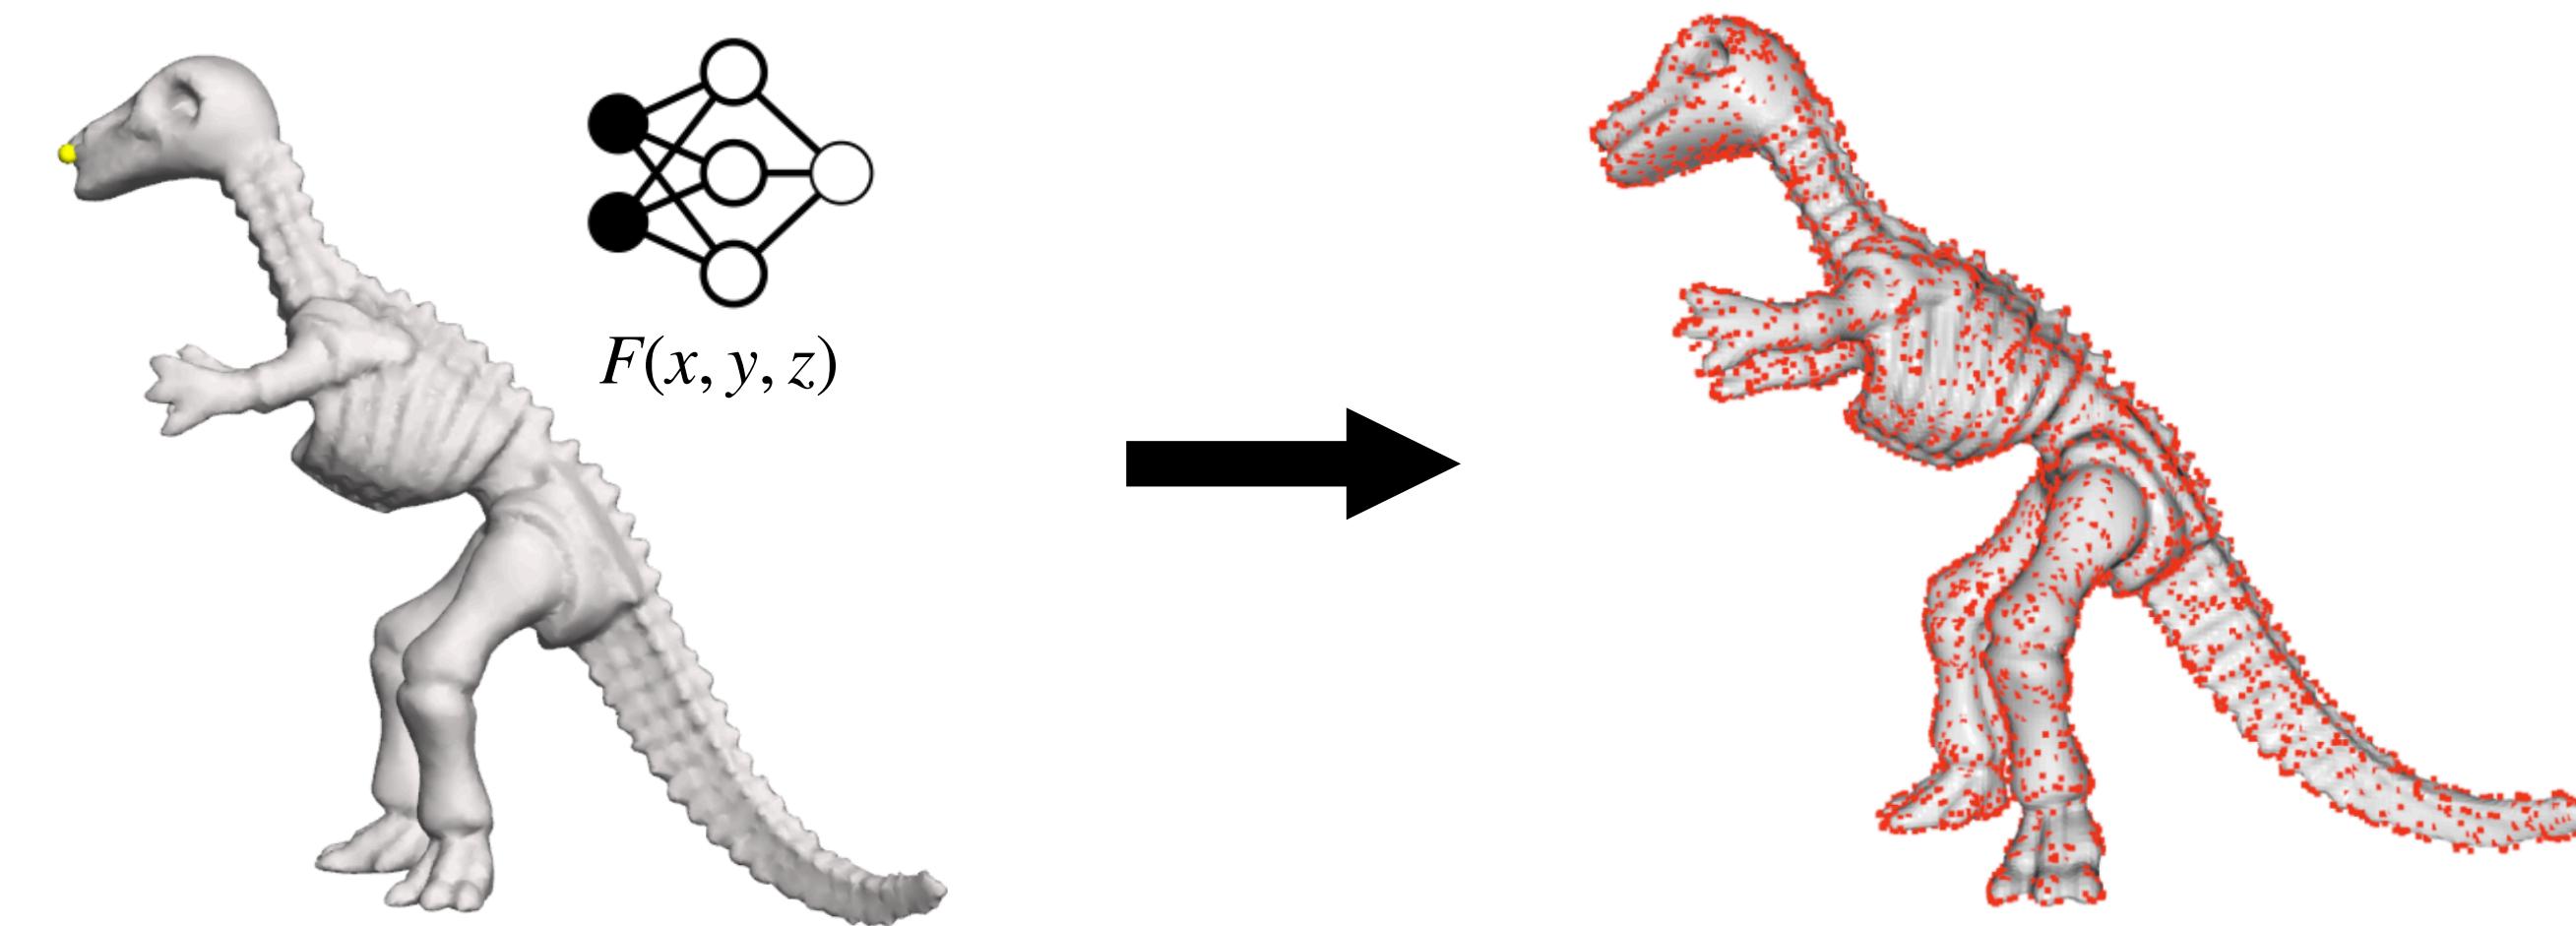
Step 1: Sampling



$$\mathbf{x}_0 \sim U([-1, 1]^3)$$

$$\hat{\mathbf{x}} = \mathbf{x}_0 - F(\mathbf{x}_0)\mathbf{n}(\mathbf{x}_0)$$

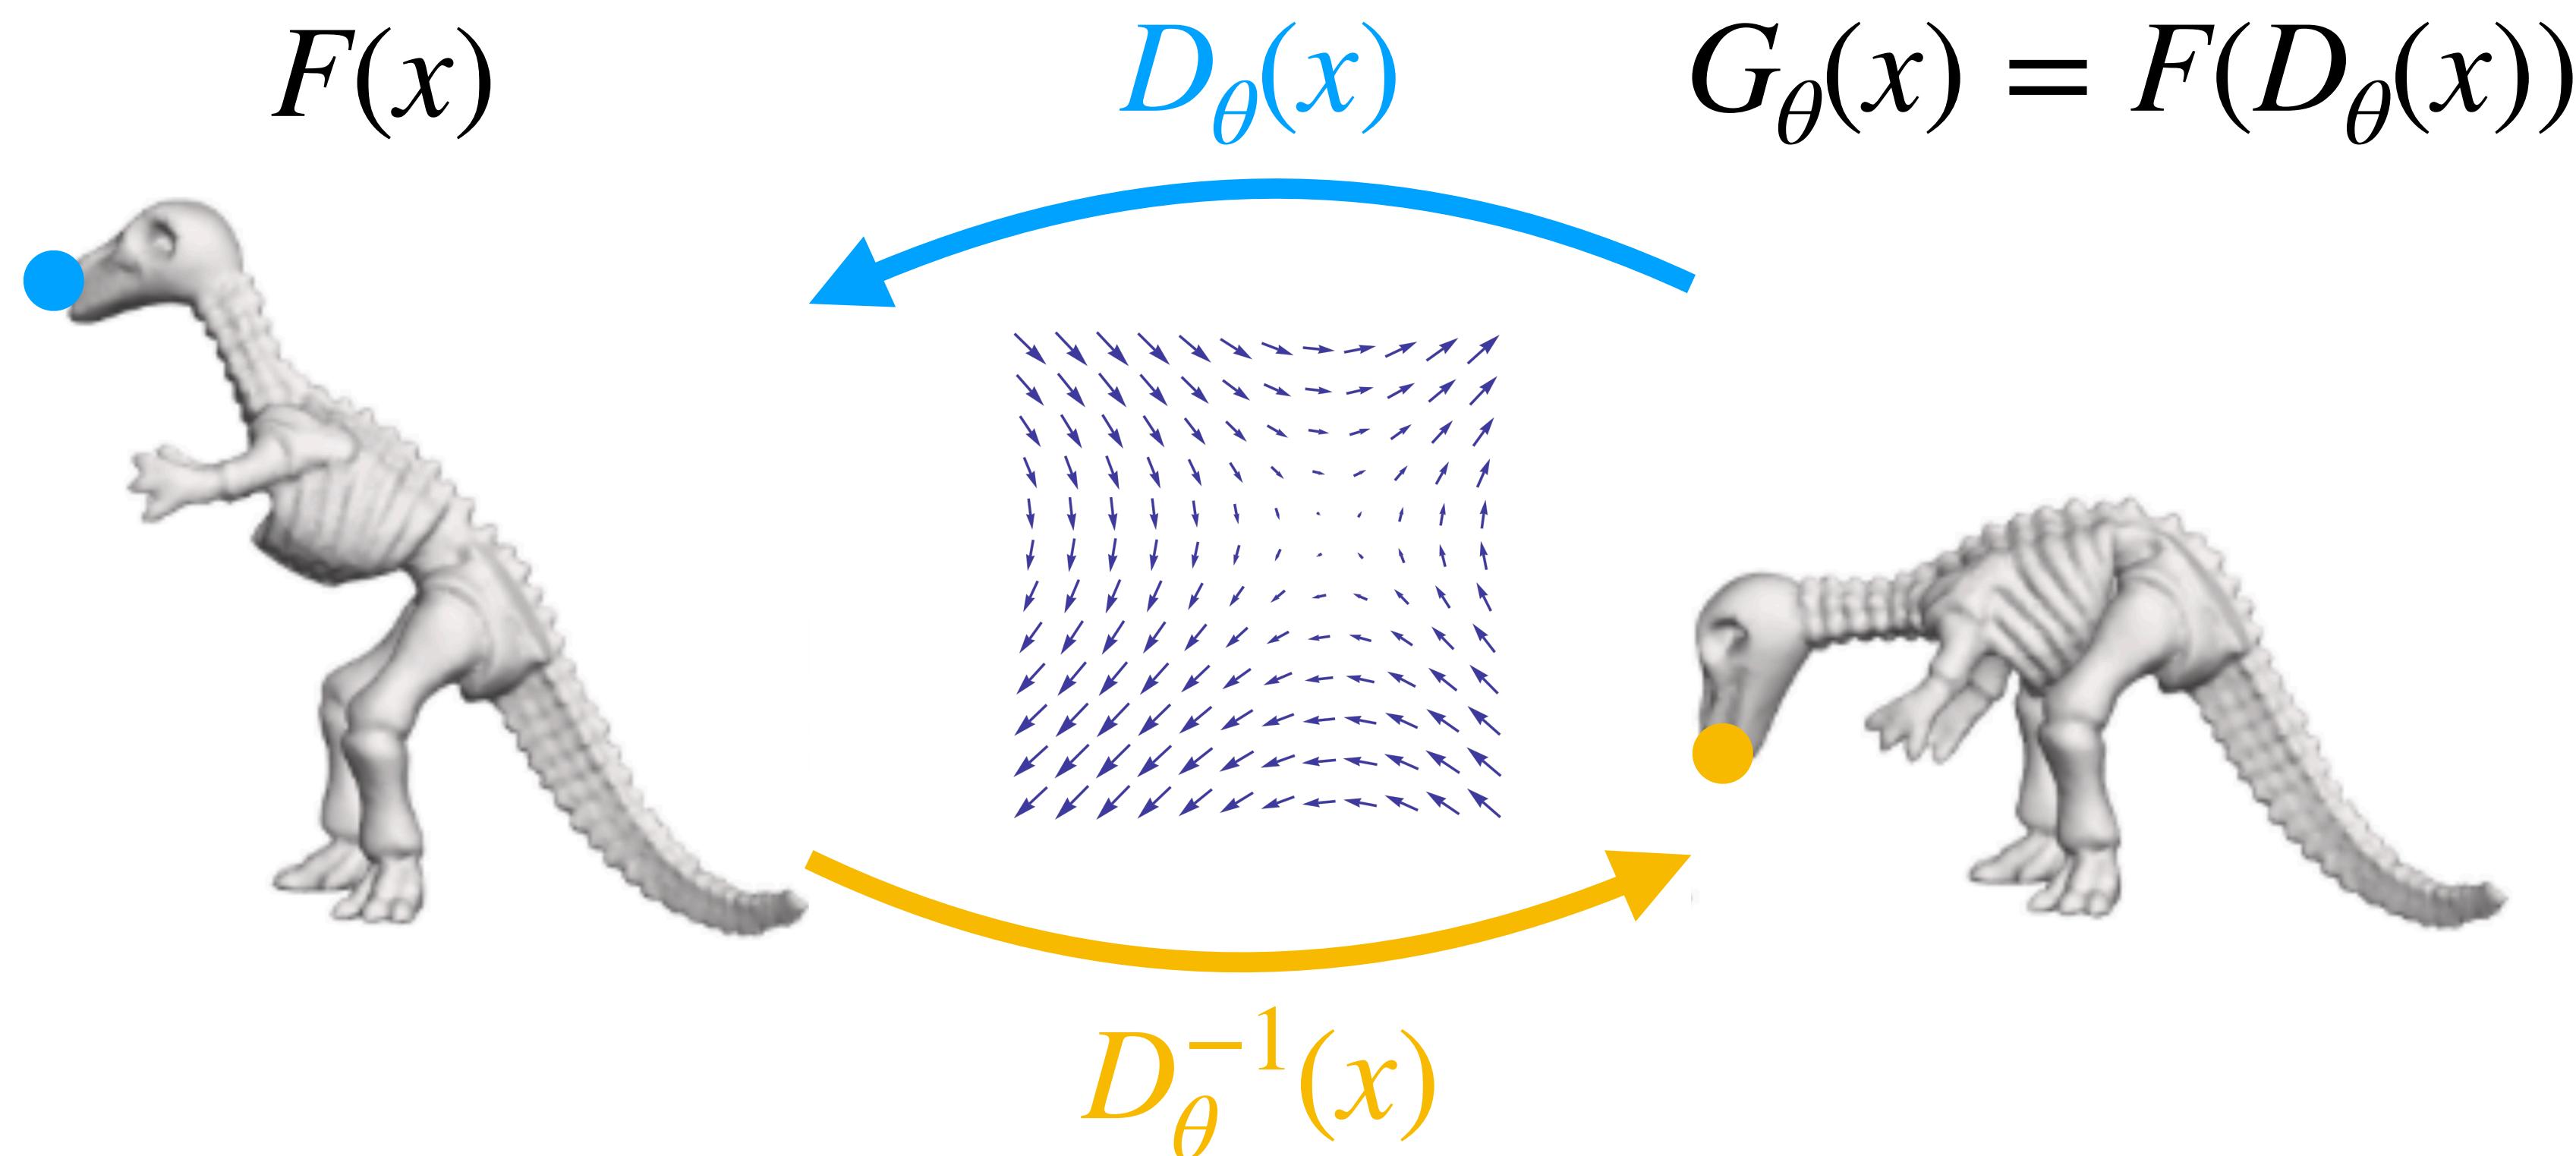
Step 1: Sampling



$$\mathbf{x}_0 \sim \{ \mathbf{x} \mid \mathbf{x} \in U([-1, 1]^3), F(\mathbf{x}) < \tau \}$$

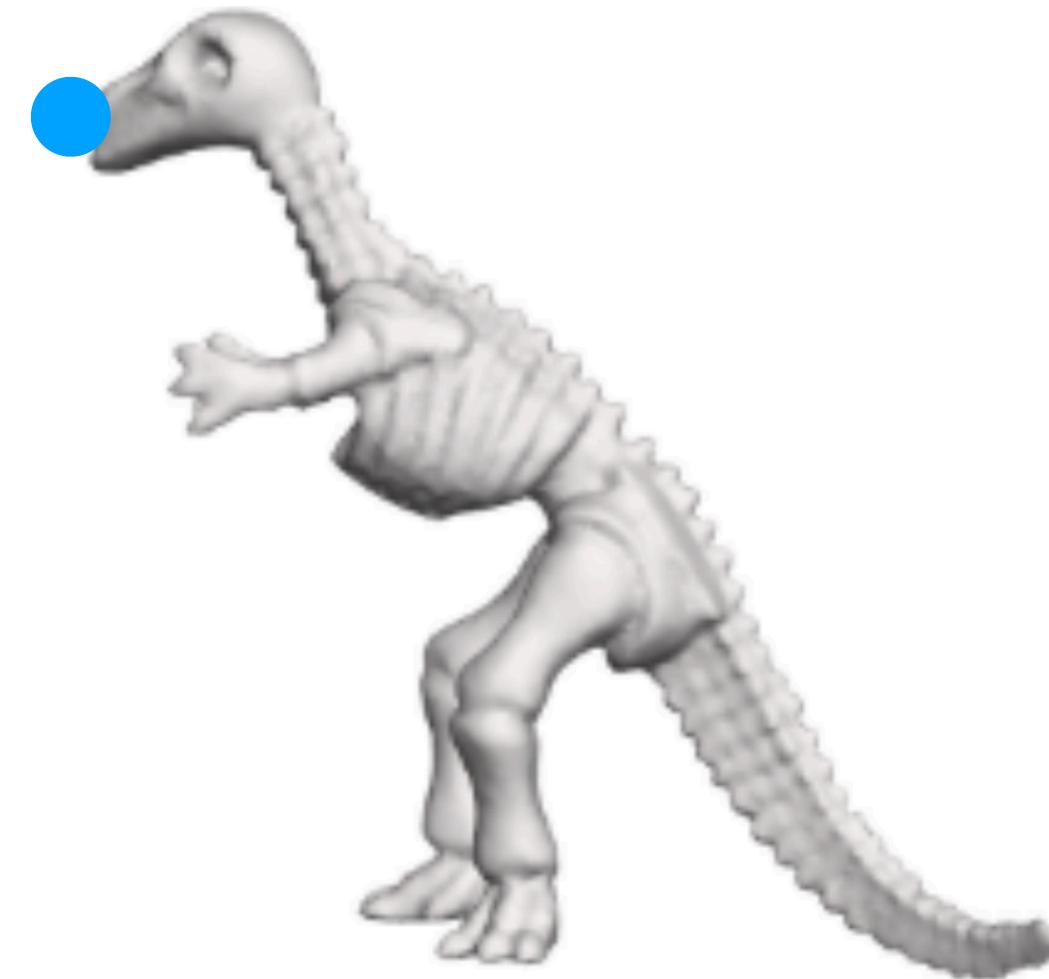
$$\hat{\mathbf{x}} = \mathbf{x}_0 - F(\mathbf{x}_0) \mathbf{n}(\mathbf{x}_0)$$

Step 2: Correspondences

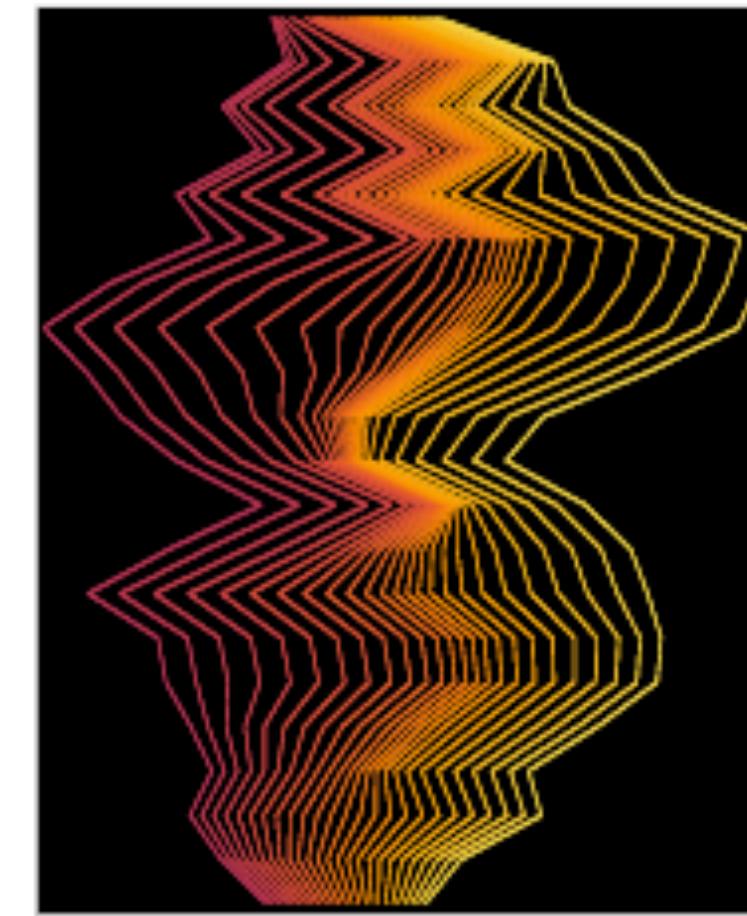


Step 2: Correspondences

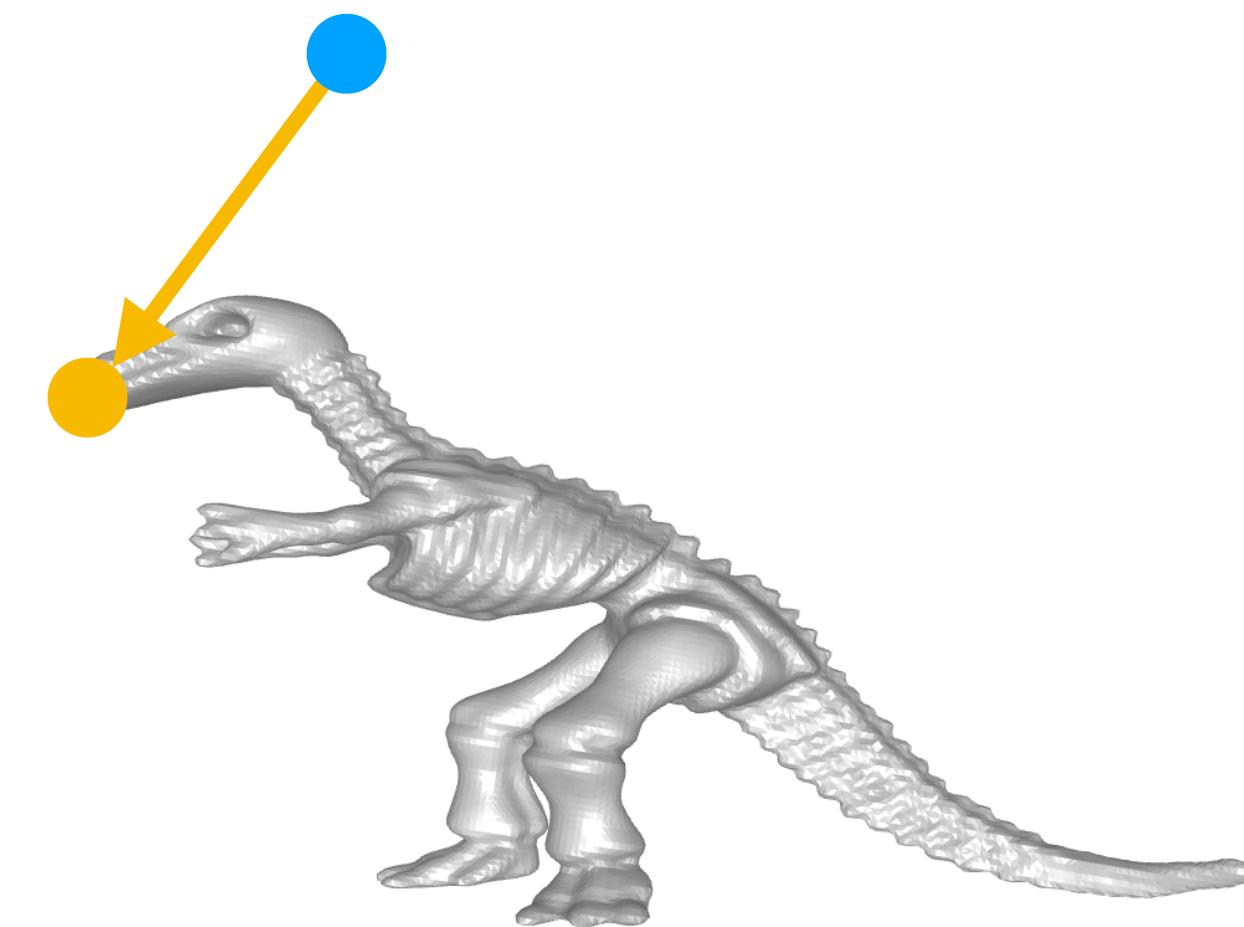
$$F(x)$$



$$D_\theta(x) = x + g_\theta(x)$$



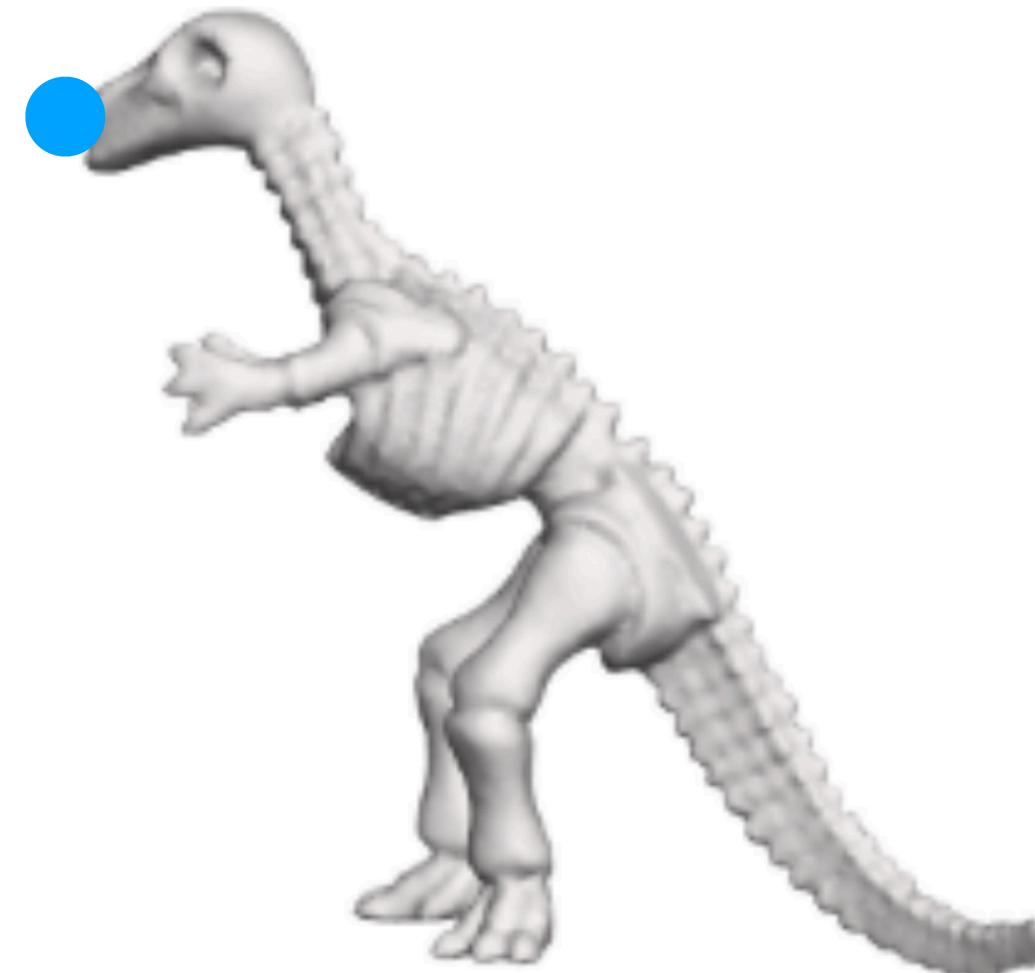
$$G_\theta(x) = F(D_\theta(x))$$



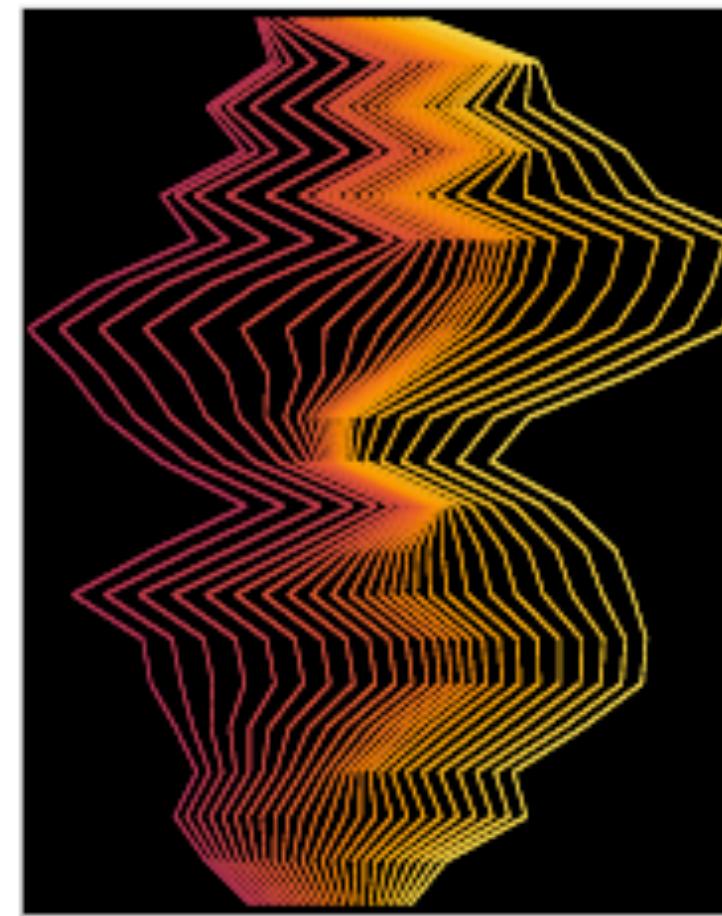
Invertible ResNet
(Behrmann et. al., 2019)

Step 2: Correspondences

$F(x)$

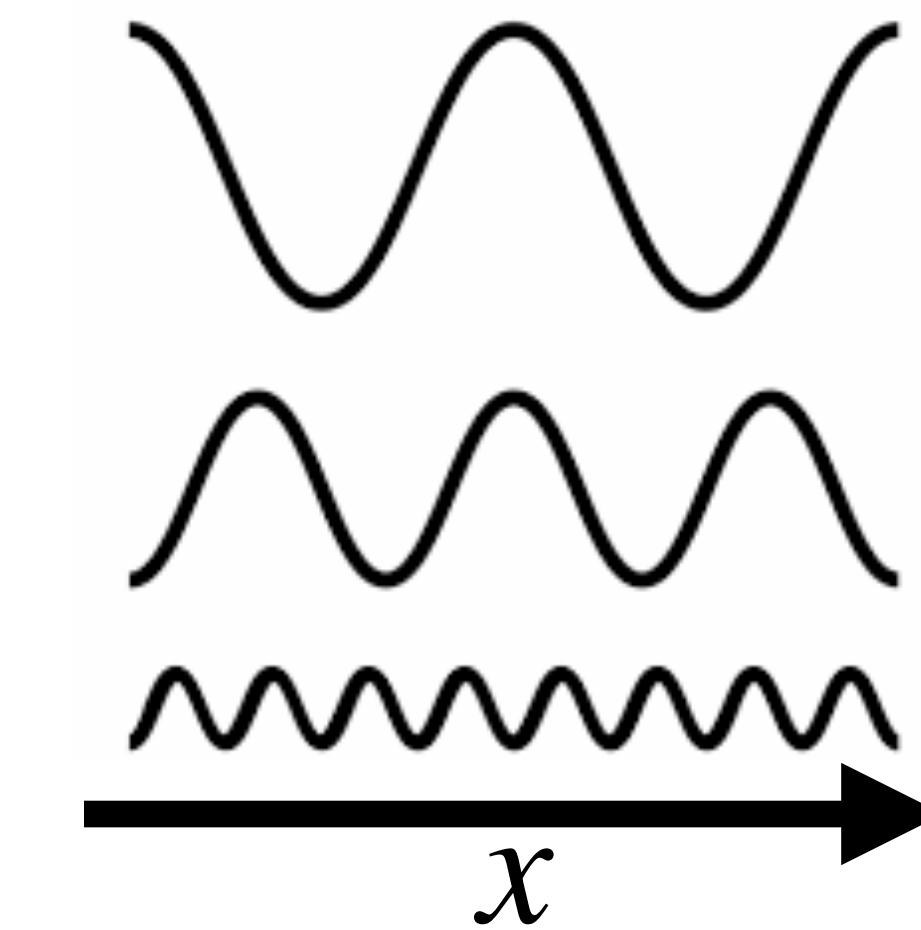


$$D_\theta(x) = x + g_\theta(x)$$



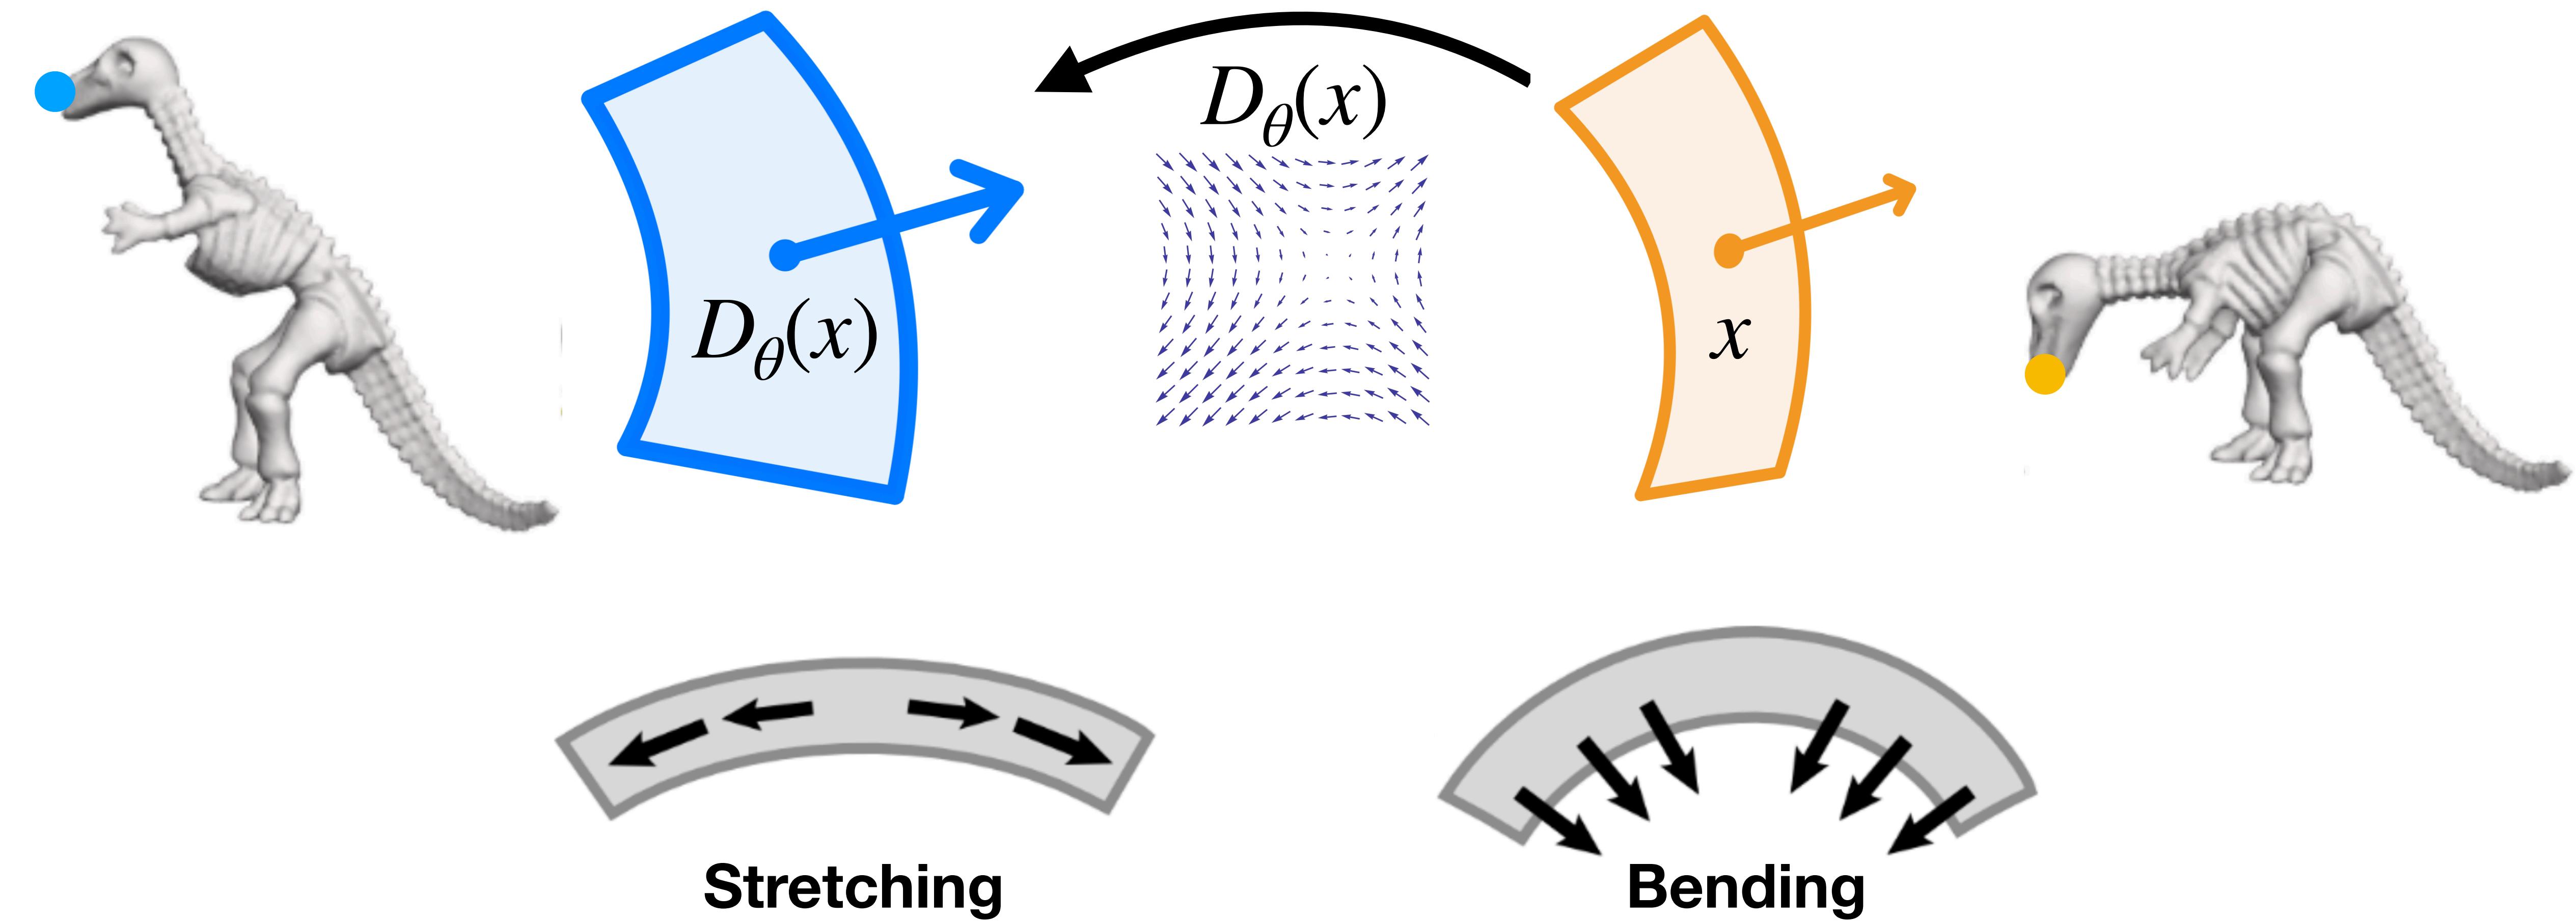
Invertible ResNet
(Behrman et. al., 2019)

$$G_\theta(x) = F(D_\theta(x))$$

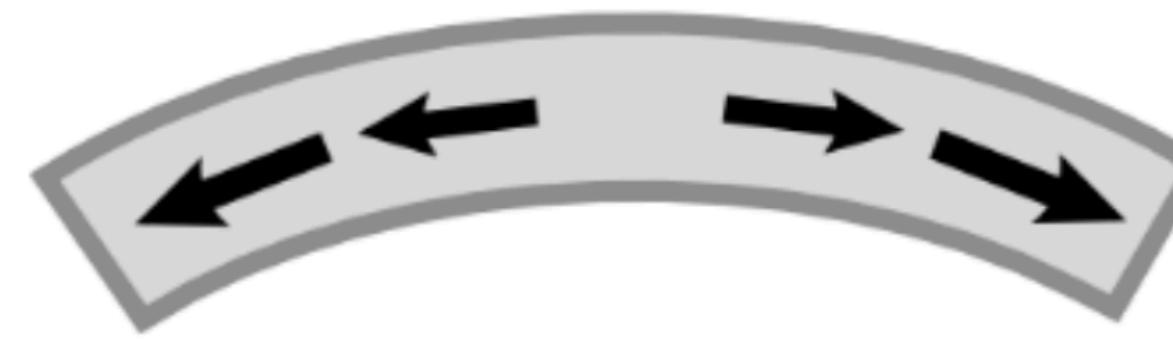


Positional Encoding

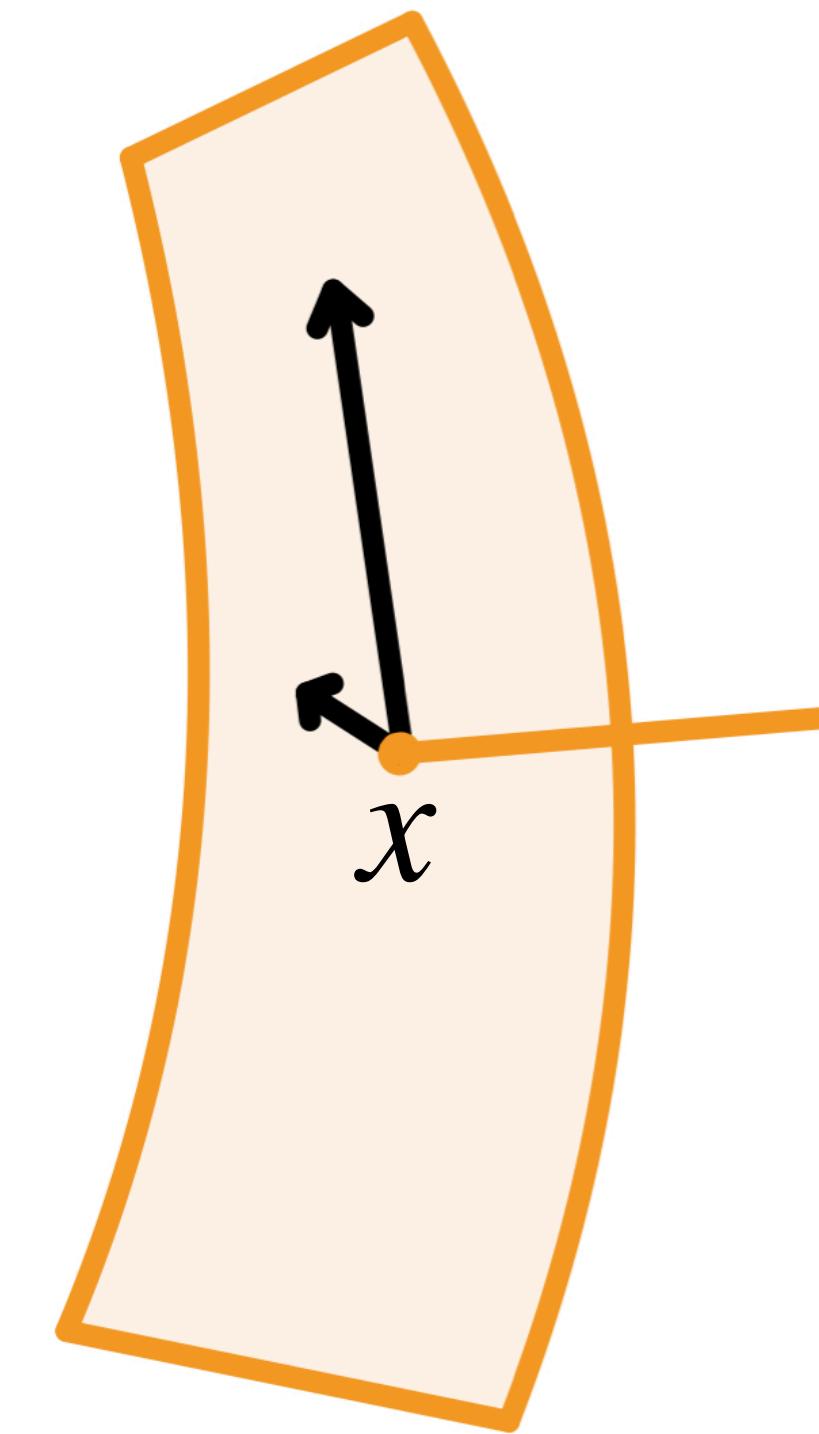
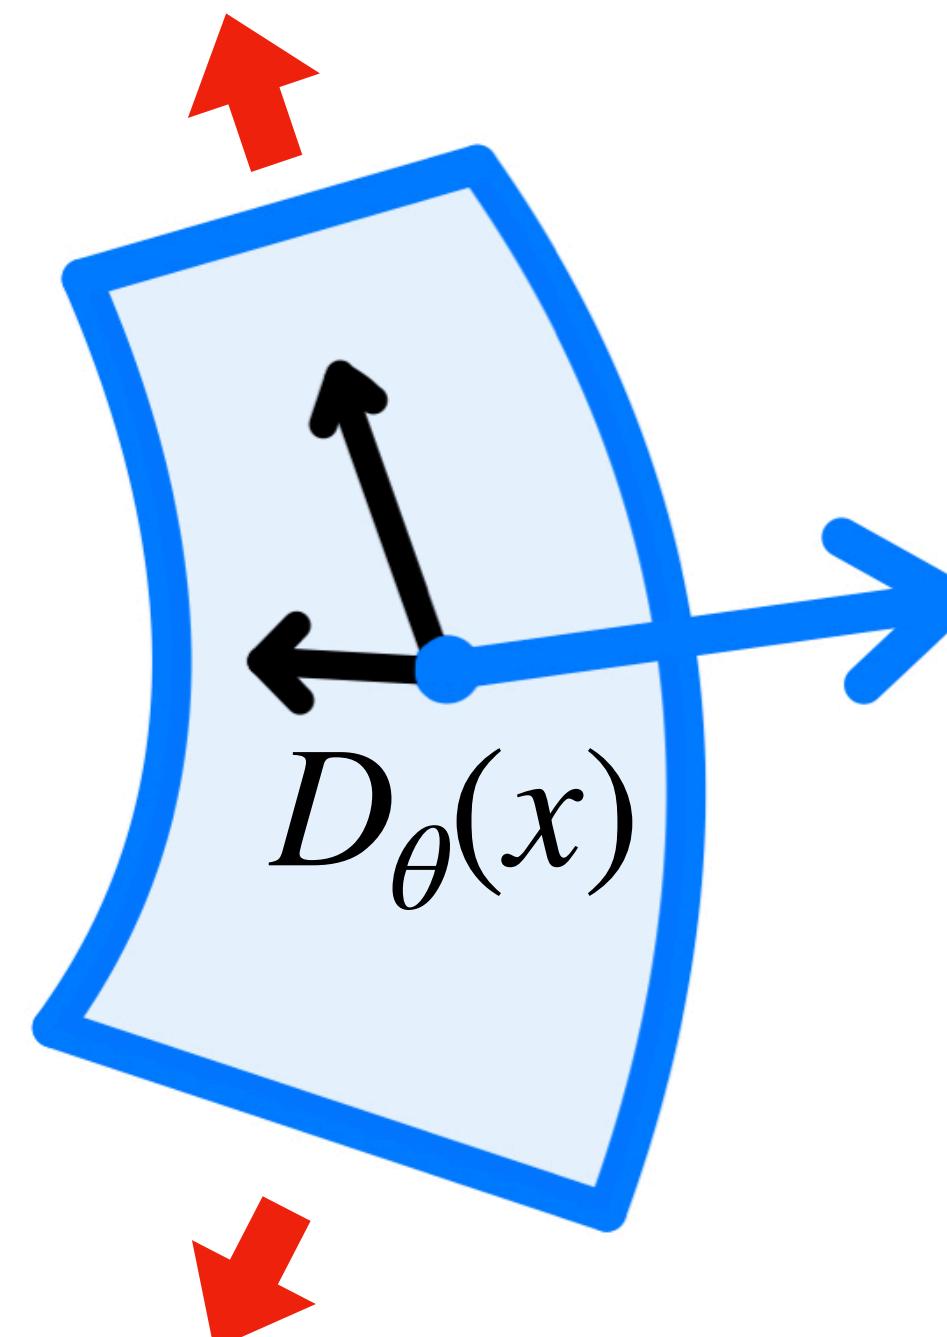
Step 3: Comparing



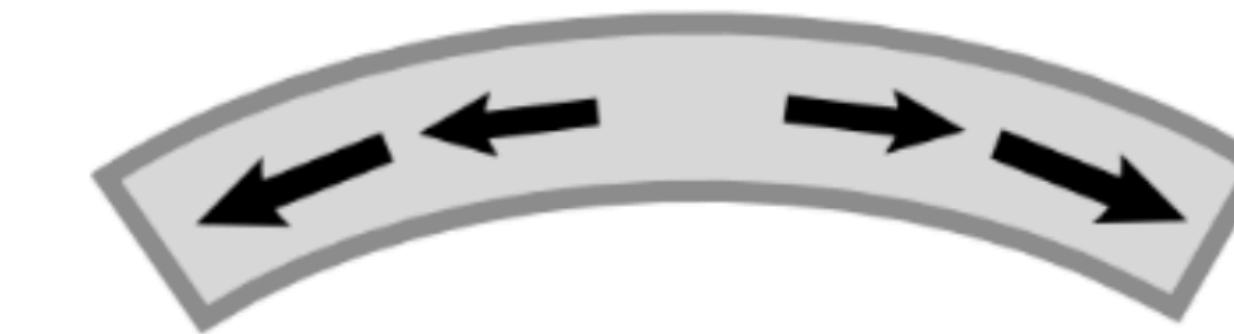
Step 3: Comparing



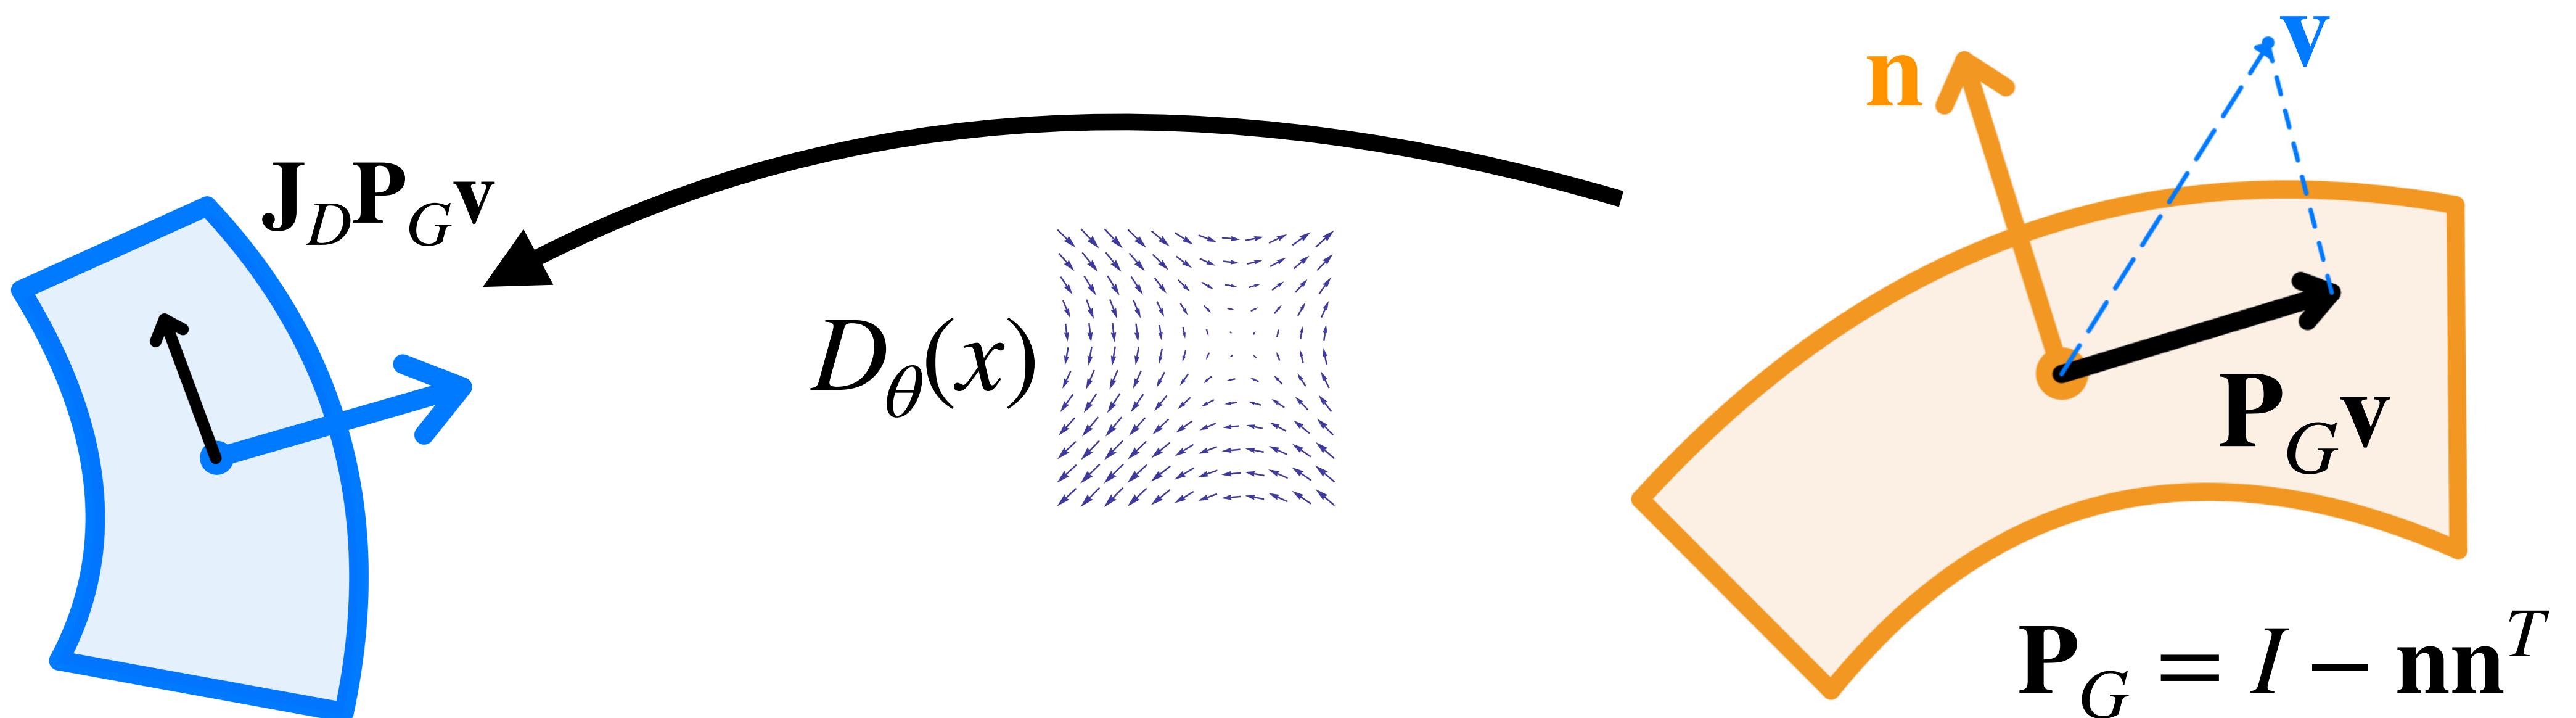
Stretch - change of tangent
dot-product



Step 3: Comparing

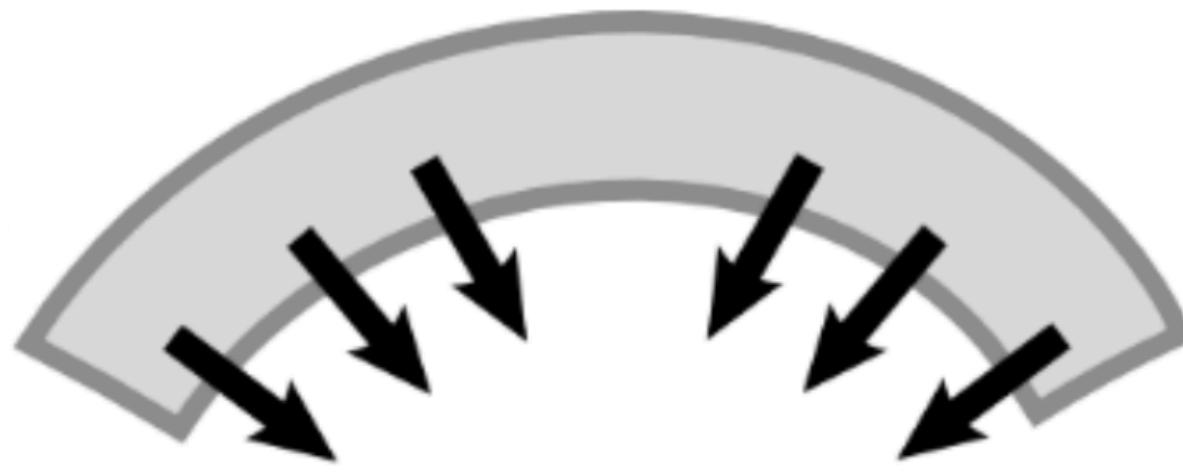


Stretch - change of tangent
dot-product

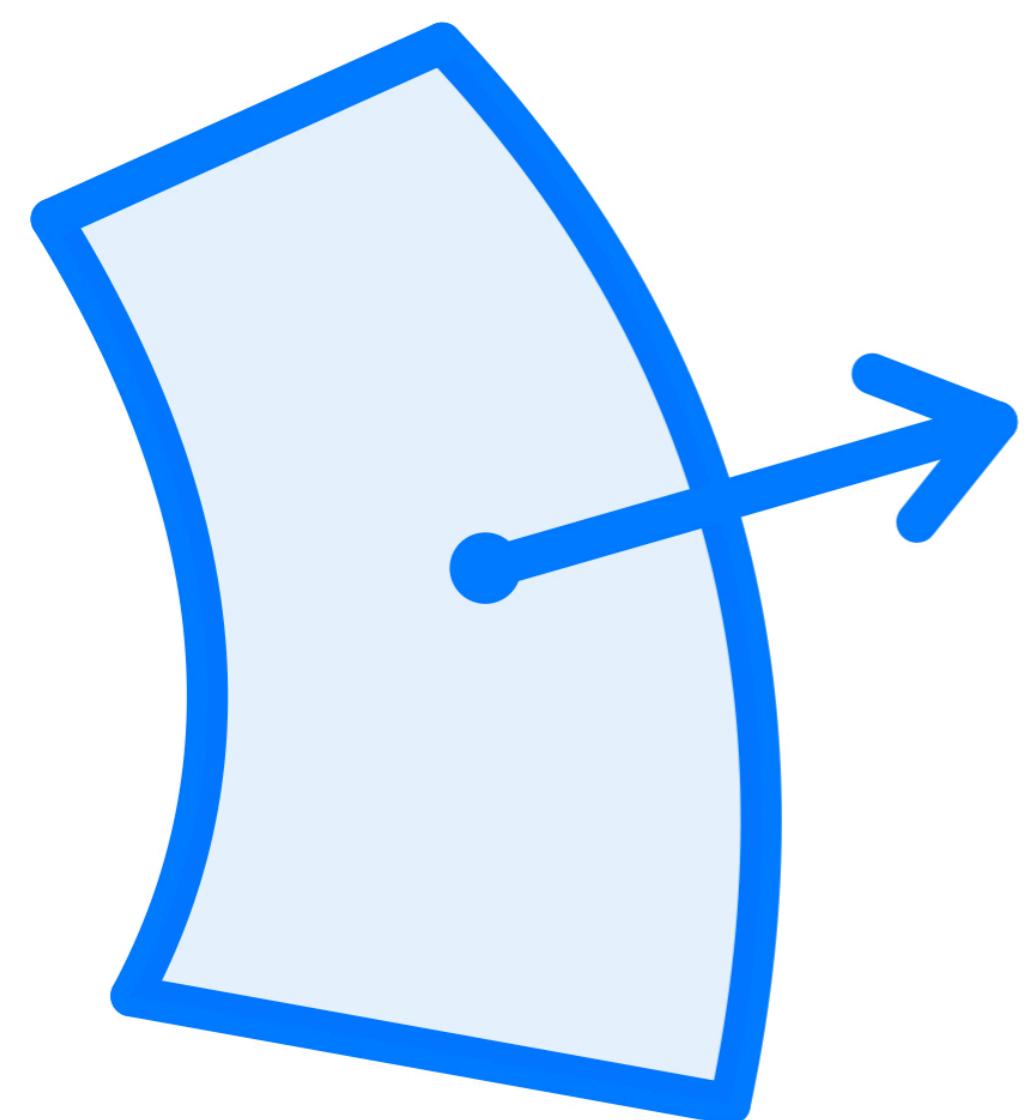


$$\mathcal{L}_s = \int \left\| \mathbf{P}_G^T (\mathbf{I} - \mathbf{J}_D^T \mathbf{J}_D) \mathbf{P}_G \right\|_F dx$$

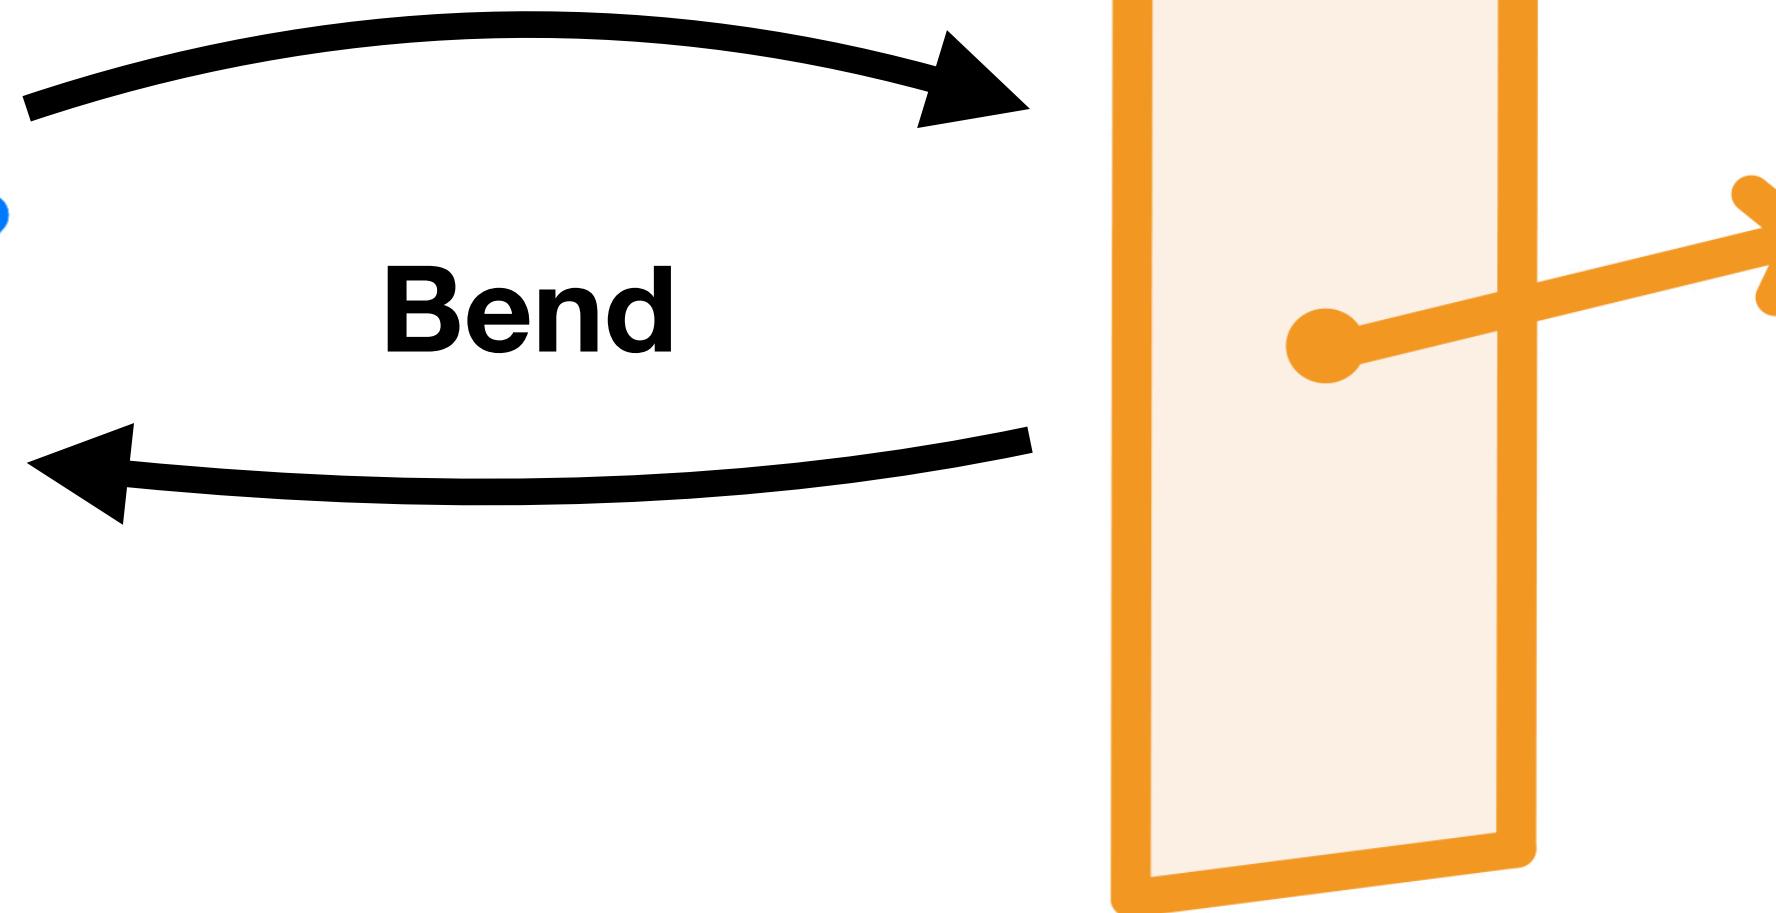
Step 3: Comparing



Bending - change of curvature



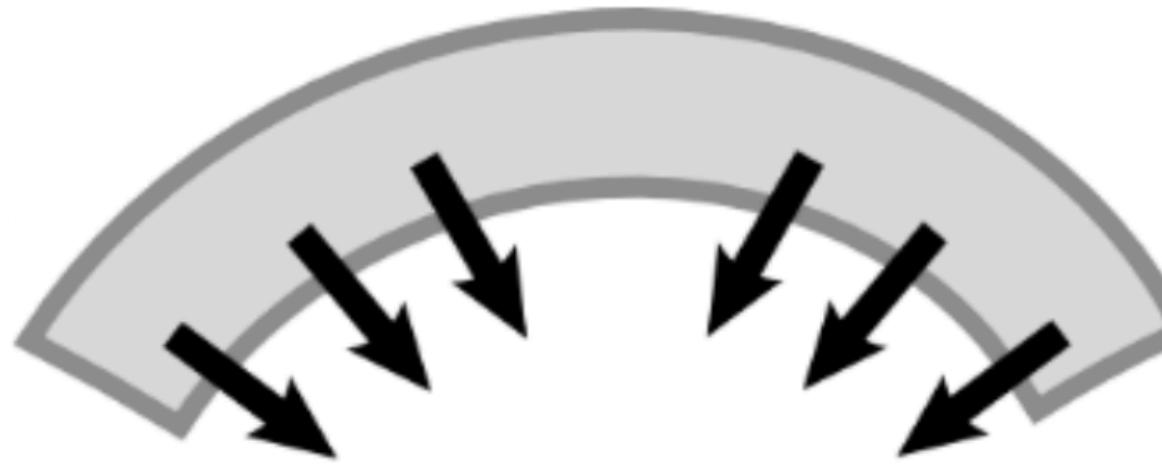
More curved



Less curved

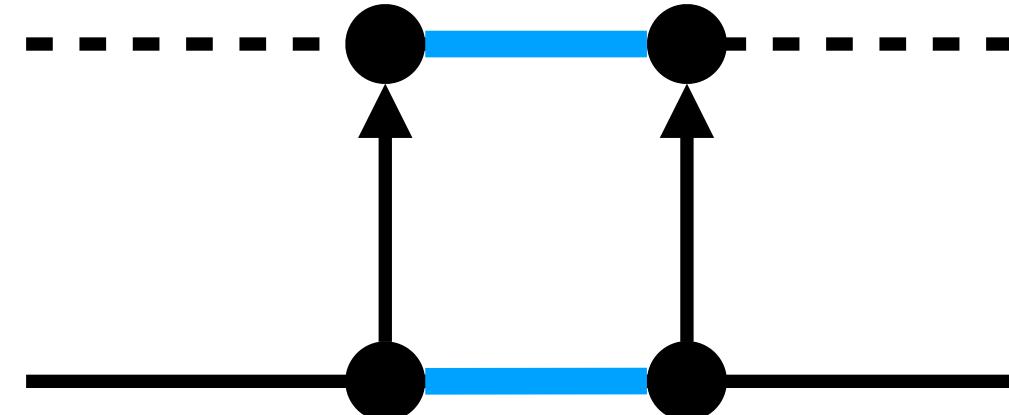
Bend

Step 3: Comparing

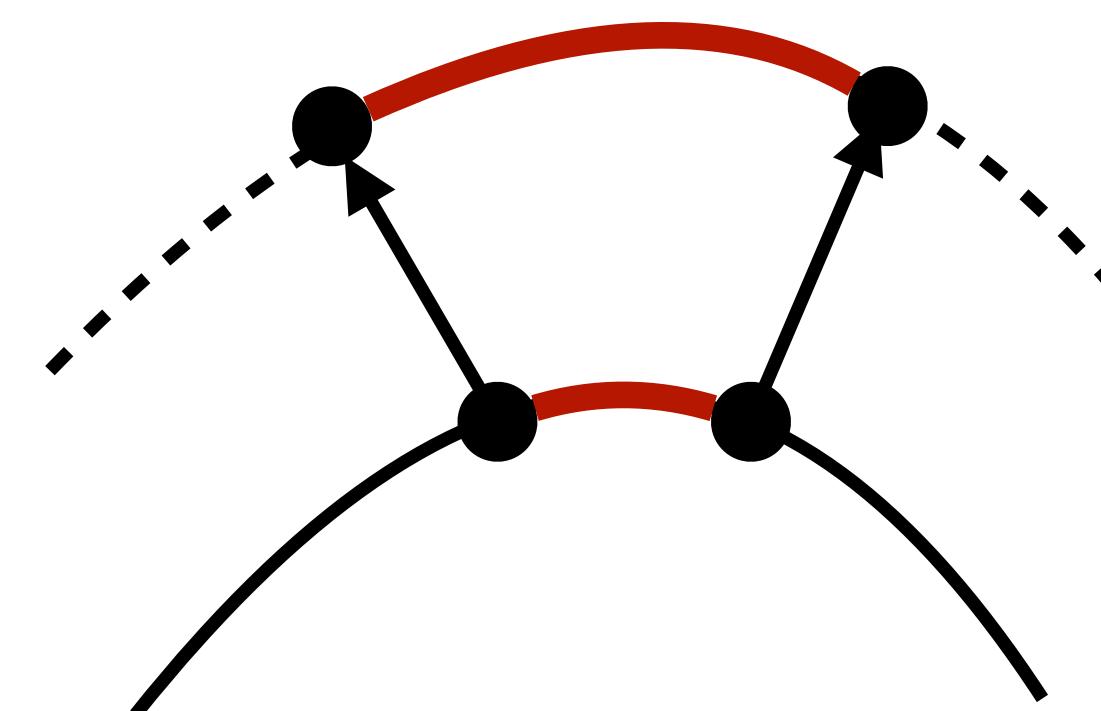


Bending - change of **curvature**

Curvature - change along normal direction



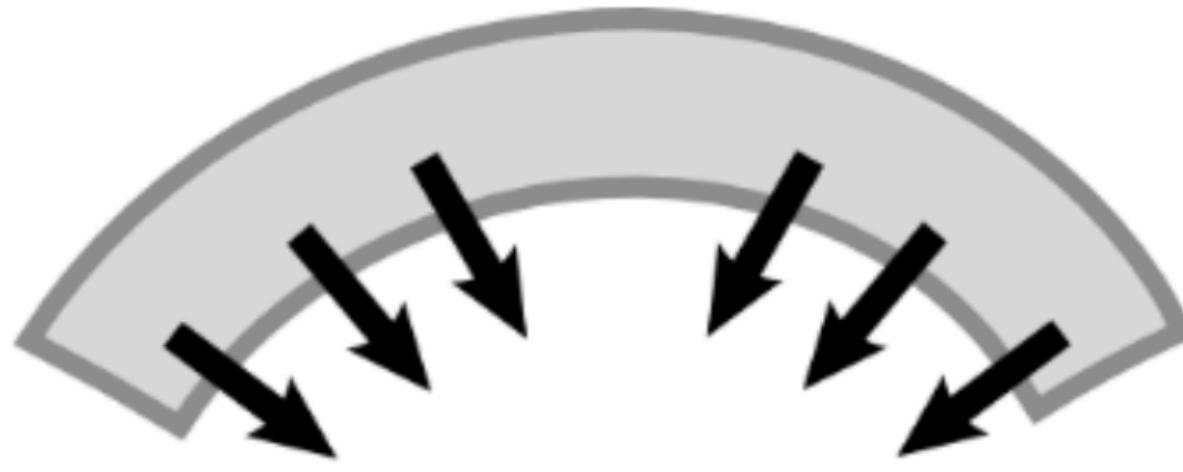
Low curvature
Little change



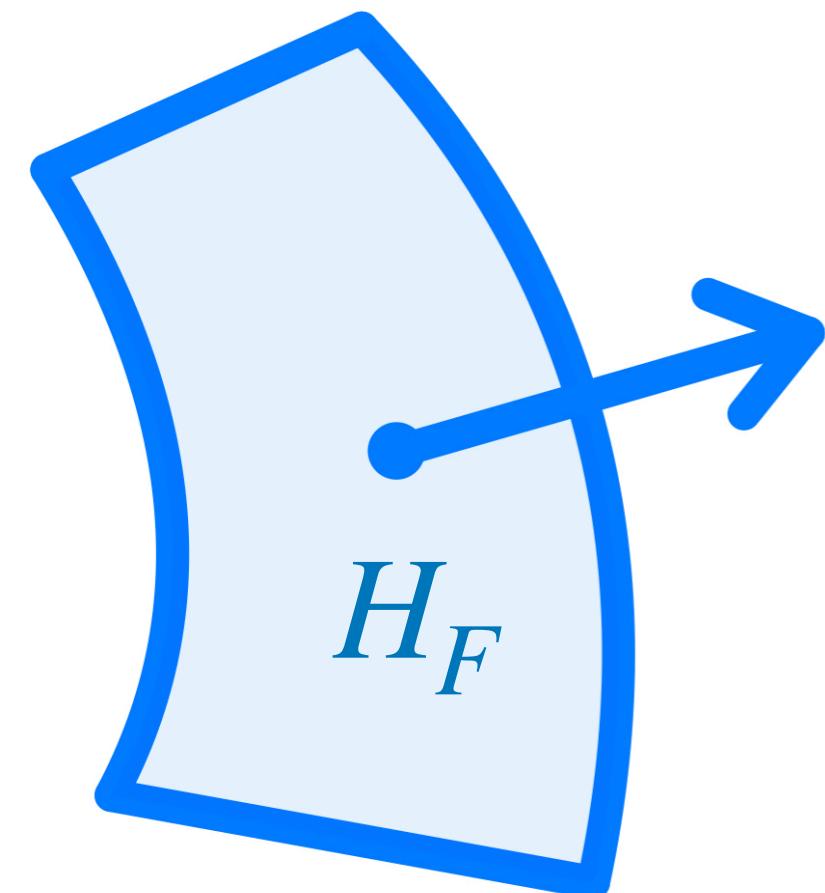
High curvature
Large change

$$H_f = \begin{bmatrix} f_{xx} & f_{xy} & f_{xz} \\ f_{yx} & f_{yy} & f_{yz} \\ f_{zx} & f_{zy} & f_{zz} \end{bmatrix}$$

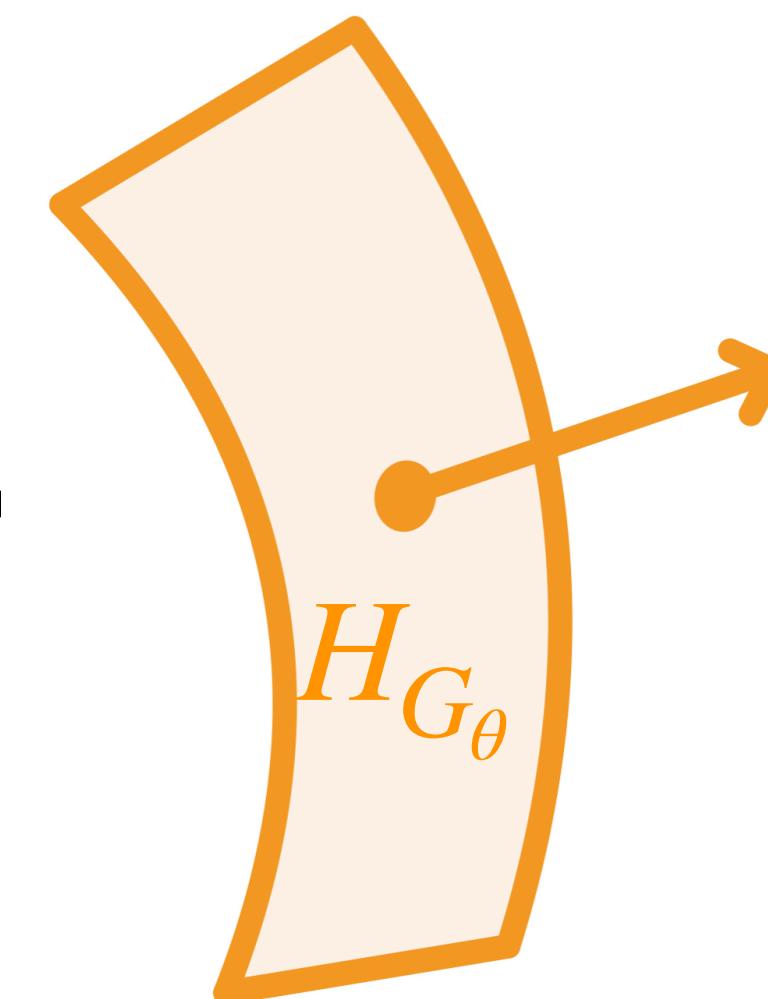
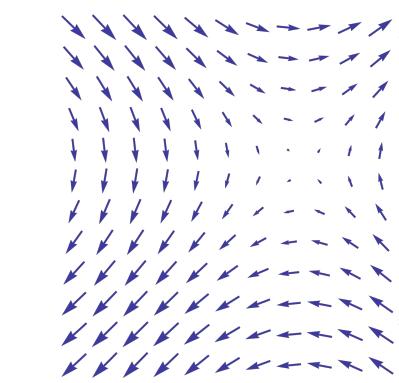
Step 3: Comparing



Bending - change of curvature



$$D_\theta(x)$$



$$\mathcal{L}_b = \int_x \left\| \mathbf{P}_G^T (\mathbf{H}_G - \mathbf{J}_D^T \mathbf{H}_F \mathbf{J}_D) \mathbf{P}_G \right\|_F dx$$

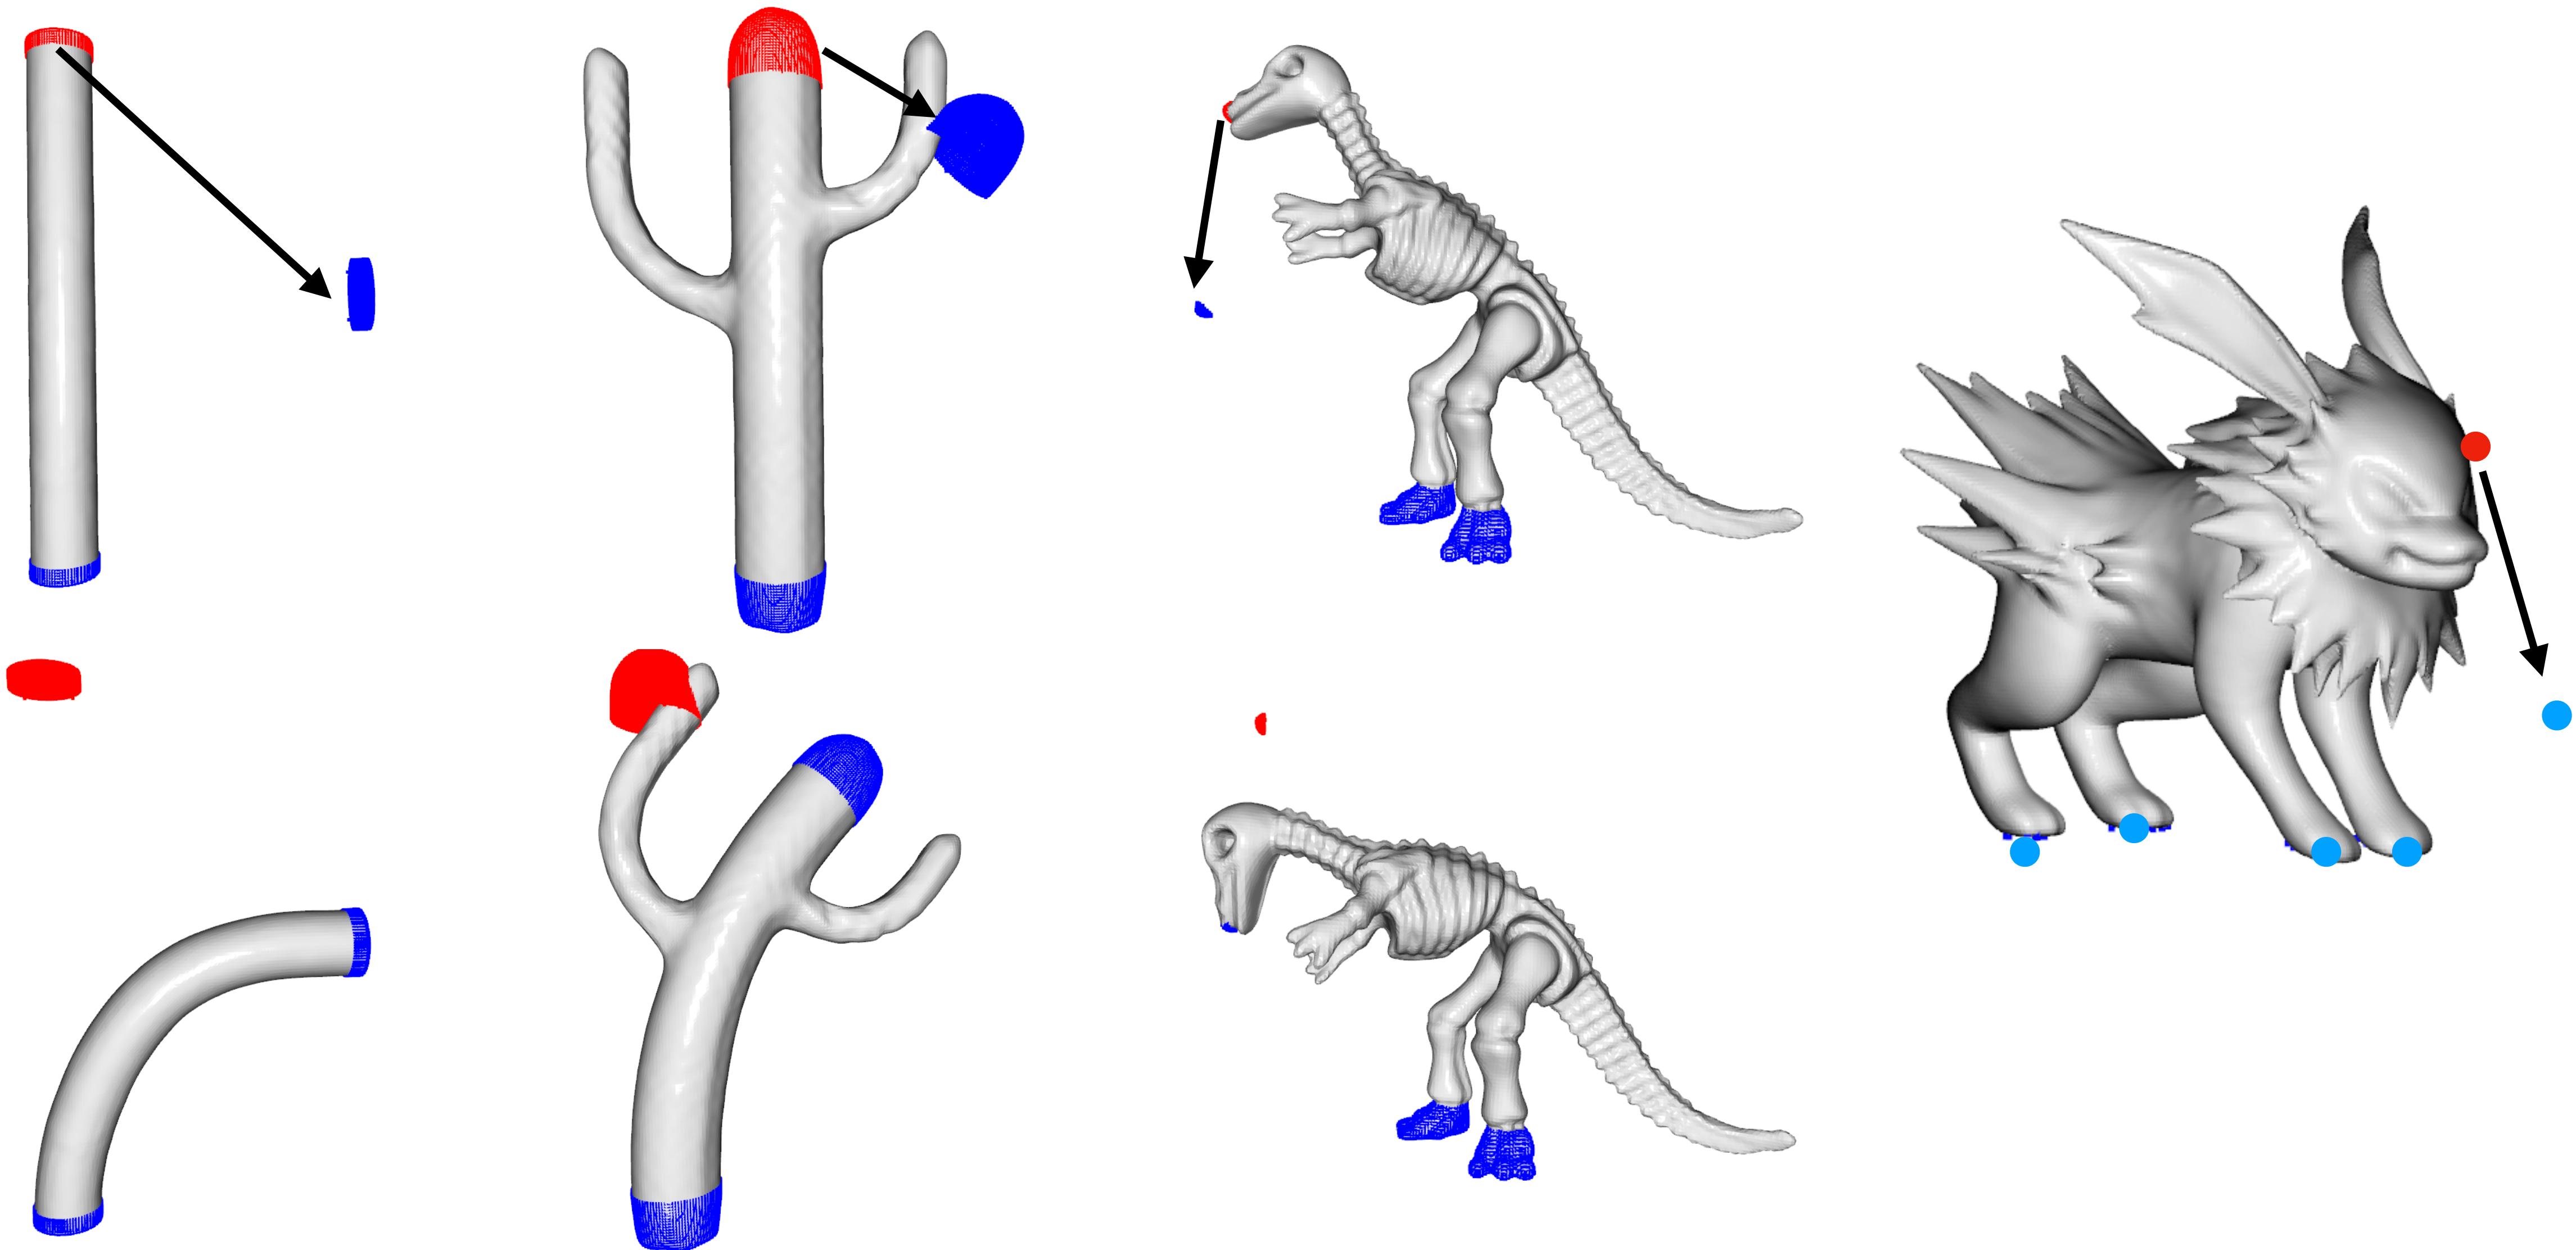
Step 3: Final Objective

$$\mathcal{L}_s = \int \left\| \mathbf{P}_G^T (\mathbf{I} - \mathbf{J}_D^T \mathbf{J}_D) \mathbf{P}_G \right\|_F dx$$

$$\mathcal{L}_b = \int_x \left\| \mathbf{P}_G^T (\mathbf{H}_G - \mathbf{J}_D^T \mathbf{H}_F \mathbf{J}_D) \mathbf{P}_G \right\|_F dx$$

$$\min_{\theta} k_s \mathcal{L}_s + k_b \mathcal{L}_b + k_c \sum_i \max(\tau, |D_{\theta}(t_i) - h_i|)$$

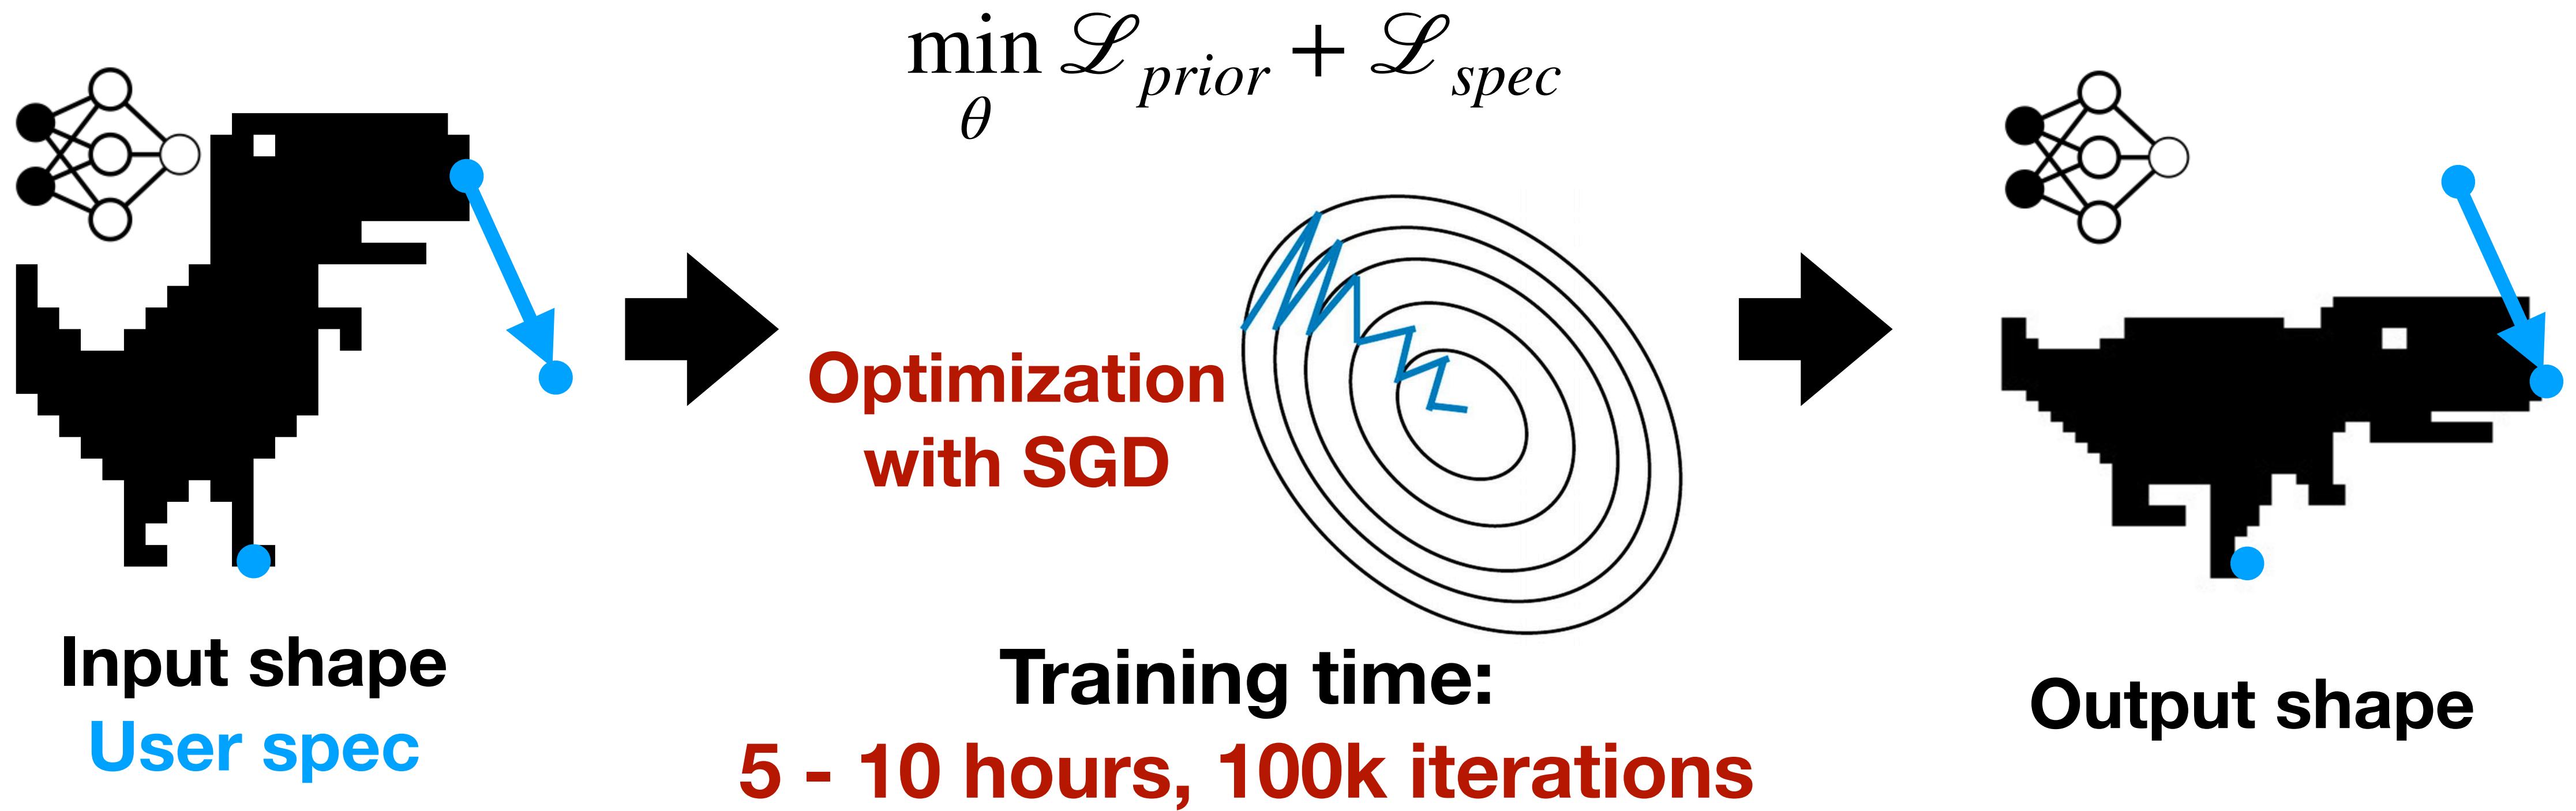
Deformation Results



Limitations & Future Works

- Speed
- Quality of the derivatives
- Field property preservation
- Physical Simulation

Speed

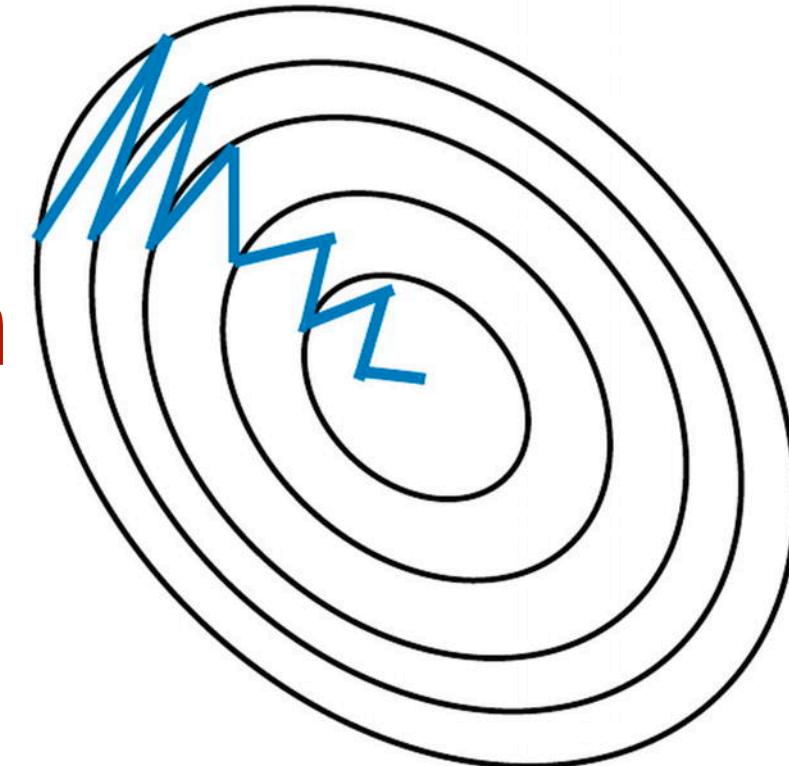


Speed

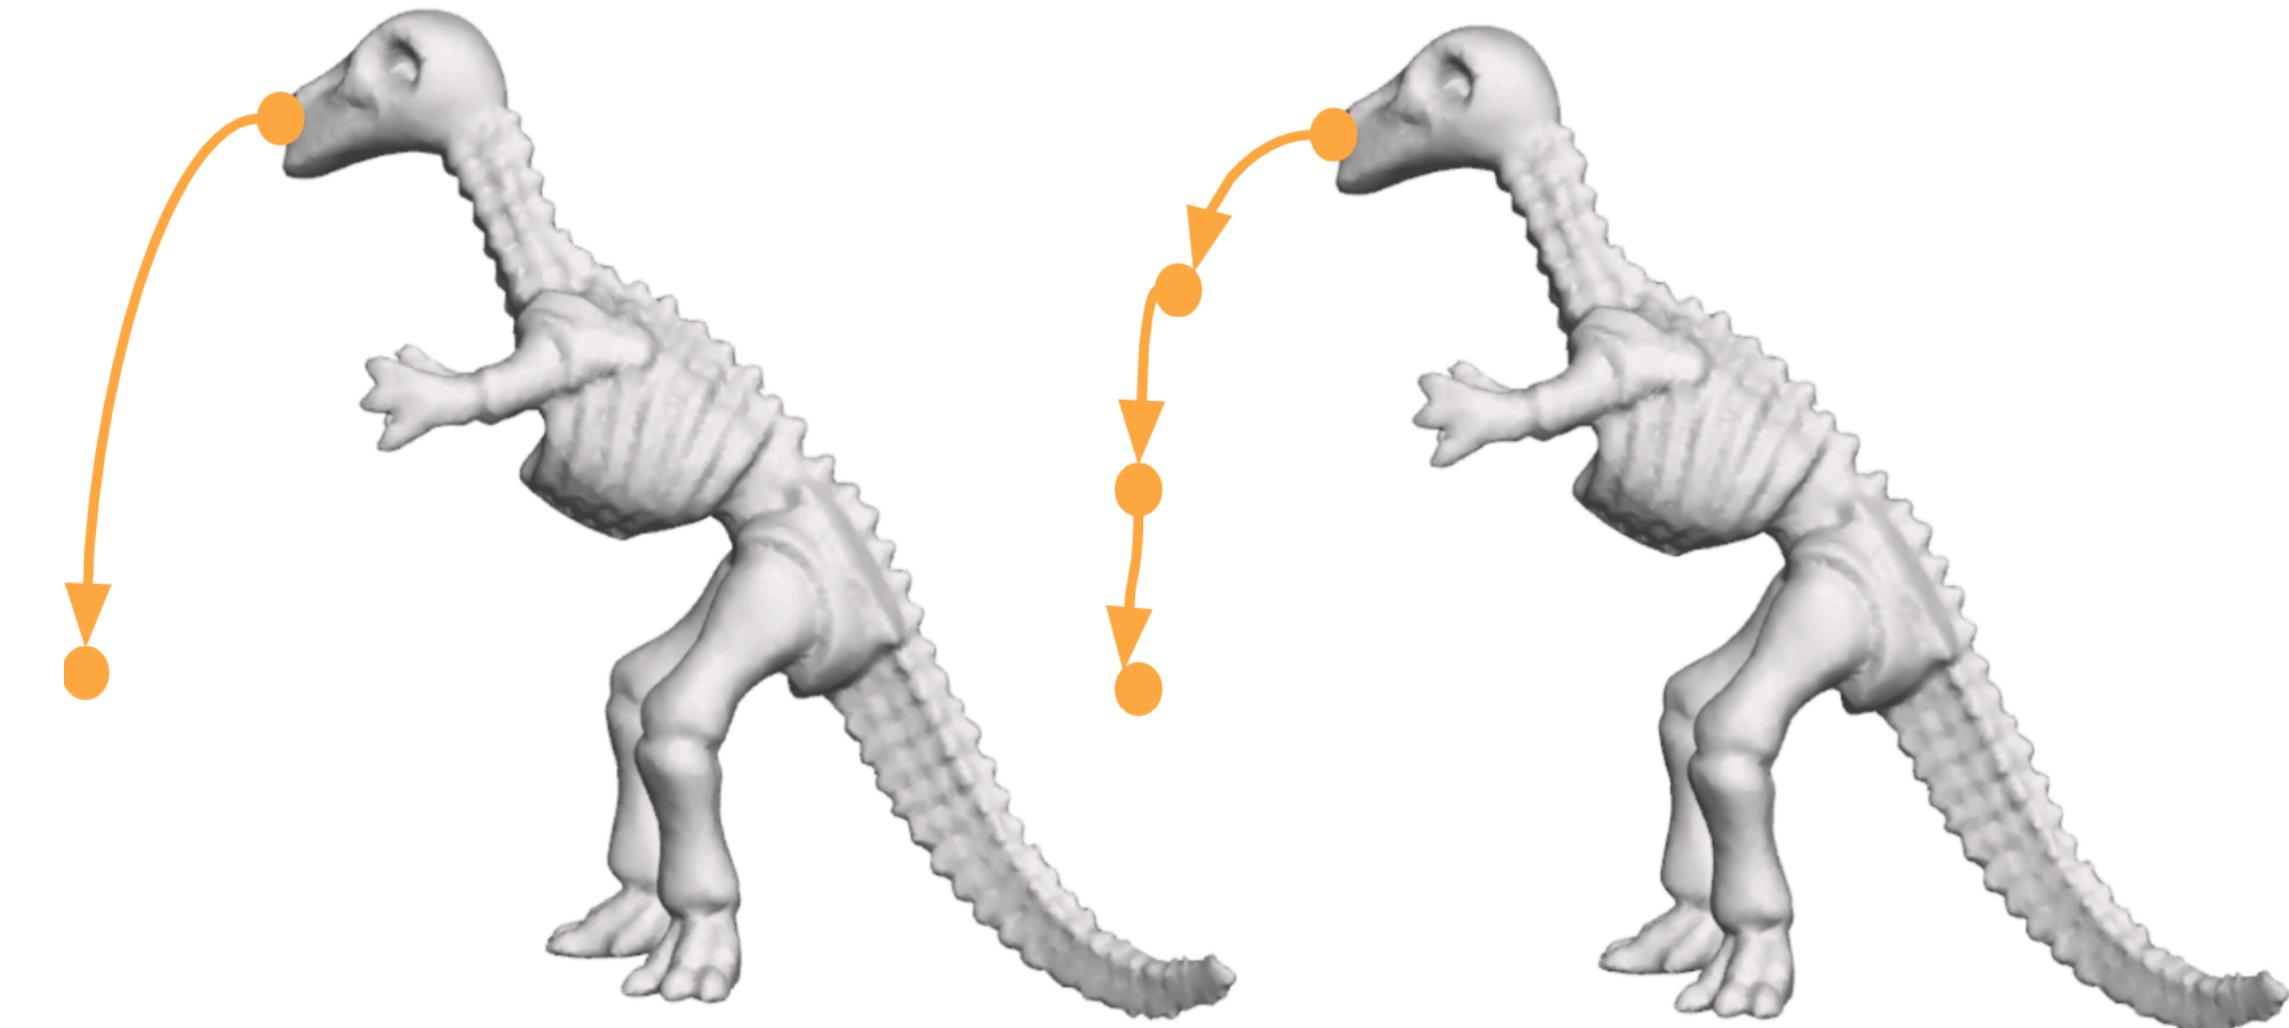
IDEA 1 : DIVIDE AND CONQUER

$$\min_{\theta} \mathcal{L}_{prior} + \mathcal{L}_{spec}$$

**Optimization
with SGD**



**Training time:
5 - 10 hours, 100k iterations**

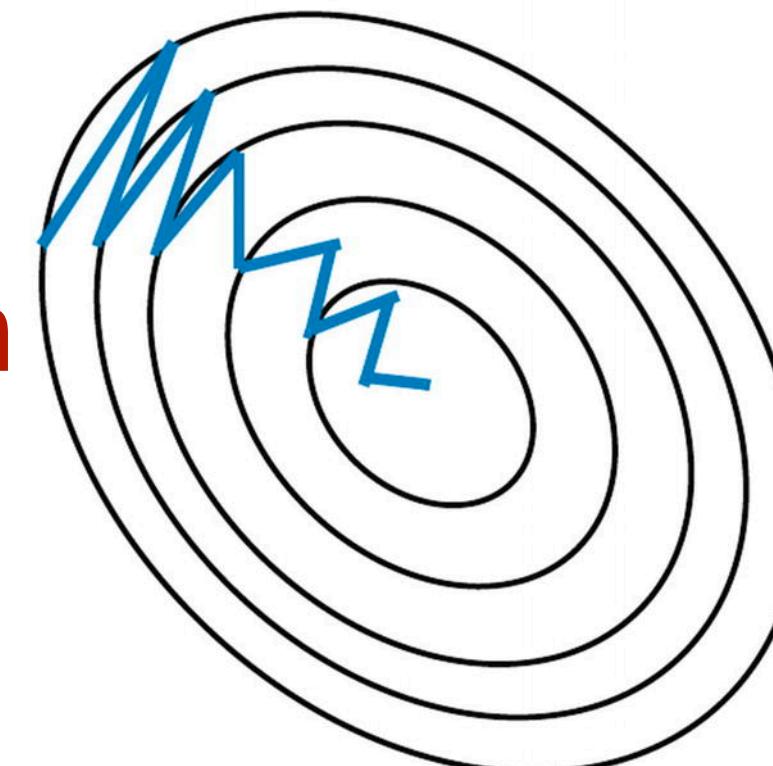


**Observation: if the deformation is
small, then it's easier to converge/
optimize**

Speed

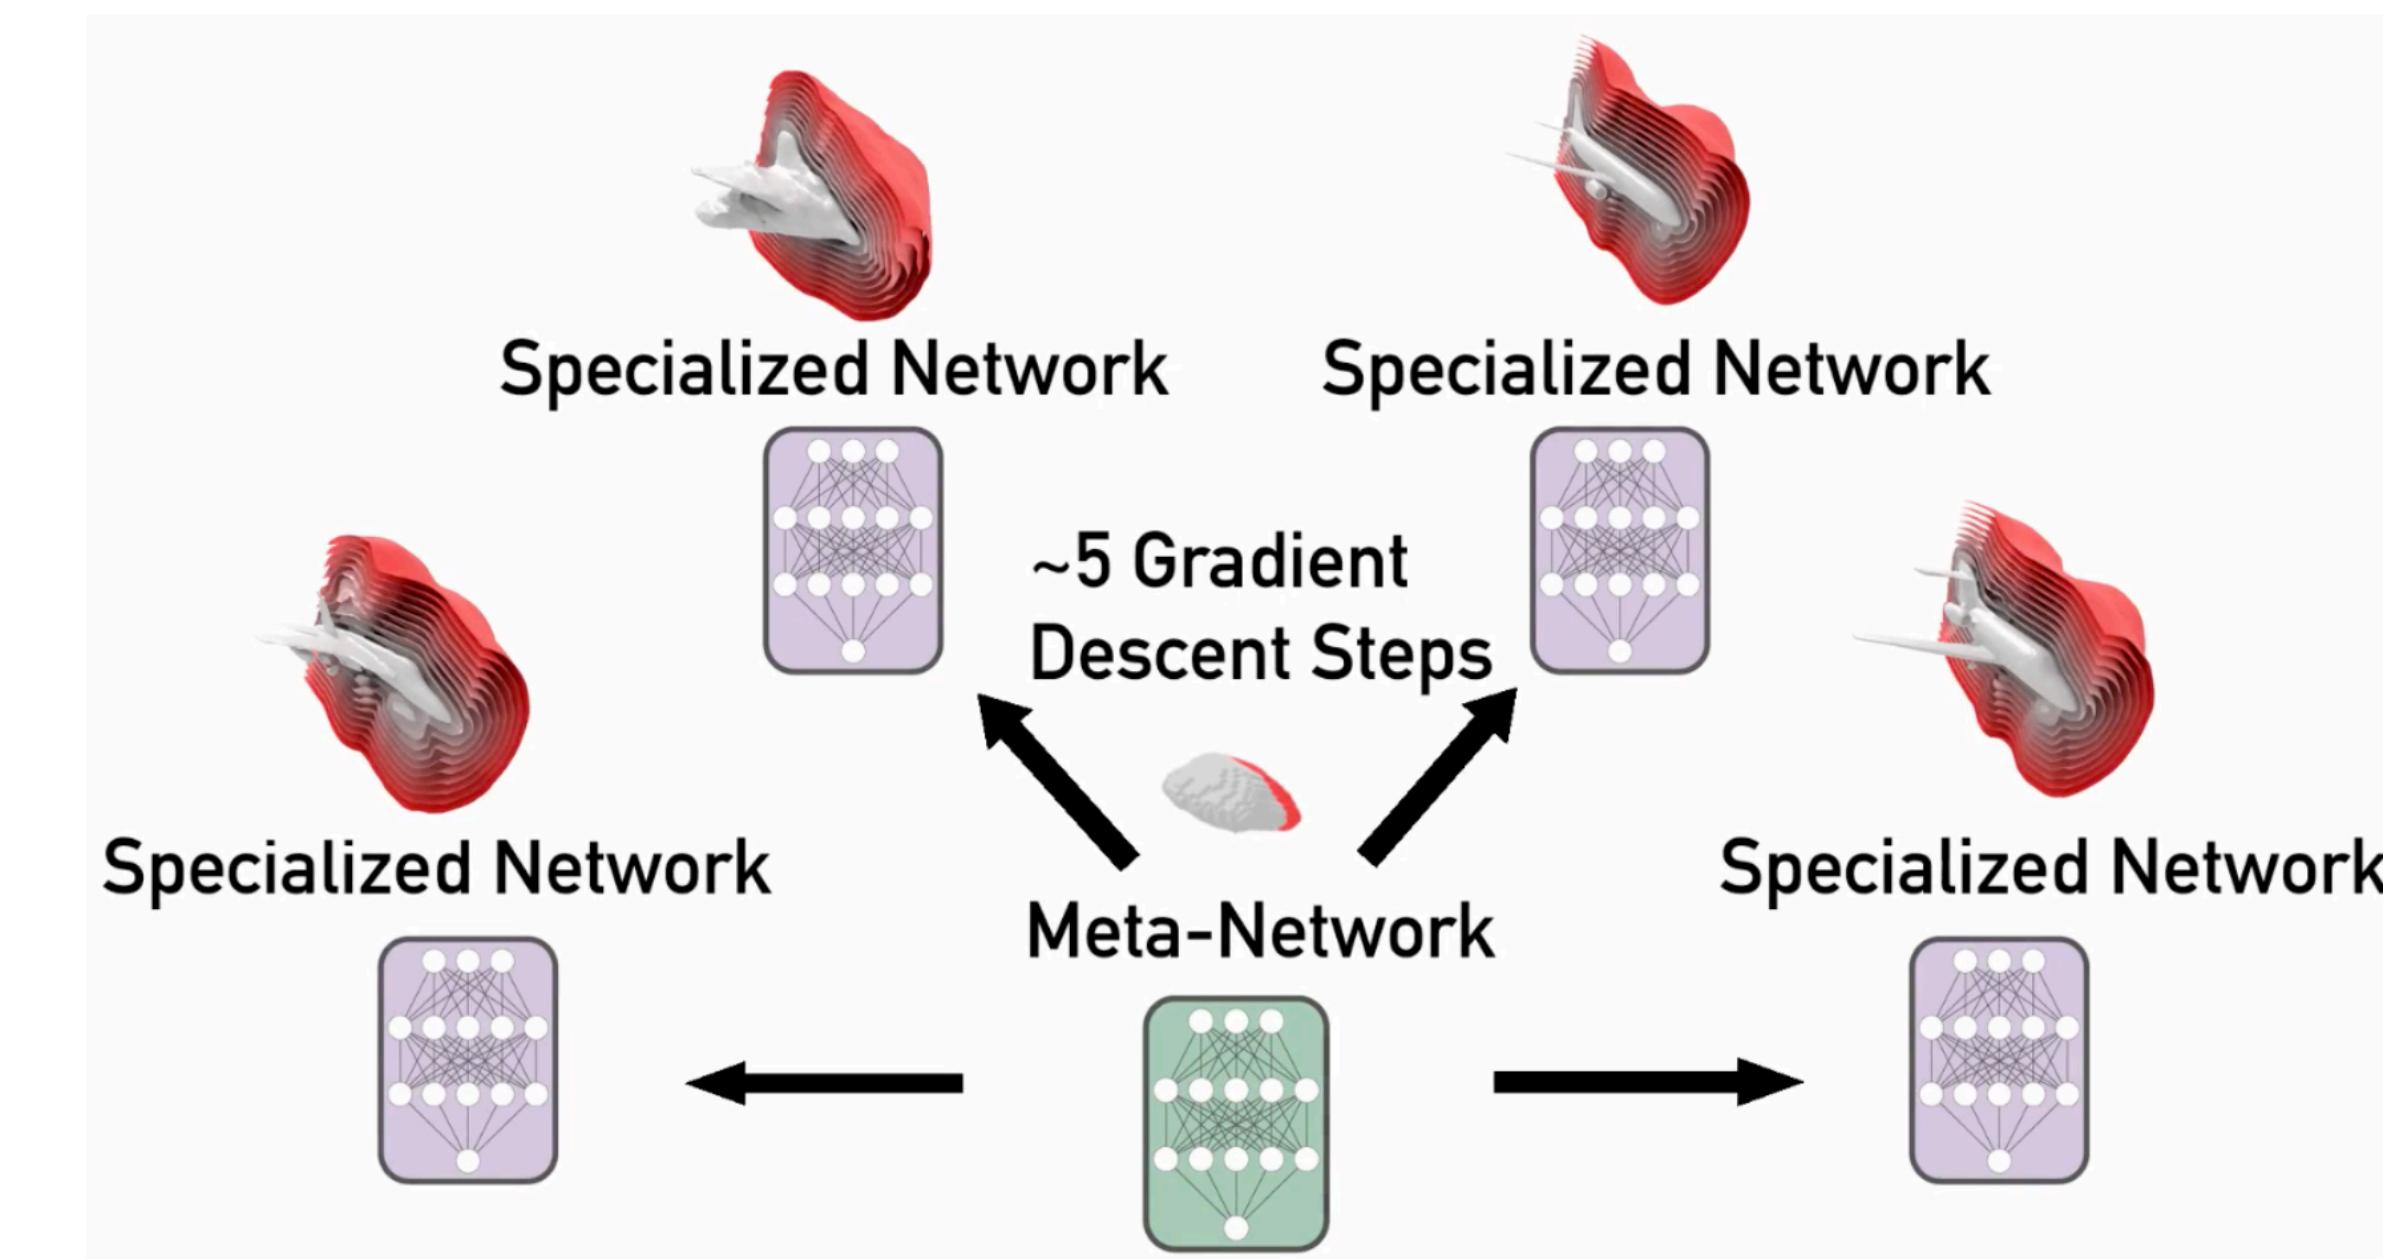
$$\min_{\theta} \mathcal{L}_{prior} + \mathcal{L}_{spec}$$

**Optimization
with SGD**



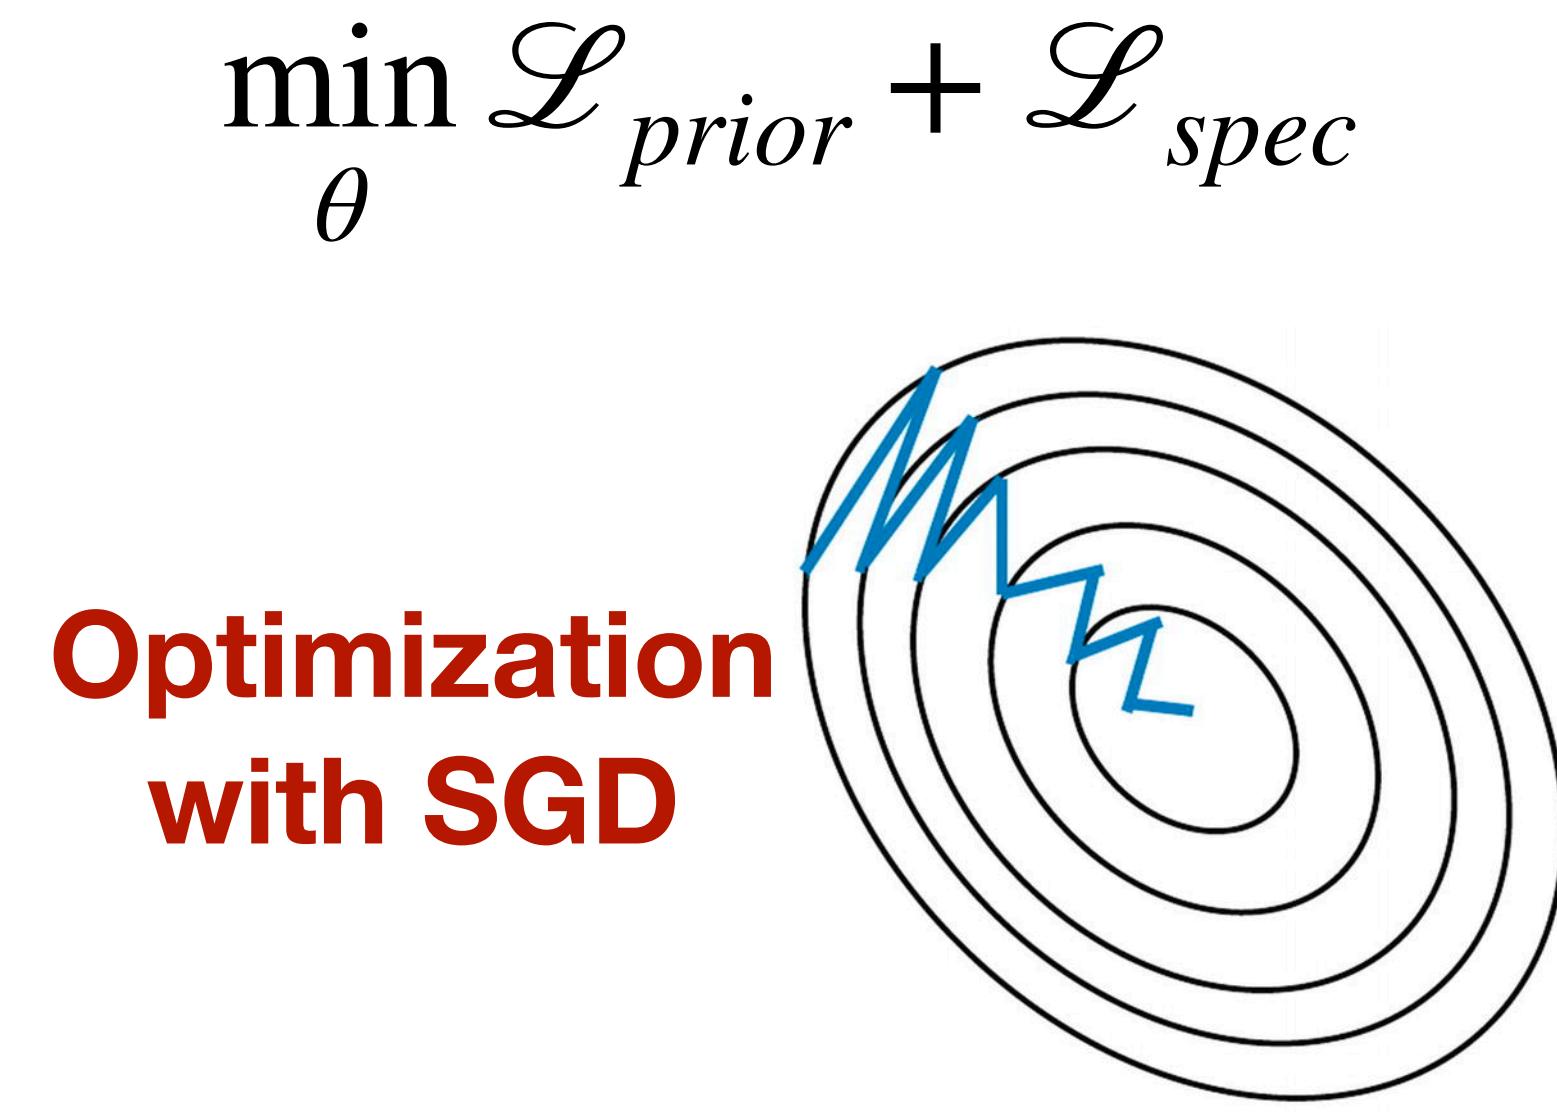
**Training time:
5 - 10 hours, 100k iterations**

IDEA 2 : Don't start from scratch



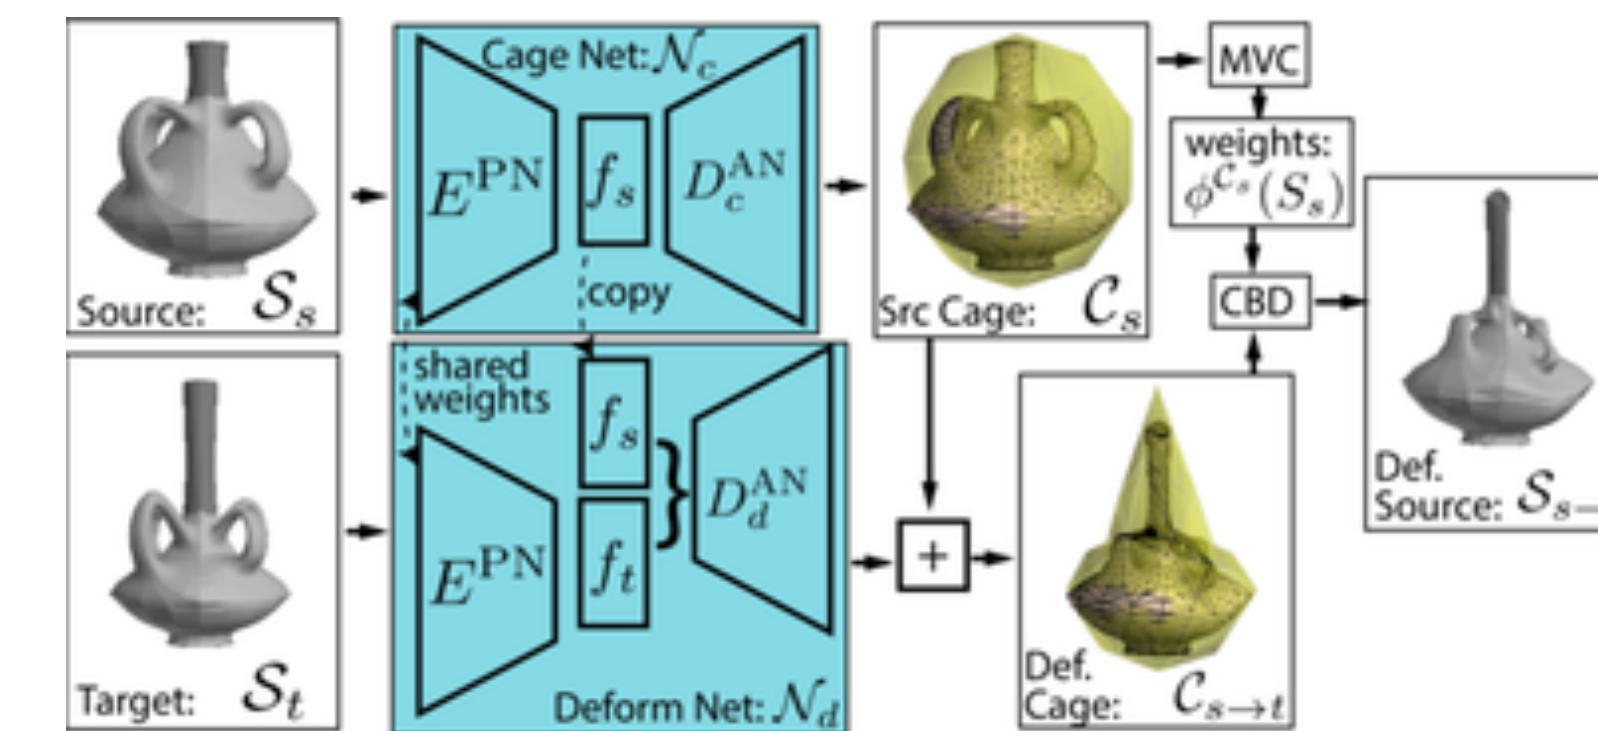
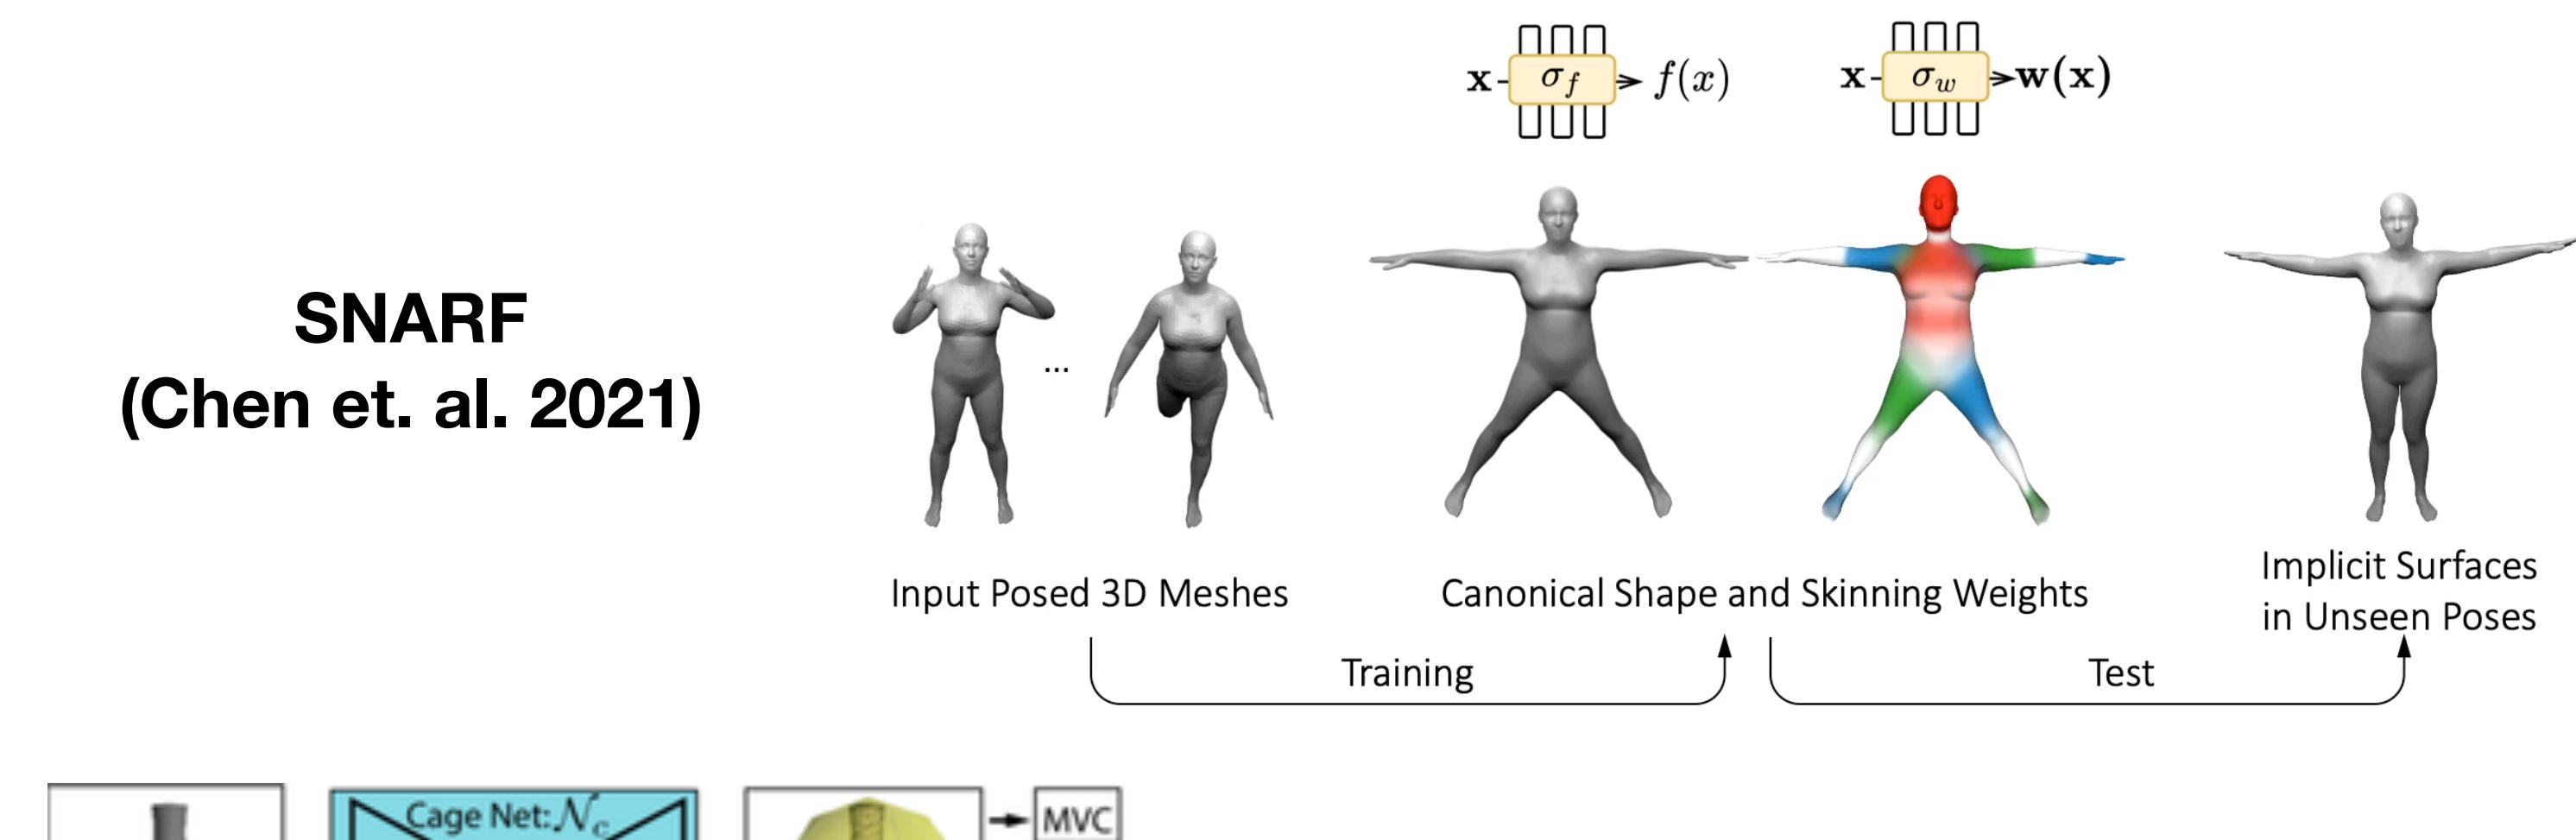
**Observation: we can make a good
guess of the output deformation field
given the user edit
MetaSDF (Sitzmann et. al. 2021)**

Speed



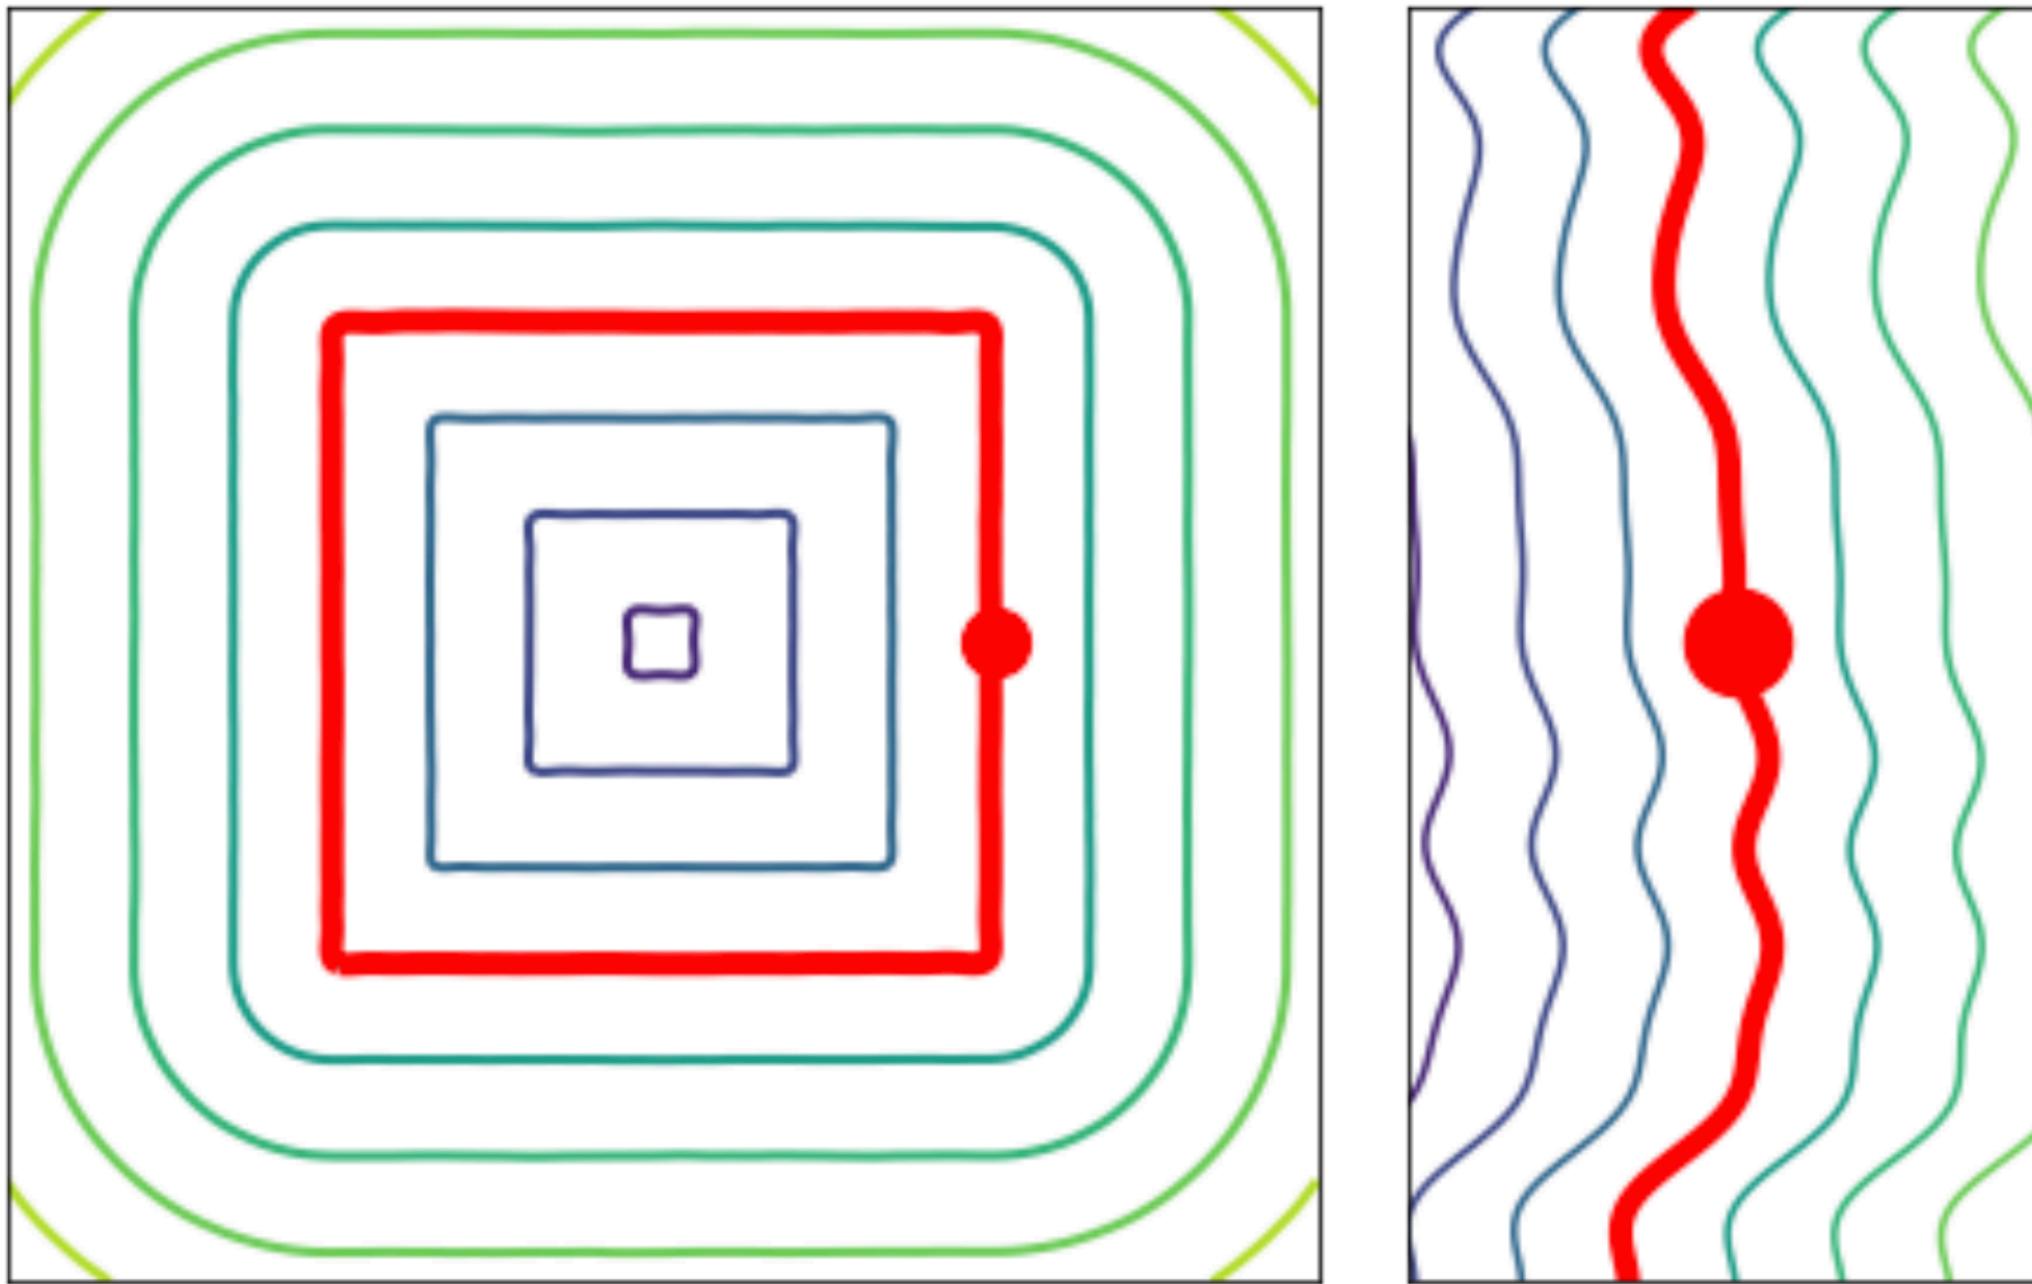
Training time:
5 - 10 hours, 100k iterations

IDEA 3: narrow down the deformation space

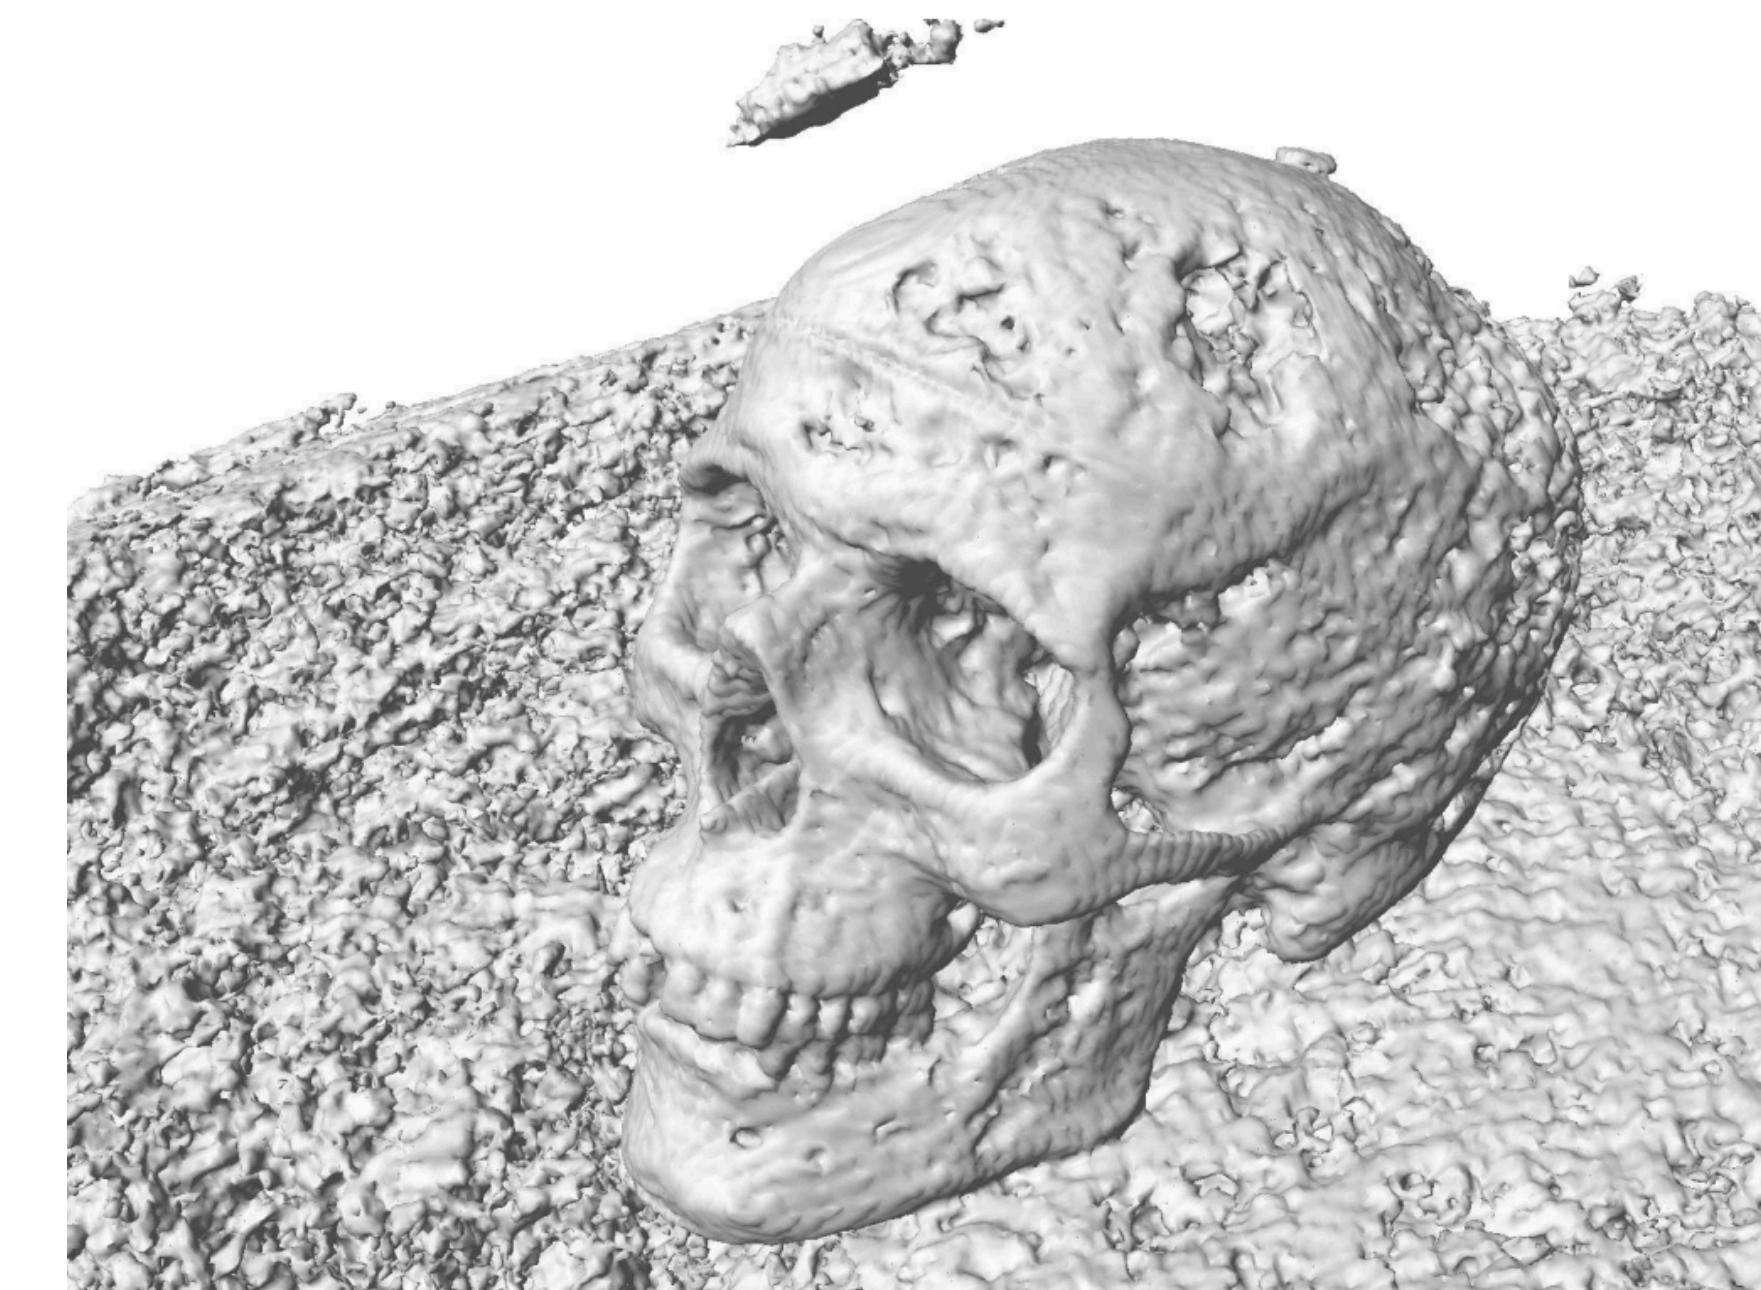


Neural Cage
(Wang et. al. 2020)

Quality of the Neural Fields Derivatives



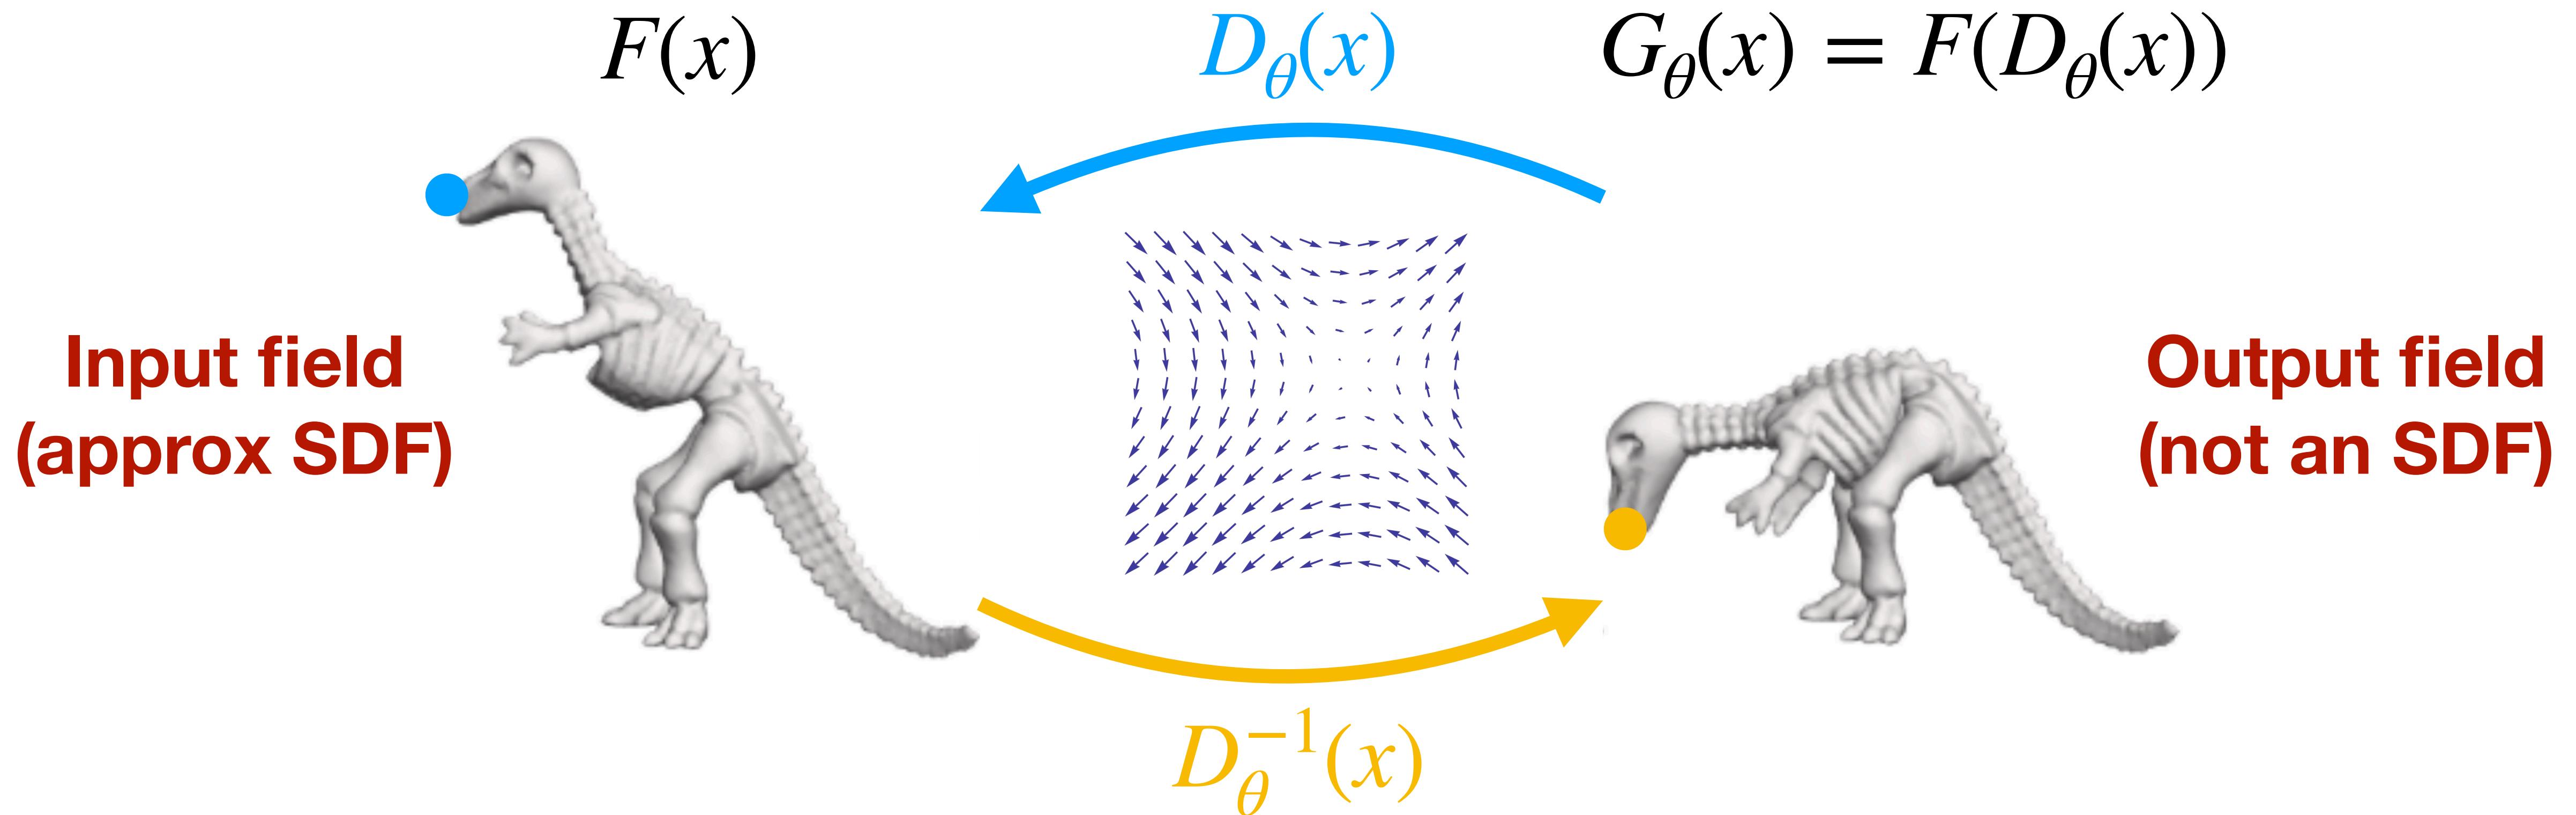
L: SIREN learned to fit the SDF of a square; R:
Zooming in on the surface.



NeRF Surface at certain threshold (not tuned)

What kind of neural fields can provide good enough gradient information to be trained for geometry processing tasks?

Preservation of Field Properties



IDEA 1: Can we incorporate a “reinitialization” schema to output an SDF?

IDEA 2: We do need the field to be an SDF in order to perform the task?

Physics Simulation

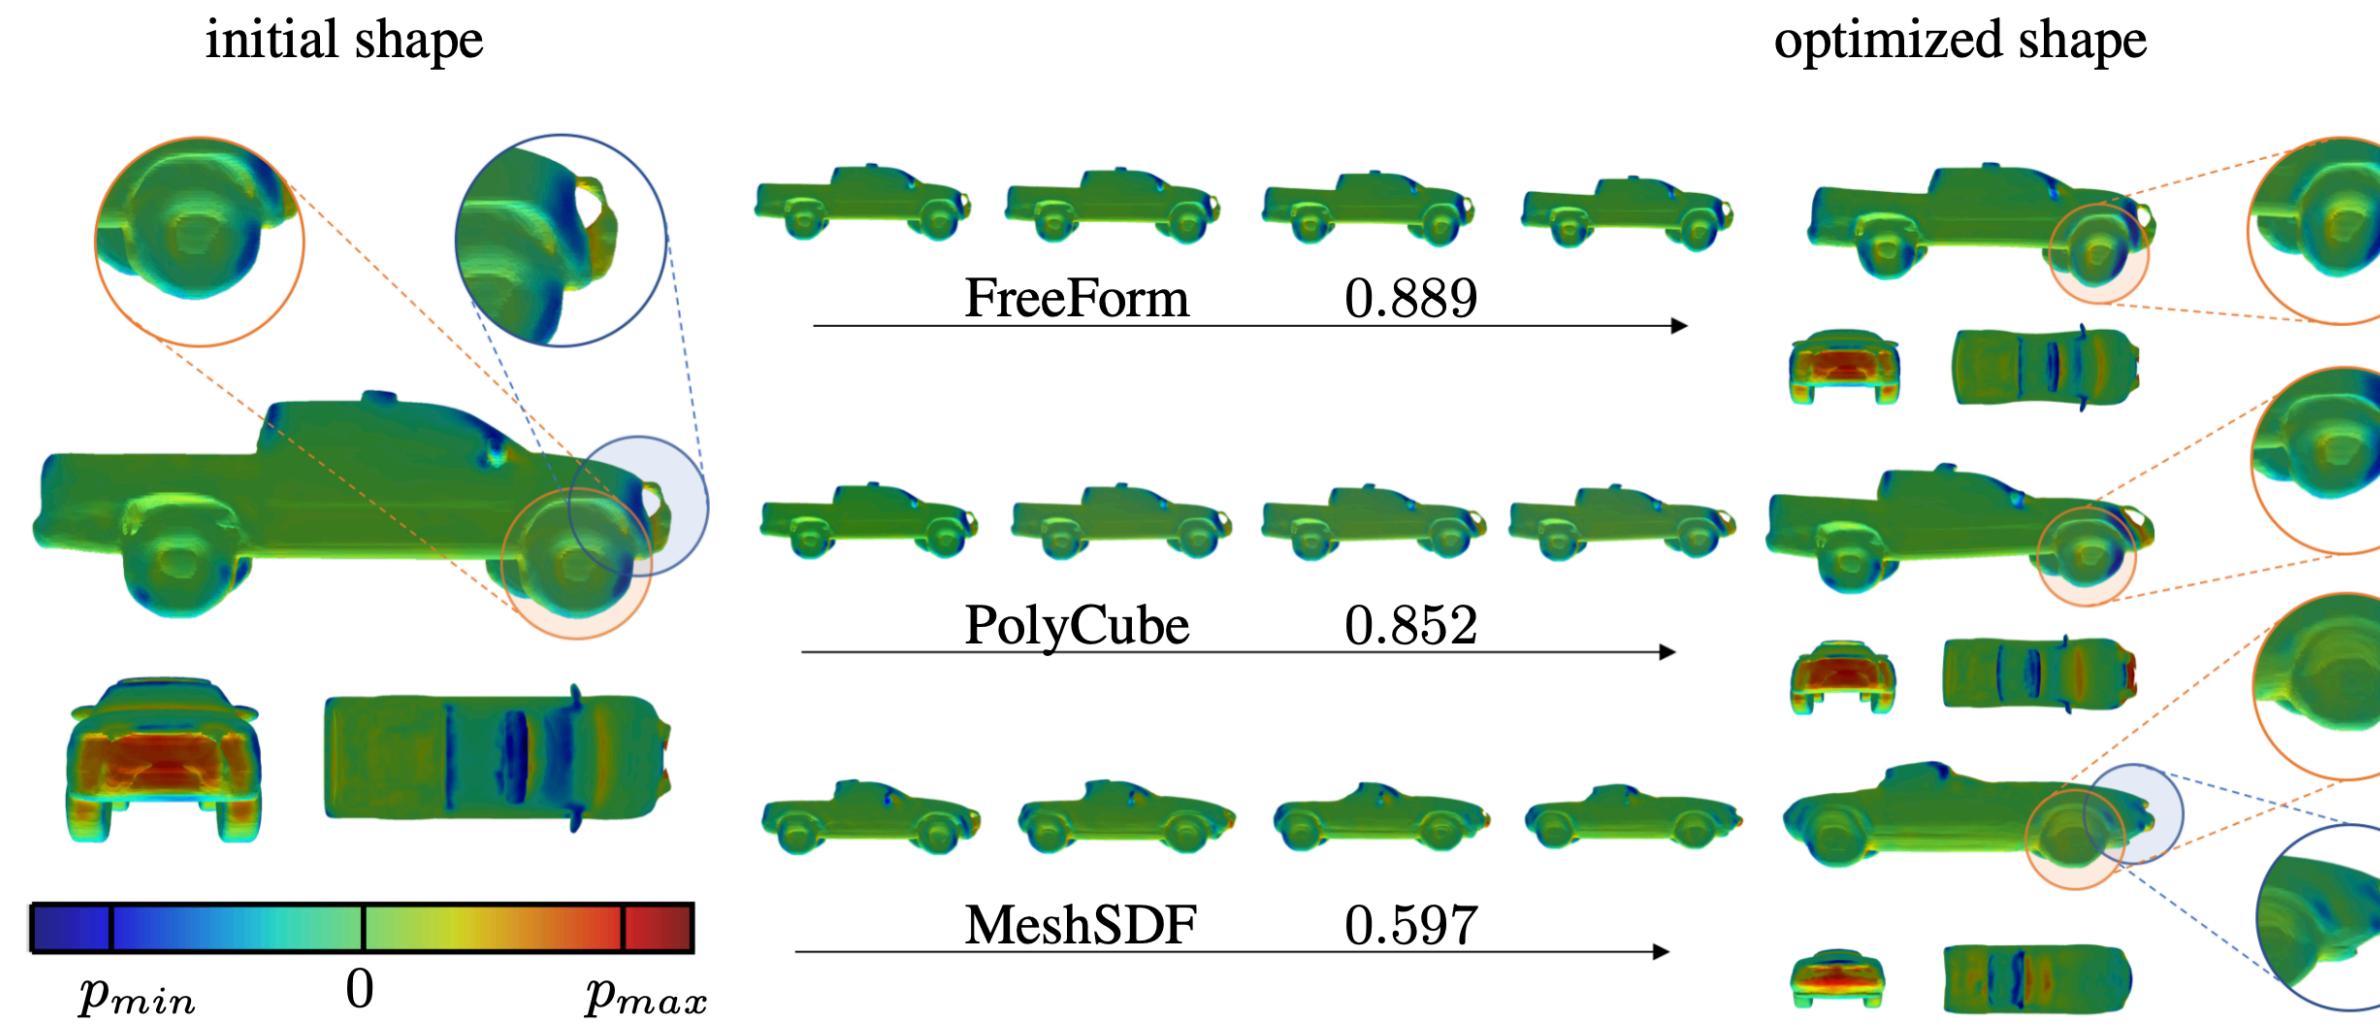


Figure 5: **Drag minimization.** Starting from an initial shape (left column), \mathcal{L}_{task} is minimized using three different parameterizations: FreeForm (top), PolyCube (middle), and our *MeshSDF* (bottom). The middle column depicts the optimization process and the relative improvements in terms of \mathcal{L}_{task} . The final result is shown in the right column. FreeForm and PolyCube lack a semantic prior, resulting in implausible details such as sheared wheels (orange inset). By contrast, *MeshSDF* not only enforces such priors but can also effect topology changes (blue inset).

T: MeshSDF (Remelli et. al. 2021)

R: NeuralODE (Chen et. al. 2019)

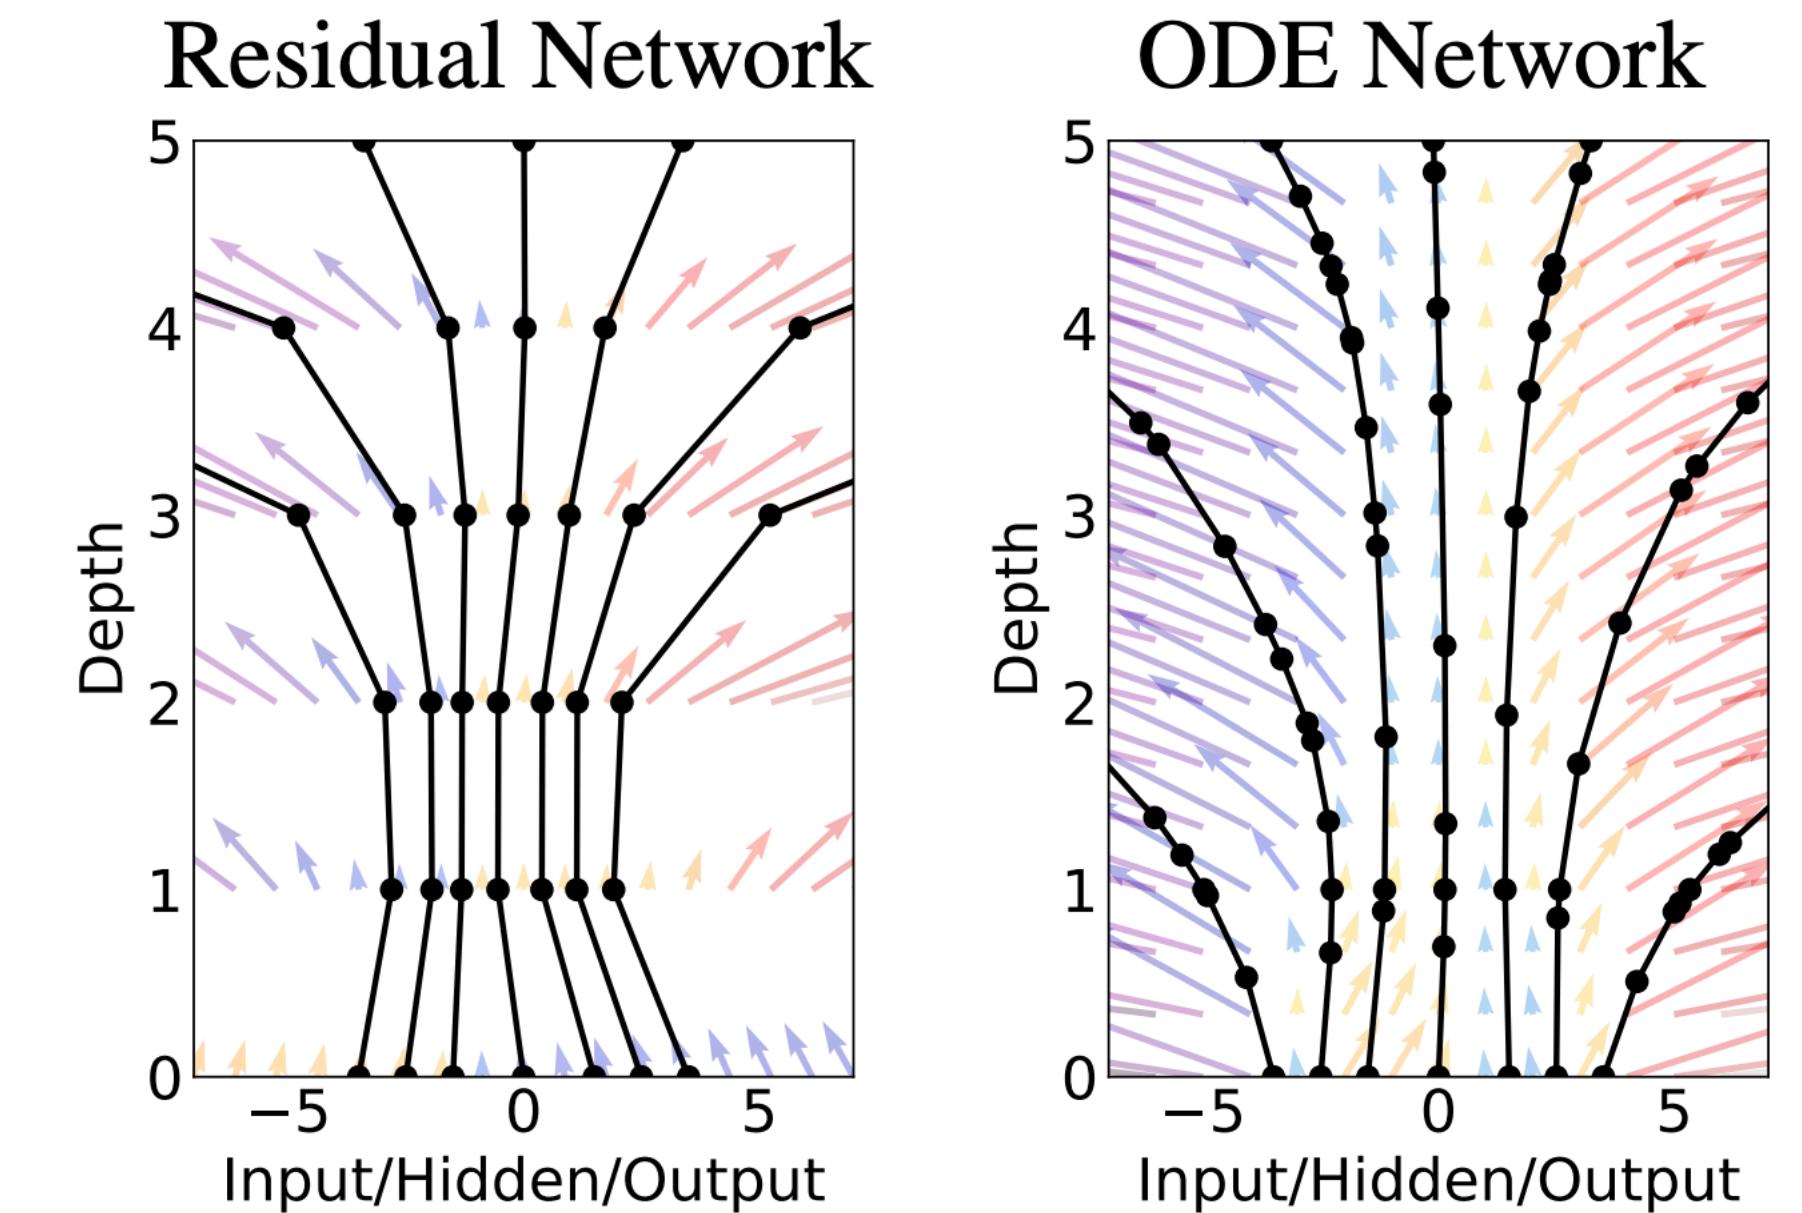


Figure 1: *Left:* A Residual network defines a discrete sequence of finite transformations. *Right:* A ODE network defines a vector field, which continuously transforms the state. *Both:* Circles represent evaluation locations.

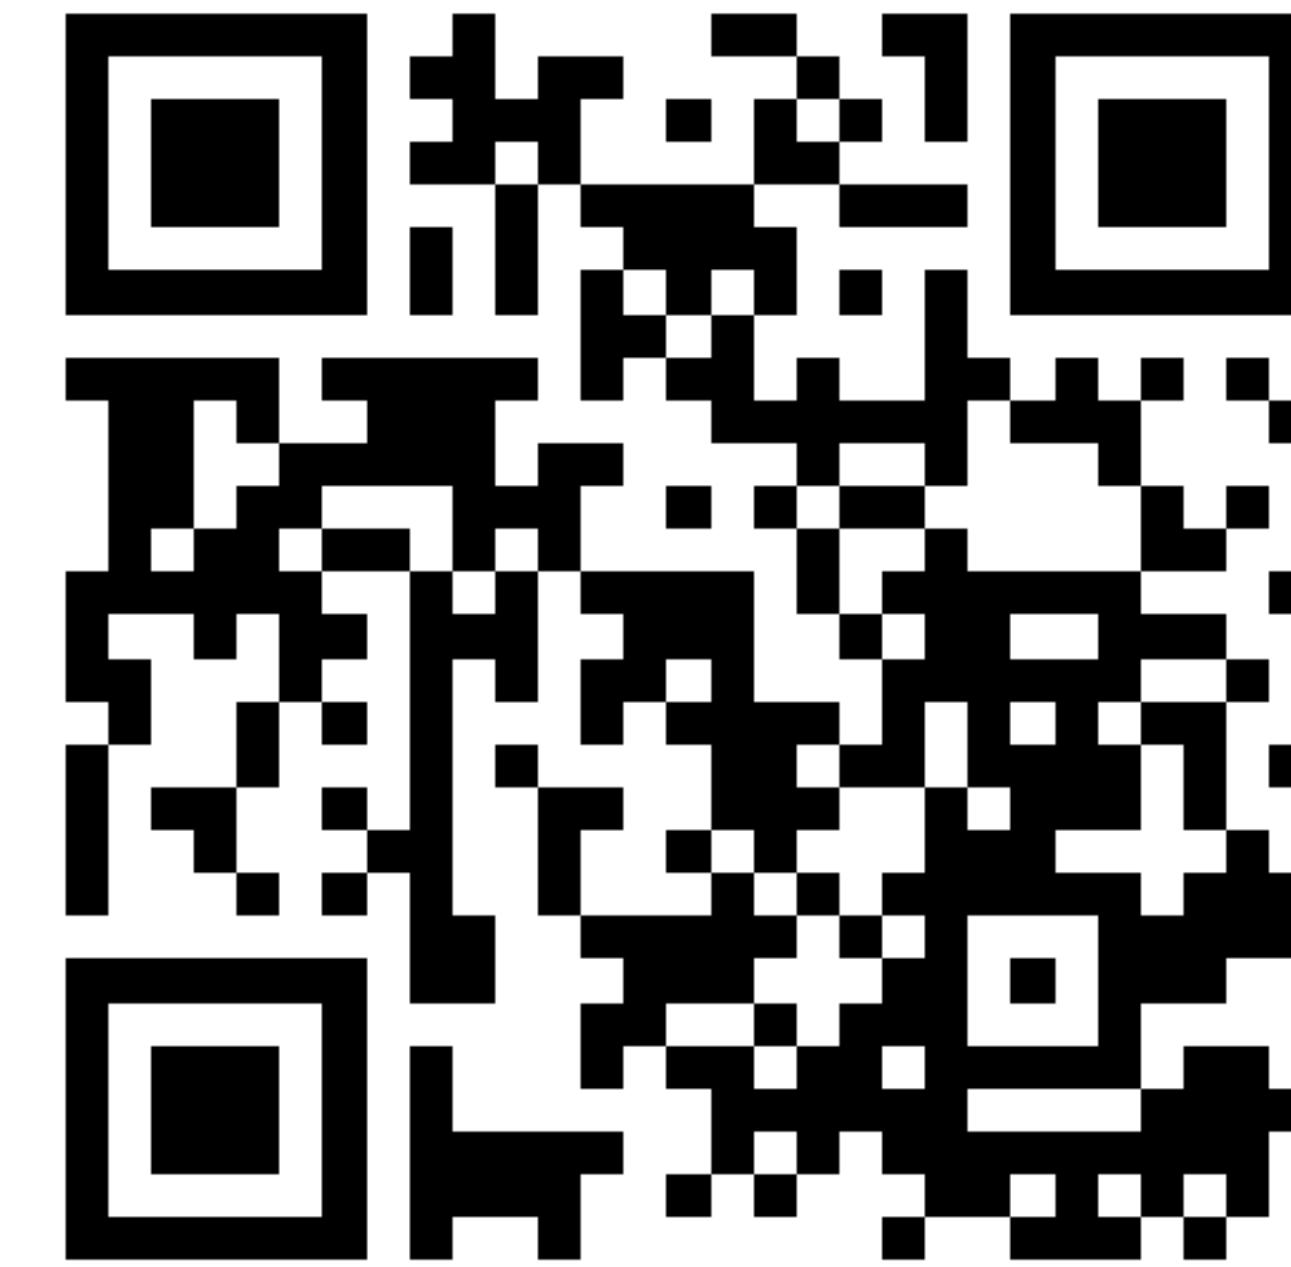
Thanks for your time!

Project Page



<https://www.guandaoyang.com/NFGP/>

Code



<https://github.com/stevenygd/NFGP>

Discussion questions

- Any application scenario where using Neural Fields can be advantageous?
(to represent the geometry, or potentially other physical quantities)
 - Need to represent a field in a compact way?
 - Need access to derivatives?
 - Can be benefited from data driven prior?
 - Might involve change of topology?

Discussion questions

- What do you see as a main obstacle to apply neural fields to problems you are interested in?
 - Optimization speed?
 - Numerical instability of the gradients?
 - Adding appearance/texture like mesh?