

# Equivariant Subgraph Aggregation Networks<sup>1</sup>

Beatrice Bevilacqua\*   Fabrizio Frasca\*   Derek Lim\*

Joint work with  
Balasubramaniam Srinivasan, Chen Cai, Gopinath Balamurugan,  
Michael M. Bronstein, Haggai Maron

December 14, 2021

---

<sup>1</sup>Project sparked from the 2021 LOGML Summer School

# Outline

1. Motivation
2. Method
3. Expressive power
4. Design choices
5. Experiments
6. Computational complexity
7. Conclusions and future work

# Outline

1. Motivation
2. Method
3. Expressive power
4. Design choices
5. Experiments
6. Computational complexity
7. Conclusions and future work

# Main idea

We define a new family of **provably powerful graph neural networks** that treat a **graph as a set of subgraphs**.

# Background: Graph Neural Networks (GNNs)

A **GNN** is a parameterized function on graphs.

## Definition (Message Passing Neural Network)

A message passing neural network (MPNN) updates node representations  $h_v^{(k)} \in \mathbb{R}^d$  by


$$\begin{aligned}\text{msg}_v^{(k)} &= \text{AGGREGATE}(\{h_u^{(k)} : u \text{ neighbor of } v\}), \\ h_v^{(k+1)} &= \text{COMBINE}\left(h_v^{(k)}, \text{msg}_v^{(k)}\right).\end{aligned}$$

**Note:** In their original formulation, they are permutation equivariant!

# MPNNs have limited expressive power

No matter how the AGGREGATE and COMBINE functions are parameterized, MPNNs are bounded in their expressive power (they are not universal!).

$\Rightarrow$  They cannot distinguish between non-isomorphic graphs beyond 1-WL [Xu et al., 2019]




# The 1-WL isomorphism test

A polynomially fast heuristic to disambiguate non-isomorphic graphs:

- Iterate node color refinements

$$c_v \leftarrow \text{HASH}(c_v, \{c_w\}_{w \in \mathcal{N}(v)})$$




# The 1-WL isomorphism test

A polynomially fast heuristic to disambiguate non-isomorphic graphs:

- Iterate node color refinements

$$c_v \leftarrow \text{HASH}(c_v, \{c_w\}_{w \in \mathcal{N}(v)})$$

- Represent graphs as color histograms






# The 1-WL isomorphism test

A polynomially fast heuristic to disambiguate non-isomorphic graphs:

- Iterate node color refinements

$$c_v \leftarrow \text{HASH}(c_v, \{c_w\}_{w \in \mathcal{N}(v)})$$

- Represent graphs as color histograms
- *distinct* histograms: *non-isomorphic* graphs



# The 1-WL isomorphism test


A polynomially fast heuristic to disambiguate non-isomorphic graphs:

- Iterate node color refinements

$$c_v \leftarrow \text{HASH}(c_v, \{c_w\}_{w \in \mathcal{N}(v)})$$

- Represent graphs as color histograms
- *distinct* histograms: *non-isomorphic* graphs


**Problem:** only a *sufficient* condition.



# Beyond 1-WL

Why is expressive power important?


- 1 1-WL cannot distinguish some molecular graphs (e.g. *decalin* vs. *bicyclopentyl*)



# Beyond 1-WL

Why is expressive power important?

- ① 1-WL cannot distinguish some molecular graphs (e.g. *decalin* vs. *bicyclopentyl*)
- ② 1-WL cannot count non-trivial patterns (e.g. *triangles* and *cycles*) [Chen et al., 2020]




# Beyond 1-WL

Why is expressive power important?

- ① 1-WL cannot distinguish some molecular graphs (e.g. *decalin* vs. *bicyclopentyl*)
- ② 1-WL cannot count non-trivial patterns (e.g. *triangles* and *cycles*) [Chen et al., 2020]


These limitations are directly inherited by MPNNs.



# GNNs beyond 1-WL

Currently a prolific research area, with diverse emerging approaches:


- k-WL inspired architectures: k-GNNs  
[Morris et al., 2019, Maron et al., 2019]
  - non-local
  - computationally intractable for  $k \geq 3$



# GNNs beyond 1-WL

Currently a prolific research area, with diverse emerging approaches:



- k-WL inspired architectures: k-GNNs  
[Morris et al., 2019, Maron et al., 2019]
  - non-local
  - computationally intractable for  $k \geq 3$
- Random node features / identifiers  
[Abboud et al., 2020, Sato et al., 2021]
  - no permutation equivariance
  - difficulties in generalisation



# GNNs beyond 1-WL

Currently a prolific research area, with diverse emerging approaches:

- k-WL inspired architectures: k-GNNs  
[Morris et al., 2019, Maron et al., 2019]
  - non-local
  - computationally intractable for  $k \geq 3$
- Random node features / identifiers  
[Abboud et al., 2020, Sato et al., 2021]
  - no permutation equivariance
  - difficulties in generalisation
- Substructure-aware message passing  
[Bouritsas et al., 2020, Bodnar et al., 2021]
  - some form of domain-knowledge is required






# Outline

1. Motivation
2. Method
3. Expressive power
4. Design choices
5. Experiments
6. Computational complexity
7. Conclusions and future work


# Graphs as sets of subgraphs



## Desiderata:

- Provable expressive power ( $> 1\text{-WL}$ )
- (Almost) no engineering required
- Computationally tractable
- Equivariance to permutations

# Graphs as sets of subgraphs



## Desiderata:

- Provable expressive power ( $> 1\text{-WL}$ )
- (Almost) no engineering required
- Computationally tractable
- Equivariance to permutations

**Intuition:** Even if two graphs are indistinguishable to MPNNs, they may contain subgraphs which are, instead, (easily) separated.

# Equivariant Subgraph Aggregation Networks (ESAN)

## Recipe?

- 1 Map a graph into a set of subgraphs (bag) via a *selection policy*:  $G \mapsto \{G_1, \dots, G_m\}$
- 2 Process the bag in a principled way: respecting the *inherent symmetries* of such object


# Equivariant Subgraph Aggregation Networks (ESAN)

## Recipe?

- ① Map a graph into a set of subgraphs (bag) via a *selection policy*:  $G \mapsto \{G_1, \dots, G_m\}$
- ② Process the bag in a principled way: respecting the *inherent symmetries* of such object

### Policies:

- Node Deletion (ND)
- Edge Deletion (ED)
- Ego-Nets (EGO)
- ...



# Background: Equivariance

Let  $f : \mathbb{R}^{d_1} \rightarrow \mathbb{R}^{d_2}$  and  $H$  be a symmetry group acting on  $\mathbb{R}^{d_1}$  and  $\mathbb{R}^{d_2}$ .

## Definition

$f$  is **invariant** if  $f(h \cdot x) = f(x)$  for all  $h \in H$ .

$f$  is **equivariant** if  $f(h \cdot x) = h \cdot f(x)$  for all  $h \in H$ .

# Background: Equivariance

Let  $f : \mathbb{R}^{d_1} \rightarrow \mathbb{R}^{d_2}$  and  $H$  be a symmetry group acting on  $\mathbb{R}^{d_1}$  and  $\mathbb{R}^{d_2}$ .

## Definition

$f$  is **invariant** if  $f(h \cdot x) = f(x)$  for all  $h \in H$ .

$f$  is **equivariant** if  $f(h \cdot x) = h \cdot f(x)$  for all  $h \in H$ .


## Example (Symmetries in Graphs)

Let  $G$  a graph with adjacency matrix  $A \in \mathbb{R}^{n \times n}$  and node features  $X \in \mathbb{R}^{n \times d}$ .

$S_n$  (permutations on  $n$  elements) acts on  $G$  by renaming node  $i$  to  $\sigma(i)$ :

$$(\sigma \cdot A)_{ij} = A_{\sigma^{-1}(i), \sigma^{-1}(j)} \quad (\sigma \cdot X)_{ij} = X_{\sigma^{-1}(i), j} \quad (1)$$

# Required equivariance



We want equivariance to  $S_m \times S_n$ .

$\tau \in S_m$  permutes graphs in the set.

$\sigma \in S_n$  permutes nodes in each subgraph.

**Main idea: node alignment.**  $\sigma \in S_n$  same across all subgraphs. Node  $i$  in each subgraph corresponds to node  $i$  in original graph.



# Our models

## Definition (DS-GNN)

DS-GNN is Siamese:

$$\text{DeepSets}(\{\text{MPNN}(A_1, X_1), \dots, \text{MPNN}(A_m, X_m)\}) \quad (2)$$

## Definition (DSS-GNN)

DSS-GNN defines layers  $L$  where:

$$L(\mathbf{A}, \mathbf{X})_i = \text{MPNN}_1(A_i, X_i) + \text{MPNN}_2\left(\sum_{j=1}^m A_j, \sum_{j=1}^m X_j\right) \quad (3)$$

**Main idea:** DSS preserves node alignment ( $S_m \times S_n$ ), DS-GNN does not ( $S_m \wr S_n$ ).

# Outline

1. Motivation
2. Method
3. Expressive power
4. Design choices
5. Experiments
6. Computational complexity
7. Conclusions and future work

# “DSS-WL”

- Get a bag of subgraphs

$$\mathcal{G} \leftarrow \pi(G)$$

# “DSS-WL”

- Get a bag of subgraphs

$$\mathcal{G} \leftarrow \pi(G)$$


- Augmented node color refinements

$$c_v^S \leftarrow \text{HASH}(c_v^S, N_v^S, \textcolor{violet}{C}_v, \textcolor{teal}{M}_v)$$

$$N_v^S = \{c_w^S\}_{w \in \mathcal{N}^S(v)} \quad \triangleleft \text{adjacent colors on } S$$

$$\textcolor{violet}{C}_v = \{c_v^R\}_{R \in \mathcal{G}} \quad \triangleleft \textit{needle-color}$$

$$\textcolor{teal}{M}_v = \{C_w\}_{w \in \mathcal{N}(v)} \quad \triangleleft \text{adjacent \textit{needle}-colors on } G$$



# “DSS-WL”

- Get a bag of subgraphs

$$\mathcal{G} \leftarrow \pi(G)$$

- Augmented node color refinements

$$c_v^S \leftarrow \text{HASH}(c_v^S, N_v^S, C_v, M_v)$$


$$N_v^S = \{c_w^S\}_{w \in \mathcal{N}^S(v)} \quad \triangleleft \text{adjacent colors on } S$$

$$C_v = \{C_v^R\}_{R \in \mathcal{G}} \quad \triangleleft \text{needle-color}$$

$$M_v = \{C_w\}_{w \in \mathcal{N}(v)} \quad \triangleleft \text{adjacent needle-colors on } G$$

- Subgraph color: HASH the multiset of node colors

$$c_S \leftarrow \text{HASH}(\{c_v^S\}_{v \in S})$$



# “DSS-WL”

- Get a bag of subgraphs

$$\mathcal{G} \leftarrow \pi(G)$$

- Augmented node color refinements

$$c_v^S \leftarrow \text{HASH}(c_v^S, N_v^S, \mathbf{C}_v, \mathbf{M}_v)$$


$$N_v^S = \{c_w^S\}_{w \in \mathcal{N}^S(v)} \quad \triangleleft \text{adjacent colors on } S$$

$$\mathbf{C}_v = \{c_v^R\}_{R \in \mathcal{G}} \quad \triangleleft \text{needle-color}$$

$$\mathbf{M}_v = \{C_w\}_{w \in \mathcal{N}(v)} \quad \triangleleft \text{adjacent needle-colors on } G$$

- Subgraph color: HASH the multiset of node colors

$$c_S \leftarrow \text{HASH}(\{c_v^S\}_{v \in S})$$



Eventually, represent the original graph as the *histogram of subgraph colors*.

# WL-variants and expressive power

Variant	Policy	Hash inps.	Neural
DSS-WL	(any)	$c_v^S, N_v^S, C_v, M_v$	DSS-GNN
DS-WL	(any)	$c_v^S, N_v^S$	DS-GNN
1-WL	$G \mapsto \{G\}$	$c_v^S, N_v^S$	GNN

- DS-WL: disable cross-bag info sharing
  - independent 1-WL on each subgraph
- 1-WL: employ trivial policy  $G \mapsto \{G\}$

<sup>2</sup>on families of bounded-sized graphs; DSS: edge set-preserving policies

# WL-variants and expressive power

Variant	Policy	Hash inps.	Neural
DSS-WL	(any)	$c_v^S, N_v^S, C_v, M_v$	DSS-GNN
DS-WL	(any)	$c_v^S, N_v^S$	DS-GNN
1-WL	$G \mapsto \{G\}$	$c_v^S, N_v^S$	GNN

- DS-WL: disable cross-bag info sharing
  - independent 1-WL on each subgraph
- 1-WL: employ trivial policy  $G \mapsto \{G\}$

**Theorem:** There exist policies such that DS(S)-WL  $>$  1-WL.

**Theorem:** DS(S)-GNN can implement DS(S)-WL<sup>2</sup>. Also: DS-GNN(MPNN)  $\leq$  DS-WL.

**Corollary:** There exist policies such that DS(S)-GNN  $>$  MPNN.

<sup>2</sup>on families of bounded-sized graphs; DSS: edge set-preserving policies



# Expressiveness experimental validation

**Experiments:** DSS-GNN and DS-GNN are perfect on RNI and CSL datasets (which require  $> 1$ -WL power).


	EXP	CEXP
GIN [Xu et al., 2019]	51.1 $\pm$ 2.1	70.2 $\pm$ 4.1
<b>DS-GNN (GIN) (ED/ND/EGO/EGO+)</b>	100 $\pm$ 0.0	100 $\pm$ 0.0
<b>DSS-GNN (GIN) (ED/ND/EGO/EGO+)</b>	100 $\pm$ 0.0	100 $\pm$ 0.0
GRAPHCONV [Morris et al., 2019]	50.3 $\pm$ 2.6	72.9 $\pm$ 3.6
<b>DS-GNN (GraphConv) (ED/ND/EGO/EGO+)</b>	100 $\pm$ 0.0	100 $\pm$ 0.0
<b>DSS-GNN (GraphConv) (ED/ND/EGO/EGO+)</b>	100 $\pm$ 0.0	100 $\pm$ 0.0

# Outline

1. Motivation
2. Method
3. Expressive power
4. Design choices
5. Experiments
6. Computational complexity
7. Conclusions and future work

# DSS-GNN > DS-GNN

**Proposition:** DSS-GNN > DS-GNN for certain policies.



Subgraph policy: map  $G$  to set of single edges.

DSS-GNN can reconstruct graph  $A = \sum_j A_j$  because of node alignment.  
 DS-GNN only sees number of edges (no alignment).

# DSS-GNN > DS-GNN

**Experiments:** DSS-GNN tends to outperform DS-GNN.

DSS-GNN is better than base encoder **91%** of the time.


DS-GNN is better than base encoder **75%** of the time.

# Subgraph policy matters

**Proposition:** On the family of *strongly regular graphs*:

$$\text{Edge-deletion} > \text{Node-deletion} = \text{Depth-}n \text{ EGO}(+) = 3\text{-WL}.$$

*Strongly regular graphs*: highly-symmetric graphs that are hard cases for graph isomorphism.



## Subgraph policy choice in the experiments

No strong correlation between the performance of a policy and the application domain.  
The EGO(+) policies were found generally more consistent in their results across datasets.

## Subgraph policy choice in the experiments

No strong correlation between the performance of a policy and the application domain.  
The EGO(+) policies were found generally more consistent in their results across datasets.

	ZINC ↓
<b>DS-GNN (GIN) (ED)</b>	0.172±0.008
<b>DS-GNN (GIN) (ND)</b>	0.171±0.010
<b>DS-GNN (GIN) (EGO)</b>	0.126±0.006
<b>DS-GNN (GIN) (EGO+)</b>	0.116±0.009
<b>DSS-GNN (GIN) (ED)</b>	0.172±0.005
<b>DSS-GNN (GIN) (ND)</b>	0.166±0.004
<b>DSS-GNN (GIN) (EGO)</b>	0.107±0.005
<b>DSS-GNN (GIN) (EGO+)</b>	0.102±0.003

## Subgraph policy choice in the experiments


No strong correlation between the performance of a policy and the application domain.  
The EGO(+) policies were found generally more consistent in their results across datasets.

	ZINC ↓
DS-GNN (GIN) (ED)	0.172±0.008
DS-GNN (GIN) (ND)	0.171±0.010
DS-GNN (GIN) (EGO)	0.126±0.006
DS-GNN (GIN) (EGO+)	0.116±0.009
DSS-GNN (GIN) (ED)	0.172±0.005
DSS-GNN (GIN) (ND)	0.166±0.004
DSS-GNN (GIN) (EGO)	0.107±0.005
DSS-GNN (GIN) (EGO+)	0.102±0.003



# Base encoder matters

**Proposition:** Using a 3-WL base encoder is stronger than using a 1-WL (MPNN) base encoder.



# Outline

1. Motivation
2. Method
3. Expressive power
4. Design choices
- 5. Experiments**
6. Computational complexity
7. Conclusions and future work

# ESAN outperforms the base encoder

	OGBG-MOLHIV	OGBG-MOLTOX21
BASE	75.58 $\pm$ 1.40	74.91 $\pm$ 0.51
<b>ESAN</b>	78.00 $\pm$ 1.42	77.95 $\pm$ 0.40

**ESAN<sup>3</sup> generally improves the base encoder accuracy** (e.g. 2.4% improvement on OGBG-MOLHIV and 3% improvement on OGBG-MOLTOX21)

<sup>3</sup>In this section, we report the performance of our best ESAN model

# ESAN is competitive with SoTA

	MUTAG	PTC	PROTEINS	NCI1	NCI109	IMDB-B
GIN [Xu et al., 2019]	89.4±5.6	64.6±7.0	76.2±2.8	82.7±1.7	82.2±1.6	75.1±5.1
GRAPHCONV [Morris et al., 2019]	90.5±4.6	64.9±10.4	73.9±6.1	82.4±2.7	81.7±1.0	76.1±3.9
PPGNs [Maron et al., 2019]	90.6±8.7	66.2±6.6	77.2±4.7	83.2±1.1	82.2±1.4	73.0±5.8
GSN [Bouritsas et al., 2020]	92.2±7.5	68.2±7.2	76.6±5.0	83.5±2.0	N/A	<b>77.8</b> ±3.3
CIN [Bodnar et al., 2021]	<b>92.7</b> ±6.1	68.2±5.6	77.0±4.3	83.6±1.4	<b>84.0</b> ±1.6	75.6±3.7
<b>ESAN</b>	92.0±5.0	<b>69.2</b> ±6.5	<b>77.3</b> ±3.8	<b>83.8</b> ±2.4	83.1±0.8	77.1±3.0

ESAN achieves excellent results with respect to SoTA methods

# ESAN on ZINC

**Best performing model amongst all provably expressive, domain agnostic GNNs**

	ZINC ↓
PNA [Corso et al., 2020]	$0.188 \pm 0.004$
DGN [Beaini et al., 2021]	$0.168 \pm 0.003$
SMP [Vignac et al., 2020]	$0.138 \pm ?$
GIN [Xu et al., 2019]	$0.252 \pm 0.017$
HIMP [Fey et al., 2020]	$0.151 \pm 0.006$
GSN [Bouritsas et al., 2020]	$0.108 \pm 0.018$
CIN-SMALL [Bodnar et al., 2021]	$0.094 \pm 0.004$
<b>ESAN</b>	<b><math>0.102 \pm 0.003</math></b>

# ESAN on ZINC

**Best performing** model amongst all provably expressive, **domain agnostic** GNNs  
**Competitive with** (provably powerful) GNNs employing **domain specific** structural information, often outperforming them.

	ZINC ↓
PNA [Corso et al., 2020]	0.188±0.004
DGN [Beaini et al., 2021]	0.168±0.003
SMP [Vignac et al., 2020]	0.138±?
GIN [Xu et al., 2019]	0.252±0.017
HIMP [Fey et al., 2020]	0.151±0.006
GSN [Bouritsas et al., 2020]	0.108±0.018
CIN-SMALL [Bodnar et al., 2021]	0.094±0.004
<b>ESAN</b>	<b>0.102±0.003</b>

# Outline

1. Motivation
2. Method
3. Expressive power
4. Design choices
5. Experiments
6. Computational complexity
7. Conclusions and future work

# Complexity

**Table:** Complexity of graph networks that are more expressive than 1-WL.  $\Delta_{\max}$  denotes the maximum degree over all nodes.

	PPGN	3-IGN	3-GNN	Ours
Time	$\mathcal{O}(n^3)$	$\mathcal{O}(n^3)$	$\mathcal{O}(n^4)$	$\mathcal{O}( S n\Delta_{\max})$
Space	$\mathcal{O}(n^2)$	$\mathcal{O}(n^3)$	$\mathcal{O}(n^3)$	$\mathcal{O}( S (n + n\Delta_{\max}))$

**Table:** Bag sizes and bag construction complexity.

	$ S $	Time complexity
ND	$\Theta(n)$	$\mathcal{O}(n^2\Delta_{\max})$
ED	$\mathcal{O}(n\Delta_{\max})$	$\mathcal{O}((n\Delta_{\max})^2)$
EGO(+)	$\Theta(n)$	$\mathcal{O}(n(n + n\Delta_{\max}))$



# Sampling approach

To reduce computational cost, sample a **subset of the subgraphs** for each graph.  
At inference: majority voting on the outputs of  $k$  independent samplings ( $k = 5$ ).

# Sampling approach

To reduce computational cost, sample a **subset of the subgraphs** for each graph.  
At inference: majority voting on the outputs of  $k$  independent samplings ( $k = 5$ ).

Benefits:

- ① Enables running on larger graphs
- ② Allows increasing the batch size resulting in faster runtime

# Sampling experiments

Sampling does not significantly reduce accuracy, and in some cases even **increases accuracy**

		OGBG-MOLHIV	OGBG-MOLTOX21	EXP	CEXP
GIN		$75.58 \pm 1.40$	$74.91 \pm 0.51$	$51.2 \pm 2.1$	$70.2 \pm 4.1$
DS-GNN (GIN) (ED)	100%	$76.43 \pm 2.12$	$75.12 \pm 0.50$	$100 \pm 0.0$	$100 \pm 0.0$
	50%	$76.29 \pm 1.33$	$74.59 \pm 0.71$	$100 \pm 0.0$	$100 \pm 0.0$
	20%	$76.57 \pm 1.48$	$75.67 \pm 0.89$	$100 \pm 0.0$	$99.9 \pm 0.2$
	5%	$77.82 \pm 1.00$	$76.39 \pm 1.11$	$99.7 \pm 0.4$	$99.9 \pm 0.2$
DS-GNN (GIN) (ND)	100%	$76.19 \pm 0.96$	$75.34 \pm 1.21$	$100 \pm 0.0$	$100 \pm 0.0$
	50%	$77.23 \pm 1.32$	$74.82 \pm 1.05$	$100 \pm 0.0$	$99.9 \pm 0.2$
	20%	$77.65 \pm 0.84$	$75.66 \pm 0.46$	$100 \pm 0.0$	$99.9 \pm 0.2$
	5%	$78.26 \pm 1.02$	$76.51 \pm 1.04$	$97.2 \pm 1.1$	$99.8 \pm 0.8$

# Runtime

In practice, the additional computational complexity is empirically tractable.

**Table:** Timing comparison (in seconds) per epoch on a RTX 2080 GPU.

	NCI1			ZINC	
	BASELINE	100% SUBGRAPHS	20% SUBGRAPHS	BASELINE	100% SUBGRAPHS
GIN	1.00±0.05	-	-	1.45±0.01	-
<b>DS-GNN (GIN) (ED)</b>	-	2.07±0.01	1.73±0.01	-	3.56±0.03
<b>DS-GNN (GIN) (ND)</b>	-	2.08±0.03	1.71±0.01	-	3.42±0.02
<b>DS-GNN (GIN) (EGO)</b>	-	1.96±0.01	1.72±0.01	-	3.02±0.04
<b>DS-GNN (GIN) (EGO+)</b>	-	2.01±0.02	1.73±0.01	-	3.09±0.04
<b>DSS-GNN (GIN) (ED)</b>	-	3.26±0.07	2.94±0.01	-	4.25±0.01
<b>DSS-GNN (GIN) (ND)</b>	-	3.19±0.07	2.91±0.02	-	4.12±0.03
<b>DSS-GNN (GIN) (EGO)</b>	-	3.14±0.04	2.79±0.02	-	3.63±0.02
<b>DSS-GNN (GIN) (EGO+)</b>	-	3.19±0.03	2.90±0.01	-	3.69±0.05

# Outline

1. Motivation
2. Method
3. Expressive power
4. Design choices
5. Experiments
6. Computational complexity
7. Conclusions and future work

## Other similar methods

Very recently, other methods of various motivations have been proposed:

- Reconstruction GNN [Cotta et al., 2021]  $\approx$  DS-GNN (ND)
- DropGNN [Papp et al., 2021]  $\approx$  DS-GNN (ND) with sampling
- GNN-AK [Zhao et al., 2021]  $\approx$  DSS-GNN (EGO)

Our ESAN framework **unifies** and **generalizes** these methods with bags of subgraphs and equivariance!!

# Conclusions

## The ESAN framework:

- ① decomposes a graph into a set of subgraphs (bag)
- ② processes the bag respecting the emerging symmetry

Experimentally: competitive performance on various graph-wise benchmarks





ESAN naturally enjoys:

- provable expressiveness
- native equivariance
- being domain-agnostic

Future directions:


- other policies?
- structured bags
- stochastic policies & approx. equivariance


# References I


-  Abboud, R., Ceylan, I. I., Grohe, M., and Lukasiewicz, T. (2020).  
 The surprising power of graph neural networks with random node initialization.  
*arXiv preprint arXiv:2010.01179.*
-  Beaini, D., Passaro, S., Létourneau, V., Hamilton, W. L., Corso, G., and Liò, P. (2021).  
 Directional graph networks.  
*In International Conference on Machine Learning.*
-  Bodnar, C., Frasca, F., Otter, N., Wang, Y. G., Liò, P., Montúfar, G., and Bronstein, M. (2021).  
 Weisfeiler and lehman go cellular: CW networks.  
*In Advances in Neural Information Processing Systems*, volume 34.
-  Bouritsas, G., Frasca, F., Zafeiriou, S., and Bronstein, M. M. (2020).  
 Improving graph neural network expressivity via subgraph isomorphism counting.  
*arXiv preprint arXiv:2006.09252.*




# References II

 Chen, Z., Chen, L., Villar, S., and Bruna, J. (2020).  
Can graph neural networks count substructures?  
*Advances in neural information processing systems.*

 Corso, G., Cavalleri, L., Beaini, D., Liò, P., and Veličković, P. (2020).  
Principal neighbourhood aggregation for graph nets.  
*In Advances in Neural Information Processing Systems.*

 Cotta, L., Morris, C., and Ribeiro, B. (2021).  
Reconstruction for powerful graph representations.  
*In NeurIPS.*

 Fey, M., Yuen, J. G., and Weichert, F. (2020).  
Hierarchical inter-message passing for learning on molecular graphs.  
*In ICML Graph Representation Learning and Beyond (GRL+) Workshop.*

# References III



Maron, H., Ben-Hamu, H., Serviansky, H., and Lipman, Y. (2019).

Provably powerful graph networks.

In *NeurIPS*.



Morris, C., Ritzert, M., Fey, M., Hamilton, W. L., Lenssen, J. E., Rattan, G., and Grohe, M. (2019).

Weisfeiler and leman go neural: Higher-order graph neural networks.

In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 33, pages 4602–4609.



Papp, P. A., Martinkus, K., Faber, L., and Wattenhofer, R. (2021).

Dropgnn: Random dropouts increase the expressiveness of graph neural networks.

In *35th Conference on Neural Information Processing Systems (NeurIPS)*.



Sato, R., Yamada, M., and Kashima, H. (2021).

Random features strengthen graph neural networks.

In *Proceedings of the 2021 SIAM International Conference on Data Mining (SDM)*, pages 333–341.  
SIAM.

# References IV



Vignac, C., Loukas, A., and Frossard, P. (2020).

Building powerful and equivariant graph neural networks with structural message-passing.

In *NeurIPS*.



Xu, K., Hu, W., Leskovec, J., and Jegelka, S. (2019).

How powerful are graph neural networks?

In *International Conference on Learning Representations*.



Zhao, L., Jin, W., Akoglu, L., and Shah, N. (2021).

From stars to subgraphs: Uplifting any gnn with local structure awareness.